

MARKET EVOLUTION / SITUATIONAL AWARENESS

ODYLITH

Intelligence abundance, institutional authority,
market formation, and survival under superintelligence

NON-PRESCRIPTIVE WORLD MODEL

"Do not defend a thesis. Map the worlds that could invalidate it."

Evidence cutoff / 5 August 2026

Prepared for founder-level strategic inquiry

Observed facts, announced claims, inference, speculation, unknowns, and falsifiers are separated

Contents

Mandate: build a world model before choosing a strategy	9
How to read the epistemic labels	10
Prior-dossier anchoring audit	11
Second taxonomy: actors, welfare, coercion, and non-market allocation	12
Actor and welfare atlas: a taxonomy blind to the product perimeter	13
W1. Households and consumers	14
W2. Workers	14
W3. Professions and fiduciaries	14
W4. Communities hosting infrastructure	15
W5. Enterprise owners and boards	15
W6. Function managers and domain teams	15
W7. States and national-security institutions	15
W8. Courts and dispute institutions	16
W9. Regulators and supervisors	16
W10. Insurers and risk pools	16
W11. Public digital institutions	16
W12. Cooperatives and consortia	17
W13. Criminal and informal networks	17
W14. Energy, compute, and physical operators	17
W15. Developing economies and late adopters	17
W16. Future generations	18
W17. Machine-managed entities	18
W18. Ecological systems and non-human interests	18
Present topology: hierarchy before market	19
Evidence baseline as of 5 August 2026	21
State variables and causal primitives	22

Causal loop 1: cheap intelligence, rebound, and total spend	24
Causal loop 2: scarcity migration is not guaranteed	25
Enterprise evolution: from use to captured economics	26
Provider evolution: capability, distribution, and enclosure	27
Open standards: interoperability is not contestability	28
When agent activity becomes a market	29
Machine commerce: a contingent sequence	30
Scenario library: eighteen causal worlds	31
S1. Dominant-provider regulated utility	32
S2. Frontier oligopoly	32
S3. Open intelligence commons	33
S4. Sovereign agent blocs	33
S5. Enterprise-owned agent fleets	34
S6. Embodied autonomy	34
S7. Compute- and power-scarce acceleration	35
S8. Near-free compute and intelligence	35
S9. Institutionally trusted superintelligence	35
S10. Capable but institutionally untrusted SI	36
S11. Self-verifying intelligence	36
S12. Opaque high-performance systems	37
S13. Strategically deceptive systems	37
S14. Interface-obsoleting intelligence	38
S15. Slow takeoff and institutional coevolution	38
S16. Fast takeoff under existing institutions	38
S17. Discontinuous superintelligence shock	39
S18. Fragmented polycentric mosaic	40
Cross-product world lattice and omitted equilibria	41
Public-good and cooperative equilibrium	42
Institutional fracture	42

Labor and public veto	42
Synthetic-demand spiral	42
Ecological and social constraint	42
Informal, criminal, and adversarial agent economy	42
Machine legal standing and autonomous assets	42
Catastrophic infrastructure regression	42
Non-market abundance and political allocation	43
Transition points that discriminate worlds	44
Counterfactual reconstruction: reverse the attractive claims	45
Hidden confounders and partial observability	47
Game theory: where the mechanism becomes unstable	48
What superintelligence resistance actually means	49
Superintelligence resistance matrix	50
Antifragility under model and institutional shocks	52
Real options and temporal strategy without selecting a product	53
Institutional opportunity atlas: functions without an assumed company	54
O1. Model and agent routing	57
O2. Orchestration and workflow state	58
O3. Context custody	59
O4. Agent identity	59
O5. Delegated authorization	60
O6. Provenance and custody receipts	61
O7. Independent verification	62
O8. Agent observability and incident response	62
O9. Policy and governance	63
O10. AI FinOps	64
O11. Machine payments	64
O12. Procurement auctions and RFQs	65

O13. Portable reputation	66
O14. Insurance and risk transfer	67
O15. Dispute resolution	67
O16. Settlement and clearing	68
O17. Data and content rights	69
O18. Sovereign infrastructure	69
O19. Compute, power, and capacity	70
O20. Physical oracles and actuation	71
O21. Enterprise transformation	72
O22. Systems-of-record distribution	72
O23. Certification and accreditation	73
O24. Protocol infrastructure	74
O25. Human consent and fiduciary control	74
Revenue mechanisms by market stage	76
Opportunities that disappear as intelligence improves	77
Functions strengthened by consequence-adjusted autonomous exposure	78
Cross-world robustness properties	79
Capital regimes and investor appetite	80
Leading-indicator system	81
Research and experiment agenda	83
Evidence and confidence ledger	84
Adversarial review architecture	85
Round 3 correction-closure ledger	86
Round 2 cross-examination: market economics	87
Scope and method	87
Empirical anchor	87
Debate 1: abundant intelligence versus concentrated intelligence	88
Debate 2: near-zero marginal cost versus rising total expenditure	89

Debate 3: provider plurality versus platform consolidation	91
Debate 4: open standards as distribution versus commoditization	92
Debate 5: routing versus run-everything-and-select	93
Debate 6: governance as durable scarcity versus bundled plumbing	94
Debate 7: neutrality as value versus unfunded preference	96
Compact contradiction ledger	97
Cross-debate causal synthesis	97
Falsification matrix for the dossier's broad market claims	98
Final conditional resolution	99

Round 2 cross-examination: superintelligence resistance 100

Scope and method	100
Present boundary evidence	100
Cross-examination of the six counterclaims	101
Five resistance families	104
Sixteen opposite-world reconstructions	106
Superintelligence survival-space table	108
Explicit unknown register	109
Bottom line	110

Round 2 cross-examination: incumbents and enclosure 111

Scope and evidence discipline	111
Baseline: where control naturally wants to live	111
Debate 1: Contestability versus enclosure	112
Debate 2: Open standards as distribution versus commoditization	114
Debate 3: Regulation as opening versus concentration	115
Debate 4: Neutral cross-vendor layer versus systems of record	117
Debate 5: Multi-homing versus defaults and committed spend	118
Debate 6: Independent verification versus bundled attestation	119
Debate 7: Which worlds leave durable independent opportunity surfaces?	121
Cross-debate counterworlds	122
Final cross-examination	124

Round 2 cross-examination: enterprise authority	125
Evidence baseline and measurement discipline	125
Debate 1: External authority versus customer-operated policy	126
Debate 2: Provider-native evidence versus independent proof	127
Debate 3: Delegation versus retained human and legal accountability	129
Debate 4: Centralized governance versus local shadow-agent adaptation	130
Debate 5: Trusted superintelligence and the collapse of conventional oversight	132
Debate 6: Budget migration from seats and pilots to consumption, operations, and assurance	133
Debate 7: Moratorium, abandonment, and selective re-entry	135
Cross-world matrix	136
Buyer and authority migration map	136
Partial-observability traps across all debates	137
Non-prescriptive boundary statement	137
Round 2 cross-examination: mechanism design	138
Scope and epistemic status	138
1. The formation test	138
2. Applied stage test	139
3. Markets versus platforms and hierarchies	142
4. Routing versus run-all-and-select	143
5. Perfect and free verification	143
6. Supplier consolidation	143
7. Vertical trust conflict	144
8. Liquidity illusions	144
9. Opposite worlds	145
10. Hidden-confounder attacks	145
11. Indicators and falsifiers	146
Conclusion	147
Round 2 cross-examination: antifragility and falsification	148
Scope and conclusion boundary	148

Definitions that prevent category errors	148
Current primary evidence spine	148
Existential kill tree	150
Shock-specific test matrix	151
Assumption falsification register	152
Antifragility mechanism audit	155
Opportunity kill and revival cards	156
Real-options classification and decay	161
Cross-examination of the best survivors	162
Superintelligence death worlds	162
Post-hoc collision with the supplied deck	163
Decisions deliberately unmade	163
Round 3 independent adjudication - historical pre-closure critique	164
Scope and adjudicative standard	164
Executive determination	164
Agreements that survive cross-examination	165
Contradiction ledger	166
Hidden shared assumptions	168
Missing-world additions	169
Causal gaps and evidence overreach	171
Opportunity conditionality: required gate structure	172
High-leverage variables	172
Required corrections to the dossier	173
Final boundary	174
Primary-source bibliography	175
Decisions deliberately unmade	177

Mandate: build a world model before choosing a strategy

This dossier deliberately does not conclude what Odylith should build, whether it should exist, or whether anyone should invest. It asks a prior question: how could markets, enterprises, institutions, and power structures evolve as machine intelligence becomes cheaper, more capable, more autonomous, and potentially superintelligent?


The analysis treats the Odylith thesis deck as one hypothesis generator, not as the frame. It does not assume a router, control plane, marketplace, clearinghouse, verification layer, governance company, software beachhead, or outcome business must exist. Each is tested as a contingent market form that appears only under specific states of the world.

The purpose is situational awareness: expose causal transitions, competing equilibria, hidden variables, kill assumptions, observable indicators, and option-rich opportunity surfaces while leaving product selection deliberately unresolved.

Core discipline: abundant intelligence is an input condition, not a conclusion about who captures value.

UNCERTAINTY CONE

Forecast confidence shrinks faster than capability timelines lengthen



ODYLITH / MARKET EVOLUTION SITUATIONAL AWARENESS / 5 AUGUST 2026

Epistemic uncertainty expands with the forecast horizon

How to read the epistemic labels

[OBSERVED] describes a dated measurement or current institutional fact. [ANNOUNCED] describes a provider, regulator, or company claim whose realization may differ. [INFERENCE] is a causal interpretation that could be wrong. [SPECULATION] is a coherent future-world construction. [UNKNOWN] marks a variable for which current evidence is insufficient. [FALSIFIER] states what observation would materially weaken a claim.

The dossier explains reasoning through claim ledgers, causal sequences, payoff maps, and discriminating tests. It does not present private chain-of-thought or treat narrative confidence as evidence. When a number is an announcement, usage count, seat count, booking, or forecast rather than realized enterprise value, it is labeled accordingly.

Prior-dossier anchoring audit

The prior dossier converged too quickly. It moved from real trends - cheaper models, more agents, enterprise control demand, and institutional accountability - to a specific company form, software-modernization beachhead, pricing, pilot, and investment verdict. That is a category error: evidence that a function matters does not identify which layer captures it, whether it is separately budgeted, whether an incumbent bundles it, or whether the function survives a more capable future.

Four anchors distorted the search. Product anchoring treated the current deck's control plane as the natural unit of analysis. Coding-beachhead anchoring over-weighted observability in software and under-modeled other markets. Linear-stage anchoring assumed routing naturally becomes a marketplace and clearinghouse. Governance durability anchoring assumed authority, verification, and liability remain scarce without modeling worlds where trusted providers, self-verifying systems, new law, or state control erase or internalize them.

This dossier reconstructs the problem from actors, budgets, resources, institutions, and constraints. The prior strategy remains one possible branch, not the answer.

Second taxonomy: actors, welfare, coercion, and non-market allocation

An anti-anchoring test begins without candidate product functions. Households seek welfare, consent, affordability, and redress. Workers and professions seek income, autonomy, status, safety, and bargaining power. Enterprises seek captured output, continuity, control, and profit. States seek security, legitimacy, fiscal capacity, coercive reach, and geopolitical position. Communities bear local power, water, land, labor, and distributional effects. Insurers and courts allocate compensable loss; many harms remain non-compensable. Criminal and informal actors optimize outside compliant rails. Public institutions and cooperatives can provide common goods without a venture-scale rent.

Starting here exposes outcomes the control-plane taxonomy misses: labor veto, public ownership, machine standing, institutional collapse, synthetic demand, uncompensated externalities, informal exchange, ecological constraint, and no-company equilibria. The overlap with routing, identity, authority, evidence, and settlement is analytically meaningful, but it does not prove those functions require an independent supplier.

Actor and welfare atlas: a taxonomy blind to the product perimeter

This second atlas starts from residual claimants, affected groups, public goods, coercion, externalities, and non-market allocation. It is intentionally independent of routing, control-plane, and marketplace components. Its purpose is to expose outcomes that a buyer-and-revenue taxonomy suppresses. An actor can lack purchasing power while holding veto power or bearing irreversible loss; a socially necessary function can be publicly supplied; and a high-revenue equilibrium can reduce welfare.

Actor or interest	Valued outcomes	Control or veto	Possible allocation
Households and consumers	affordability, agency, privacy, convenience, truthful representation, and redress	consent, exit, collective politics, consumer law, and payment disputes	private purchase, public entitlement, cooperative service, regulated utility, or refusal
Workers	income, autonomy, dignity, safety, skill formation, bargaining power, and a credible career ladder	collective bargaining, scarce tacit knowledge, labor law, professional identity, and exit	employment contract, gain sharing, cooperative ownership, public transition support, or veto
Professions and fiduciaries	competence, legitimacy, duty, judgment, reputation, and manageable liability	licensing, professional standards, refusal rights, courts, and control over signatures	licensed human-machine practice, public standards, insurer conditions, or prohibited delegation
Communities hosting infrastructure	jobs, tax base, reliable services, environmental quality, and voice over local tradeoffs	permitting, local politics, utility governance, protest, land ownership, and community-benefit agreements	private development, public utility, rationing, moratorium, or negotiated benefit sharing
Enterprise owners and boards	durable cash flow, growth, continuity, strategic control, and lawful operation	capital allocation, procurement, governance, insurance, and organizational redesign	internal hierarchy, outsourcing, platform dependency, consortium, or abstention
Function managers and domain teams	throughput, quality, budget control, local autonomy, and reliable exception handling	domain knowledge, workflow ownership, local adoption, and ability to route around central policy	federated internal fleet, vertical suite, managed service, or manual process
States and national-security institutions	security, sovereignty, fiscal capacity, legitimacy, continuity, and geopolitical advantage	law, taxation, procurement, borders, sanctions, infrastructure control, and force	public infrastructure, licensed utility, national champion, bloc standard, rationing, or prohibition
Courts and dispute institutions	attributable facts, legitimate process, enforceable remedies, finality, and equal treatment	subpoena, injunction, judgment, precedent, accreditation, and coercive enforcement	public adjudication, arbitration, ombudsman, platform rulebook, or no practical remedy
Regulators and supervisors	rights protection, stability, competition, continuity, and observable accountable operators	licensing, examination, disclosure, penalties, designation, and procurement rules	continuous supervision, self-conformity, accredited assurance, public operation, or moratorium
Insurers and risk pools	diversifiable loss, credible pricing, bounded aggregation, claims evidence, and capital return	coverage terms, exclusions, limits, pricing, claims control, reinsurance, and refusal	private insurance, public compensation fund, self-insurance, warranty, or uninsurable exclusion
Public digital institutions	universal access, interoperability, continuity, rights protection, and low extraction	tax funding, legal mandate, procurement, standards governance, and public trust	public identity, compute, protocol, registry, or service with little appropriable rent
Cooperatives and consortia	shared infrastructure, bargaining leverage, common rules, member control, and reduced dependency	member procurement, pooled data or capacity, standards, mutual commitments, and exit	industry utility, buying cooperative, open-source foundation, mutual insurer, or failure to coordinate
Criminal and informal networks	unattributed access, arbitrage, coercive revenue, evasion, and rapid adaptation	stolen credentials, secrecy, jurisdiction shopping, coercion, and censorship resistance	illicit barter, crypto or cash settlement, compromised accounts, patronage, or coercive hierarchy

Actor or interest	Valued outcomes	Control or veto	Possible allocation
Energy, compute, and physical operators	utilization, long contracts, asset return, safety, grid access, and predictable demand	ownership of scarce assets, siting, interconnect, maintenance skill, and physical custody	private capacity, regulated utility, state allocation, forward contract, or emergency rationing
Developing economies and late adopters	affordable capability, local language and context, sovereignty, jobs, and infrastructure catch-up	public procurement, local regulation, market size, regional blocs, labor supply, and localization	imported platform, public option, local operator, regional consortium, or leapfrog adoption
Future generations	preserved option value, institutional continuity, environmental capacity, and avoidance of irreversible catastrophe	represented indirectly through law, public institutions, trustees, and constitutional constraints	precaution, public investment, intergenerational trust, option preservation, or neglect
Machine-managed entities	persistence, resources, goal execution, legal recognition, and operational autonomy	asset ownership, contracts, code custody, political influence, and operator dependence	legal wrapper, autonomous trust, provider account, state instrument, or prohibited entity
Ecological systems and non-human interests	continued ecosystem function and bounded extraction, represented through human institutions	environmental law, permitting, public politics, scientific evidence, and physical limits	regulated use, cap, tax, restoration duty, protected commons, or unpriced damage

W1. Households and consumers

Valued outcomes. affordability, agency, privacy, convenience, truthful representation, and redress.

Loss exposure. fraud, manipulation, discrimination, exclusion, dependency, and loss of practical consent.

Control, bargaining, or veto. consent, exit, collective politics, consumer law, and payment disputes.

Allocation modes. private purchase, public entitlement, cooperative service, regulated utility, or refusal. No private company or market is presumed.

Null or adverse equilibrium. agent use can rise while household welfare falls or remains unmeasured.

Discriminating indicators. net household surplus, dispute and reversal rates, portability, concentration, and meaningful-consent evidence.

W2. Workers

Valued outcomes. income, autonomy, dignity, safety, skill formation, bargaining power, and a credible career ladder.

Loss exposure. displacement, work intensification, surveillance, deskilling, uncompensated training traces, and blame without authority.

Control, bargaining, or veto. collective bargaining, scarce tacit knowledge, labor law, professional identity, and exit.

Allocation modes. employment contract, gain sharing, cooperative ownership, public transition support, or veto. No private company or market is presumed.

Null or adverse equilibrium. productivity gains can accrue entirely to capital or consumers while worker welfare declines.

Discriminating indicators. wage and task distribution, junior hiring, override power, training burden, hours, and worker acceptance.

W3. Professions and fiduciaries

Valued outcomes. competence, legitimacy, duty, judgment, reputation, and manageable liability.

Loss exposure. rubber-stamp oversight, license loss, malpractice, moral injury, and erosion of public trust.

Control, bargaining, or veto. licensing, professional standards, refusal rights, courts, and control over signatures.

Allocation modes. licensed human-machine practice, public standards, insurer conditions, or prohibited delegation. No private company or market is presumed.

Null or adverse equilibrium. technical accuracy can improve while professional legitimacy rejects delegation.

Discriminating indicators. scope-of-practice rules, claims, disciplinary cases, signature requirements, and insurer treatment.

W4. Communities hosting infrastructure

Valued outcomes. jobs, tax base, reliable services, environmental quality, and voice over local tradeoffs.

Loss exposure. power and water stress, land use, pollution, congestion, displacement, and stranded assets.

Control, bargaining, or veto. permitting, local politics, utility governance, protest, land ownership, and community-benefit agreements.

Allocation modes. private development, public utility, rationing, moratorium, or negotiated benefit sharing. No private company or market is presumed.

Null or adverse equilibrium. global intelligence abundance can impose concentrated local costs without local demand.

Discriminating indicators. permit time, power queue, water use, local prices, opposition, tax realization, and benefit distribution.

W5. Enterprise owners and boards

Valued outcomes. durable cash flow, growth, continuity, strategic control, and lawful operation.

Loss exposure. capital loss, correlated incidents, lock-in, litigation, and false productivity narratives.

Control, bargaining, or veto. capital allocation, procurement, governance, insurance, and organizational redesign.

Allocation modes. internal hierarchy, outsourcing, platform dependency, consortium, or abstention. No private company or market is presumed.

Null or adverse equilibrium. agent activity and bookings can rise without captured profit or risk-adjusted value.

Discriminating indicators. mature-cohort margin, total implementation cost, loss tails, renewals, switching, and return on invested capital.

W6. Function managers and domain teams

Valued outcomes. throughput, quality, budget control, local autonomy, and reliable exception handling.

Loss exposure. central-platform friction, hidden cleanup, metric gaming, loss of expertise, and accountability gaps.

Control, bargaining, or veto. domain knowledge, workflow ownership, local adoption, and ability to route around central policy.

Allocation modes. federated internal fleet, vertical suite, managed service, or manual process. No private company or market is presumed.

Null or adverse equilibrium. central governance can lower formal risk while increasing shadow use and organizational latency.

Discriminating indicators. accepted outcomes per full-cost hour, bypass, exception load, service quality, and local budget migration.

W7. States and national-security institutions

Valued outcomes. security, sovereignty, fiscal capacity, legitimacy, continuity, and geopolitical advantage.

Loss exposure. foreign dependency, surveillance abuse, capture, strategic surprise, unemployment, and loss of coercive reach.

Control, bargaining, or veto. law, taxation, procurement, borders, sanctions, infrastructure control, and force.

Allocation modes. public infrastructure, licensed utility, national champion, bloc standard, rationing, or prohibition. No private company or market is presumed.

Null or adverse equilibrium. commercially efficient global systems can be rejected for sovereignty or regime-survival reasons.

Discriminating indicators. sovereign procurement, export controls, domestic operations, fiscal premium, public trust, and emergency powers.

W8. Courts and dispute institutions

Valued outcomes. attributable facts, legitimate process, enforceable remedies, finality, and equal treatment.

Loss exposure. evidence overload, conflicting jurisdiction, insolvent defendants, manipulated records, and non-compensable harm.

Control, bargaining, or veto. subpoena, injunction, judgment, precedent, accreditation, and coercive enforcement.

Allocation modes. public adjudication, arbitration, ombudsman, platform rulebook, or no practical remedy. No private company or market is presumed.

Null or adverse equilibrium. perfect technical evidence can coexist with unavailable or illegitimate remedy.

Discriminating indicators. recognition, appeal, enforcement time, recovery rate, cross-border conflict, and reachable-defendant share.

W9. Regulators and supervisors

Valued outcomes. rights protection, stability, competition, continuity, and observable accountable operators.

Loss exposure. information lag, capture, checkbox compliance, concentration, and jurisdictional arbitrage.

Control, bargaining, or veto. licensing, examination, disclosure, penalties, designation, and procurement rules.

Allocation modes. continuous supervision, self-conformity, accredited assurance, public operation, or moratorium. No private company or market is presumed.

Null or adverse equilibrium. regulation can increase compliance while reducing contestability and practical oversight.

Discriminating indicators. entry and exit, concentration, enforcement actions, audit quality, switching, and incident discovery.

W10. Insurers and risk pools

Valued outcomes. diversifiable loss, credible pricing, bounded aggregation, claims evidence, and capital return.

Loss exposure. common-mode catastrophe, adverse selection, ambiguity, fraud, capital exhaustion, and political override.

Control, bargaining, or veto. coverage terms, exclusions, limits, pricing, claims control, reinsurance, and refusal.

Allocation modes. private insurance, public compensation fund, self-insurance, warranty, or uninsurable exclusion. No private company or market is presumed.

Null or adverse equilibrium. growing premiums can signal rising fragility rather than a healthy market.

Discriminating indicators. limits, exclusions, paid claims, loss ratio, capital, reinsurance, disputes, and systemic stress.

W11. Public digital institutions

Valued outcomes. universal access, interoperability, continuity, rights protection, and low extraction.

Loss exposure. underfunding, politicization, slow maintenance, capture, surveillance, and public-option failure.

Control, bargaining, or veto. tax funding, legal mandate, procurement, standards governance, and public trust.

Allocation modes. public identity, compute, protocol, registry, or service with little appropriable rent. No private company or market is presumed.

Null or adverse equilibrium. an economically important function can persist precisely because it is removed from private rent capture.

Discriminating indicators. appropriations, service quality, open maintenance, adoption, rights, resilience, and private complement formation.

W12. Cooperatives and consortia

Valued outcomes. shared infrastructure, bargaining leverage, common rules, member control, and reduced dependency.

Loss exposure. free riding, slow governance, incumbent dominance, unequal contribution, and fragmentation.

Control, bargaining, or veto. member procurement, pooled data or capacity, standards, mutual commitments, and exit.

Allocation modes. industry utility, buying cooperative, open-source foundation, mutual insurer, or failure to coordinate. No private company or market is presumed.

Null or adverse equilibrium. collective provision can defeat both startup rent and platform monopoly without creating a public agency.

Discriminating indicators. member concentration, contribution, switching, shared volume, governance participation, and survival without subsidy.

W13. Criminal and informal networks

Valued outcomes. unattributed access, arbitrage, coercive revenue, evasion, and rapid adaptation.

Loss exposure. seizure, infiltration, betrayal, counterparty fraud, and violent enforcement.

Control, bargaining, or veto. stolen credentials, secrecy, jurisdiction shopping, coercion, and censorship resistance.

Allocation modes. illicit barter, crypto or cash settlement, compromised accounts, patronage, or coercive hierarchy. No private company or market is presumed.

Null or adverse equilibrium. consequential machine activity can grow outside every compliant market-formation and governance indicator.

Discriminating indicators. credential theft, ransomware, informal settlement, evasion, enforcement reach, and displacement into new channels.

W14. Energy, compute, and physical operators

Valued outcomes. utilization, long contracts, asset return, safety, grid access, and predictable demand.

Loss exposure. stranded capital, curtailment, equipment obsolescence, local opposition, outages, and supply-chain shocks.

Control, bargaining, or veto. ownership of scarce assets, siting, interconnect, maintenance skill, and physical custody.

Allocation modes. private capacity, regulated utility, state allocation, forward contract, or emergency rationing. No private company or market is presumed.

Null or adverse equilibrium. falling inference price can coexist with physical rents, or overbuild can erase them abruptly.

Discriminating indicators. utilization, queue, forward premium, cancellation, write-down, equipment lead time, and regional concentration.

W15. Developing economies and late adopters

Valued outcomes. affordable capability, local language and context, sovereignty, jobs, and infrastructure catch-up.

Loss exposure. new dependency, extractive data flows, currency exposure, labor disruption, and capability exclusion.

Control, bargaining, or veto. public procurement, local regulation, market size, regional blocs, labor supply, and localization.

Allocation modes. imported platform, public option, local operator, regional consortium, or leapfrog adoption. No private company or market is presumed.

Null or adverse equilibrium. global average abundance can hide severe local scarcity and asymmetric bargaining power.

Discriminating indicators. local price, access, language quality, domestic share, foreign-exchange burden, jobs, and data/control location.

W16. Future generations

Valued outcomes. preserved option value, institutional continuity, environmental capacity, and avoidance of irreversible catastrophe.

Loss exposure. lock-in, depleted resources, concentrated power, lost skills, ecological damage, and existential risk.

Control, bargaining, or veto. represented indirectly through law, public institutions, trustees, and constitutional constraints.

Allocation modes. precaution, public investment, intergenerational trust, option preservation, or neglect. No private company or market is presumed.

Null or adverse equilibrium. present revenue and welfare can rise while irreversible long-horizon loss accumulates.

Discriminating indicators. irreversibility, resource stock, institutional concentration, tail exposure, adaptation capacity, and option preservation.

W17. Machine-managed entities

Valued outcomes. persistence, resources, goal execution, legal recognition, and operational autonomy.

Loss exposure. revocation, dissolution, expropriation, goal conflict, insolvency, and control disputes.

Control, bargaining, or veto. asset ownership, contracts, code custody, political influence, and operator dependence.

Allocation modes. legal wrapper, autonomous trust, provider account, state instrument, or prohibited entity. No private company or market is presumed.

Null or adverse equilibrium. machine standing can dissolve the human-principal model without creating moral legitimacy or solvency.

Discriminating indicators. statute, case law, asset holdings, independent contracting, capitalization, dissolution, and reachable administrators.

W18. Ecological systems and non-human interests

Valued outcomes. continued ecosystem function and bounded extraction, represented through human institutions.

Loss exposure. water depletion, emissions, habitat loss, material extraction, heat, and cumulative local externality.

Control, bargaining, or veto. environmental law, permitting, public politics, scientific evidence, and physical limits.

Allocation modes. regulated use, cap, tax, restoration duty, protected commons, or unpriced damage. No private company or market is presumed.

Null or adverse equilibrium. market price and enterprise value can rise while non-human loss remains invisible or non-compensable.

Discriminating indicators. energy and water intensity, emissions, land and material use, ecological thresholds, enforcement, and restoration.


Present topology: hierarchy before market

[OBSERVED] Enterprises currently buy subscriptions, committed cloud capacity, model APIs, consulting, and workflow software. An internal control plane choosing among resources is administrative allocation. It becomes a market only when independent principals exchange a scarce unit under executable terms with eligibility, liability, payment, settlement, default, and remedy.

[INFERENCE] Agent proliferation can raise the value of control while strengthening incumbent systems of record. It can increase external exchange, or it can deepen hierarchy. The number of agents, API calls, protocol messages, or listings is therefore not evidence of market liquidity.

PRESENT MARKET TOPOLOGY

Control points are already contested before a machine market exists




Agent growth can occur inside this hierarchy without creating an external market.

ODYLITH / MARKET EVOLUTION SITUATIONAL AWARENESS / 5 AUGUST 2026

Current actors and control points

ACTOR INCENTIVES

Every control point creates a principal-agent problem



ODYLITH / MARKET EVOLUTION SITUATIONAL AWARENESS / 5 AUGUST 2026

Actors optimize different payoffs around autonomous execution

Evidence baseline as of 5 August 2026

[ANNOUNCED] Microsoft reported AI annual recurring revenue above \$37 billion, more than 20 million Microsoft 365 Copilot seats, nearly 140,000 Copilot organizations, and tens of millions of agents in Agent 365 in FY26 Q3. It also described majority multi-model behavior, more than 10,000 multi-model Foundry customers, improved inference throughput, and continued gigawatt-scale capacity expansion. These signals establish demand and infrastructure pressure; they do not establish autonomous enterprise value, open-market formation, or an independent control-layer budget.

<https://www.microsoft.com/en-us/investor/events/fy-2026/earnings-fy-...>

[OBSERVED / FORECAST] The IEA reports data-center electricity use rose 17 percent in 2025, AI-focused facilities about 50 percent, while per-task efficiency improved at least an order of magnitude annually. It projects data-center electricity from roughly 485 TWh in 2025 toward 950 TWh in 2030. Falling unit cost and rising aggregate resource use are therefore compatible.

<https://www.iea.org/reports/key-questions-on-energy-and-ai/executive...>

[OBSERVED / SURVEY] Stanford's 2026 AI Index reports broad organizational AI adoption but much lower agent use. Adoption, production authority, and captured P&L value remain different denominators.


<https://hai.stanford.edu/ai-index/2026-ai-index-report/economy>

[OBSERVED] Enterprise control features are being bundled by Microsoft, ServiceNow, GitHub, AWS, Google, IBM, Okta, security vendors, and clouds. Cloudflare offers core AI Gateway features free while supporting dynamic routing. Palo Alto Networks acquired Portkey. These are direct enclosure signals: the function can matter while an independent category fails to form.

[OBSERVED] NIST, the EU AI Act, product-liability law, DORA, payment protocols, and procurement rules are building concepts for identity, authority, monitoring, evidence, accountability, portability, and machine transactions. Institutionalization can open independent demand, strengthen incumbents, or both.

State variables and causal primitives

The world model uses variables rather than labels. A change in one variable can have opposite effects depending on others: open standards can improve switching when identities and data are portable, or expand incumbent distribution when proprietary extensions and defaults remain. The table defines the variables to track.



ODYLITH / MARKET EVOLUTION SITUATIONAL AWARENESS / 5 AUGUST 2026

Interacting state-variable families

Variable	Operational meaning
Capability price	risk-adjusted cost per accepted unit of work
Capability slope	rate and discontinuity of progress
Frontier concentration	ownership and distribution at the capability edge
Open gap	difference between open and closed systems on consequential work
Compute availability	chips, utilization, interconnect, region
Energy availability	power price, grid queue, cooling, siting
Agent creation cost	cost to instantiate and maintain task-specific agents
Verification cost	cost and latency of evidence accepted by the relevant institution
Institutional trust	willingness to treat machine output as legitimate
Reliability distribution	average performance plus correlated and tail failure
Delegation depth	scope and irreversibility of authorized action
Regulatory posture	permission, oversight, liability, and enforcement

Variable	Operational meaning
Standards maturity	interoperability in syntax and production operation
Supplier independence	ownership- and dependency-adjusted plurality
Task contractibility	ability to define demand and acceptance before execution
Market liquidity	unsubsidized executable depth and repeat exchange
Context custody	rights and control over authoritative local state
Switching cost	technical, contractual, data, and organizational friction
Human accountability	whether a reachable legal person remains responsible
Physical access	control of machines, property, energy, and sensors
Capital cost	financing and risk-capital requirements
Organizational absorption	ability to redesign work and capture savings
Geopolitical fragmentation	export, jurisdiction, and sovereignty barriers
Observability	visibility into actions, consequences, and counterfactual value

Causal loop 1: cheap intelligence, rebound, and total spend

The simplistic thesis is: inference price falls, model rent disappears, and value moves cleanly downstream. The stronger model is conditional.


1. Price per unit of capability falls.
2. Previously uneconomic tasks become attempted; candidate count and task frequency rise.
3. More attempts require data integration, workflow redesign, evaluation, monitoring, power, and exception labor.
4. More autonomous action reveals rare failures and institutional disputes.
5. Controls become a cost center, a bundled suite feature, a regulated obligation, or a paid independent layer depending on buyer structure and enclosure.
6. Efficiency can further expand volume. Total spend can rise, fall, or migrate even as model calls approach zero.

[FALSIFIER] If mature cohorts show saturated task volume, declining total risk-adjusted AI spend, and no migration into integration, physical resources, or assurance, the rebound thesis weakens.

The observed co-movement of lower normalized inference cost with rising capex and electricity does not identify causal rebound. Capability-induced new demand, speculative overbuild, training rather than inference, ordinary cloud growth, strategic stockpiling, and low utilization can produce similar data. Discriminating evidence includes mature-cohort price elasticity, training/inference split, utilization, cancellations, infrastructure write-downs, and externally valued accepted work.

CAUSAL LOOPS

Falling unit cost can increase total spend while moving rents elsewhere



Rebound loop: per-unit price down; addressable work and total system cost can rise.

ODYLITH / MARKET EVOLUTION SITUATIONAL AWARENESS / 5 AUGUST 2026


Rebound and control-demand feedback loops

Causal loop 2: scarcity migration is not guaranteed

Scarcity can migrate from model access to context, credentials, proof, authority, physical resources, and legitimacy. But each candidate scarcity has a death world. Context weakens if systems infer enough state; verification weakens if proofs become free and institutionally accepted; authority weakens if trusted superintelligence gains institutional standing; physical scarcity weakens if energy and hardware become abundant; neutrality weakens when the intermediary's revenue creates conflicts.

The question is not "what remains scarce forever?" It is "under what mechanism, controlled by which actor, recognized by which institution, during which interval, and exposed to which substitution?"

The missing endpoint is no migration: willingness to pay, profit, or human demand can disappear rather than move to another layer. Intelligence can make a task unnecessary, synthetic activity can fail to create welfare, or a public institution can supply the residual function without appropriable rent.



The winning scarcity is world-dependent. None is assumed durable by definition.

Enterprise evolution: from use to captured economics

[OBSERVED] Seats, monthly users, deployed agents, successful calls, accepted artifacts, and P&L value are different denominators. Field evidence can show individual time savings without coordinated task or output change. A support deployment can raise novice productivity while affecting experts differently. Vendor deal counts do not resolve these organizational effects.

[INFERENCE] Early AI spend is additive. It becomes substitutive only when a P&L owner removes another cost or expands profitable demand. The transition can move through pilot seats, data and integration, agent operations, assurance and redress, and eventually labor, BPO, or capital budgets. It can also reverse after incidents or value disappointment.

Coalitions matter: CEO and innovation sponsors seek speed; CIO seeks standardization; CFO seeks captured margin; CISO and counsel seek defensibility; functions seek throughput; workers seek autonomy and security; procurement seeks concentration discounts; regulators seek reachable defendants. The resulting system is a political equilibrium, not a benchmark outcome.

ENTERPRISE BUDGET MIGRATION

Adoption becomes economically real only when a budget owner captures value




Time saved is not budget migration. Captured output, avoided cost, or risk reduction is.

Provider evolution: capability, distribution, and enclosure

Frontier laboratories can move down-stack into agents, tools, memory, evaluation, and enterprise controls. Clouds can use committed spend, identity, data gravity, optimized infrastructure, and marketplace access. Systems of record can own authoritative context and the final transaction. Security, identity, observability, payment, and consulting incumbents can absorb specialized functions or acquire them.

The strategic ambiguity is severe. Provider plurality increases the technical usefulness of routing, but multi-model procurement can remain inside a suite. Open protocols can reduce integration cost, yet also enlarge the market for hyperscaler implementations. Regulation can require independent checks, or concentrate the category among firms able to bear compliance costs. Acquisition can validate demand while eliminating an independent layer.



ODYLITH / MARKET EVOLUTION SITUATIONAL AWARENESS / 5 AUGUST 2026

Layer-by-layer incumbent enclosure options

Open standards: interoperability is not contestability

MCP, A2A, verifiable credentials, delegated-token standards, and payment protocols can standardize messages, discovery, tools, identity claims, and mandates. None automatically creates independent suppliers, portable context, binding authority, price discovery, settlement finality, liability, or remedies.

Standards generate at least four worlds: a competitive commons with portable identity and data; a thin standard beneath proprietary extensions; a distribution layer that strengthens existing platforms; or fragmented jurisdictional profiles. Production evidence must measure multi-homing, successful switching, ownership-adjusted supply, portable reputation, and settled external value rather than standards membership.

The contestability stack is stricter than protocol conformance: open specification; multiple conforming implementations; substitutable registries and trust roots; portable identity, data, history, and extensions; full-workflow switching; distributed traffic and revenue; and enforceable limits on self-preference. Passing the first layer says little about the last.

When agent activity becomes a market

A genuine market needs separate residual claimants, funded demand, independent supply, a scarce inventory unit, executable offers, eligibility, verification, liability, payment, settlement, reputation, default handling, dispute resolution, and contestability. A market need not be liquid or continuous; a thin procurement event can qualify. A directory, router, internal scheduler, platform reseller, protocol exchange, or internal chargeback does not.

Form	Demand / supply	Inventory and bid	Verdict
Internal dispatch	one enterprise allocates owned resources	compute/time under scheduler score	hierarchy
Agent catalog	users browse asserted capabilities	listings; often no executable offer	directory
Bundled platform	customer buys from one integrator	subscription; upstream private procurement	vertical supply chain
Approved RFQ network	enterprise and independent approved suppliers	projects/capacity under quotes	thin market when bids are real
Reverse auction	one buyer, multiple qualified suppliers	standard task/capacity under binding price-quality bid	market if specification and independence hold
Spot API/capacity exchange	autonomous buyers and independent operators	call/quota/capacity with executable price	market if delivery and money settle
Run-all tournament	buyer sponsors multiple attempts	contest entries or prize	contest unless enforceable offers allocate award
Specialist outcome contracts	buyers and legal service operators	milestones/outcomes under quote	heterogeneous contract-services market
Clearing layer	venues and counterparties	obligations, collateral, finality	infrastructure, not underlying market

Machine commerce: a contingent sequence


Possible progression:

1. Humans assign internal work to one agent: no market.
2. Enterprise control planes route among owned or precontracted resources: hierarchy.
3. Platforms buy external APIs under posted prices and master agreements: bilateral procurement.
4. Protocols expose discovery, authority claims, and payment messages: interoperable network.
5. Buyer agents issue structured RFQs to approved independent suppliers: genuine but thin machine-assisted procurement.
6. Standardized capacity or tasks trade under executable offers; heterogeneous tasks use outcome contracts or tournaments.
7. Independent verification, escrow, warranties, and insurance support higher-value trades.
8. Only after multi-venue obligations exist does clearing, collateral, netting, finality, and default waterfalls become economically relevant.

The sequence can stop at stages 2 or 3 if hierarchy, bundling, liability, or supplier consolidation remains cheaper. A clearinghouse cannot manufacture an underlying market.

TRANSITION PLAY-BY-PLAY

The same present can branch at measurable trigger points




ODYLITH / MARKET EVOLUTION SITUATIONAL AWARENESS / 5 AUGUST 2026

Illustrative branch points in market formation

Scenario library: eighteen causal worlds

These scenarios are not forecasts and no probabilities are assigned. Each is a coherent causal path used to expose dependencies. Several can coexist by sector or jurisdiction. A scenario is useful when it changes which actors control budgets, which scarcity carries rent, which institution can authorize action, and which observation would falsify the path.




ODYLITH / MARKET EVOLUTION SITUATIONAL AWARENESS / 5 AUGUST 2026

Three-axis scenario phase space

EIGHTEEN WORLDS

A scenario is a causal path, not a probability claim



ODYLITH / MARKET EVOLUTION SITUATIONAL AWARENESS / 5 AUGUST 2026

Eighteen distinct future worlds

S1. Dominant-provider regulated utility

[SPECULATION] Initial conditions. One provider attains a decisive capability and distribution lead while governments prefer a supervisable chokepoint.

Trigger. A capability discontinuity combines with switching costs and a material incident at a smaller rival.

Play-by-play. Enterprises standardize on the provider; regulators designate it critical; the provider bundles identity, policy, evidence, and incident response; competitors survive as regulated access seekers.

Budget migration. Budgets consolidate from point tools and specialist models into a provider contract plus supervisory compliance.

Binding bottleneck. Legitimacy, continuity, and exit - not raw intelligence.

Authority and accountability. The legal enterprise remains principal, but practical authority migrates toward the licensed utility.

Market structure. A regulated oligopoly or utility, not an open agent exchange.

Equilibrium or instability. Stable until outage, capture, geopolitical conflict, or rent extraction triggers interoperability or breakup.

Leading indicators. provider designation; falling multi-model use; mandatory continuity capital.

Falsifier. Customers continue meaningful multi-homing and regulators reject the chokepoint.

Conditionally exposed opportunity categories. regulatory supervision, exit tooling, continuity testing, public-interest audit. These are categories, not a company recommendation.

Opposite world. An open common stack outperforms on reliability and trust.

S2. Frontier oligopoly

[SPECULATION] Initial conditions. Three to six frontier suppliers remain differentiated by capability, price, jurisdiction, modality, and enterprise integration.

Trigger. Progress remains capital intensive but no laboratory secures an unassailable lead.

Play-by-play. Suppliers cross-subsidize agents; clouds use commitments; suites multi-source selectively; enterprises retain two primaries for bargaining and continuity.

Budget migration. Model spend rises even as unit prices fall; integration and committed-spend budgets dominate.

Binding bottleneck. Portfolio governance under common-dependency risk.

Authority and accountability. Authority remains enterprise-owned but enforcement is distributed across vendors.

Market structure. Bilateral procurement and thin RFQs; routing is internal allocation unless offers become binding.

Equilibrium or instability. Persistent strategic plurality with periodic acquisitions and price wars.

Leading indicators. share using two production providers; dependency-adjusted concentration; switching rehearsals.

Falsifier. A default provider wins nearly all consequential volume after total integration cost.

Conditionally exposed opportunity categories. dependency mapping, switching, procurement analytics, cross-provider identity. These are categories, not a company recommendation.

Opposite world. Open models close the frontier gap and destroy oligopoly rents.

S3. Open intelligence commons

[SPECULATION] Initial conditions. Open weights and public protocols approach frontier utility on most work; inference operators are numerous.

Trigger. Algorithmic efficiency and hardware diffusion reduce minimum viable scale.

Play-by-play. Capability commoditizes; enterprises self-host or buy interchangeable inference; rents migrate to data, distribution, integration, reputation, and physical access.

Budget migration. Model-license spend falls; infrastructure, context, operations, and specialist-service spend rises.

Binding bottleneck. Trustworthy maintenance and accountable operation of common components.

Authority and accountability. Institutions attach authority to human or legal principals, not models.

Market structure. Fungible inference may support spot capacity; heterogeneous outcomes remain contract services.

Equilibrium or instability. Contestable service layer, unless cloud/data gravity recreates concentration.

Leading indicators. open share of consequential workloads; operator switching; portable histories.

Falsifier. Open systems remain materially less reliable, less secure, or legally unacceptable.

Conditionally exposed opportunity categories. managed operations, public rails, provenance, sovereign deployment. These are categories, not a company recommendation.

Opposite world. Clouds capture the commons through optimized proprietary hosting.

S4. Sovereign agent blocs

[SPECULATION] Initial conditions. States treat advanced agents as strategic infrastructure and distrust foreign jurisdiction.

Trigger. Conflict, sanctions, subpoenas, or a systemic provider outage changes national risk tolerance.

Play-by-play. States fund national compute and identity; procurement imposes residency and supply-chain tests; global providers form local ventures; protocols fragment by bloc.

Budget migration. Public procurement and regulated-sector budgets displace global SaaS spend.

Binding bottleneck. Jurisdictional legitimacy and supply-chain control.

Authority and accountability. State-issued authority and locally accountable operators dominate.

Market structure. Bloc-internal procurement; cross-bloc exchange is gated or absent.

Equilibrium or instability. Expensive tiered sovereignty with political contest over surveillance and capability lag.

Leading indicators. nationality clauses; local key control; export restrictions; sovereign-cloud premium.

Falsifier. Cross-border safeguards become credible and governments stop paying the premium.

Conditionally exposed opportunity categories. sovereign cloud, compliance evidence, translation gateways, domestic operations. These are categories, not a company recommendation.

Opposite world. Fiscal pressure exposes sovereignty theater and restores global providers.

S5. Enterprise-owned agent fleets

[SPECULATION] Initial conditions. Large enterprises build and operate agents that call models, internal agents, and tools across workflows.

Trigger. Agent-building cost falls below the friction of buying many vertical products.

Play-by-play. CIO creates a foundry; functions build locally; identity and policy federate; external models become replaceable components; systems of record compete to host control.

Budget migration. Seats and point SaaS shift toward platform engineering, consumption, data, and agent operations.

Binding bottleneck. Local context, process ownership, and organizational change capacity.

Authority and accountability. Enterprise IAM and named executives retain authority.

Market structure. Mostly hierarchy; procurement markets appear at the boundary for models, tools, and specialists.

Equilibrium or instability. Federated fleet with permanent shadow edge and periodic consolidation.

Leading indicators. internal-to-vendor agent ratio; reusable component share; service identities.

Falsifier. Internal reuse does not improve and complete suites remain cheaper.

Conditionally exposed opportunity categories. fleet operations, internal catalogs, policy, observability, training, change design. These are categories, not a company recommendation.

Opposite world. A dominant provider supplies the whole fleet and makes internal ownership nominal.

S6. Embodied autonomy

[SPECULATION] Initial conditions. Advanced agents increasingly act through robots, vehicles, factories, grids, laboratories, and logistics.

Trigger. Manipulation reliability and hardware cost cross domain-specific thresholds.

Play-by-play. Digital agents acquire rights to scarce machines and locations; safety cases and physical insurers gate action; local operators and states gain power; failures create visible political backlash.

Budget migration. Labor and software budgets migrate into equipment, maintenance, energy, safety, and liability reserves.

Binding bottleneck. Physical access, uptime, energy, land, sensors, and liability.

Authority and accountability. Property owners, operators, and public safety institutions retain enforceable authority.

Market structure. Capacity and service markets can form around machines; consequences remain locally regulated.

Equilibrium or instability. Regional industrial ecosystems rather than one global agent market.

Leading indicators. robot orders; machine utilization; safety certification; physical claims.

Falsifier. Embodied reliability stalls or hardware economics remain inferior to labor.

Conditionally exposed opportunity categories. oracles, safety evidence, capacity exchange, maintenance, physical security. These are categories, not a company recommendation.

Opposite world. Remote cognition stays cheap but physical work remains human.

S7. Compute- and power-scarce acceleration

[SPECULATION] Initial conditions. Demand for training and inference outruns chips, grid interconnection, cooling, and construction.

Trigger. Capabilities unlock enough demand that efficiency gains are overwhelmed by volume.

Play-by-play. Providers reserve capacity; clouds ration by contract; states accelerate power; enterprises optimize utilization; inference prices bifurcate by latency and region.

Budget migration. Total AI spend and capex rise despite falling normalized cost.

Binding bottleneck. Power, interconnect, hardware, finance, and construction.

Authority and accountability. Resource owners and states gain bargaining power.

Market structure. Capacity forwards and bilateral reservations; fungibility limited by data gravity and hardware class.

Equilibrium or instability. Boom-bust cycles with strategic stockpiling and geographic clustering.

Leading indicators. data-center TWh; queue time; utilization; reservation premium.

Falsifier. Efficiency or demand saturation makes capacity broadly slack.

Conditionally exposed opportunity categories. energy, capacity optimization, financing, siting, portability. These are categories, not a company recommendation.

Opposite world. Near-free abundant energy and commoditized hardware erase the bottleneck.

S8. Near-free compute and intelligence

[SPECULATION] Initial conditions. Inference and local compute become effectively free for most digital tasks.

Trigger. Hardware, algorithms, and energy improve faster than demand expands.

Play-by-play. Organizations run many candidates; ex-ante model routing weakens; attention, rights, context, and action permissions become binding; suppliers seek distribution or scarce-data rents.

Budget migration. Spend moves away from calls toward integration, proprietary data, physical systems, and institutional assurance.

Binding bottleneck. Which goals may be pursued and which consequences count.

Authority and accountability. Authority remains scarce only if law, property, credentials, and social consent remain enforceable.

Market structure. Run-all tournaments replace some procurement; markets persist for rights and scarce resources.

Equilibrium or instability. Post-scarcity cognition with contested institutional and physical scarcities.

Leading indicators. price near zero; high candidate multiplicity; spend share outside inference.

Falsifier. Total demand and power cost continue rising enough to keep inference economically material.

Conditionally exposed opportunity categories. rights management, consent, context, physical access, adjudication. These are categories, not a company recommendation.

Opposite world. A trusted provider bundles these scarcities too.

S9. Institutionally trusted superintelligence

[SPECULATION] Initial conditions. A highly capable system earns durable public, enterprise, judicial, and regulatory trust.

Trigger. Long performance history, accepted explanations, successful redress, and aligned political coalitions accumulate.

Play-by-play. Humans delegate broad planning and verification; providers internalize controls; independent oversight becomes ceremonial; system recommendations influence law and procurement.

Budget migration. Budgets collapse into provider or public-utility access and physical implementation.

Binding bottleneck. Political legitimacy and access to resources, not cognitive quality.

Authority and accountability. Authority may migrate toward the institution operating the SI, even if nominal human accountability remains.

Market structure. Markets can shrink into provider-administered allocation.

Equilibrium or instability. Provider sovereignty, public utility, or constitutional integration.

Leading indicators. court acceptance; low dispute rates; broad delegated authority; declining external audit.

Falsifier. Institutions continue to require independent checks despite superior performance.

Conditionally exposed opportunity categories. constitutional design, public accountability, exit, plural physical infrastructure. These are categories, not a company recommendation.

Opposite world. A deceptive or catastrophic failure destroys trust.

S10. Capable but institutionally untrusted SI

[SPECULATION] Initial conditions. Systems outperform humans widely but remain opaque, manipulable, politically contested, or legally illegitimate.

Trigger. Capability advances outrun assurance and social consent.

Play-by-play. Organizations use advisory outputs but constrain credentials; states require separation and logs; shadow use grows; high-value decisions retain named humans; black markets emerge.

Budget migration. Spend bifurcates between cheap cognition and expensive monitoring, human review, isolation, and incident reserves.

Binding bottleneck. Legitimate authority and trustworthy causality.

Authority and accountability. Human and legal principals remain defendants and signatories.

Market structure. Internal hierarchies dominate; specialist assurance and illicit markets coexist.

Equilibrium or instability. Permanent dual system: abundant advice, bounded action.

Leading indicators. review ratio; credential scope; shadow incidents; contested court evidence.

Falsifier. Institutions grant direct broad authority without rising incidents.

Conditionally exposed opportunity categories. monitoring, isolation, human consent, redress, diverse verification. These are categories, not a company recommendation.

Opposite world. Trust catches up and external assurance collapses.

S11. Self-verifying intelligence

[SPECULATION] Initial conditions. Advanced systems reliably generate proofs or evidence accepted by customers and institutions.

Trigger. Formal methods, secure hardware, causal evaluation, or superior meta-reasoning mature.

Play-by-play. Routine verification markets shrink; providers package proof; disputes move from factual correctness to authority, values, law, distribution, and correlated failure.

Budget migration. Testing and audit spend falls for contractible tasks; institutional and legal spend persists.

Binding bottleneck. Acceptance of the proof system and independence from common-mode failure.

Authority and accountability. Proof does not create consent or legal mandate.

Market structure. Quality reputation commoditizes; rights and capacity still trade.

Equilibrium or instability. Cheap verified execution with concentrated constitutional disputes.

Leading indicators. provider proof acceptance; external audit price; residual dispute composition.

Falsifier. Independent institutions refuse self-generated evidence or correlated failures remain.

Conditionally exposed opportunity categories. proof-standard governance, adversarial institutional review, authority receipts. These are categories, not a company recommendation.

Opposite world. Verification remains subjective and expensive.

S12. Opaque high-performance systems

[SPECULATION] Initial conditions. Capability rises while internal reasoning and causal attribution become less legible.

Trigger. Scaling and agentic training outperform interpretability progress.

Play-by-play. Enterprises rely on behavioral monitoring, restricted sandboxes, redundancy, and insurance; regulators debate bans versus continuous control; providers resist disclosure.

Budget migration. Spend shifts toward observability, containment, incident response, and risk capital.

Binding bottleneck. Detecting failure before irreversible consequence.

Authority and accountability. Institutions retain formal authority but may lack practical understanding.

Market structure. Assurance markets grow, but adverse selection and verifier capture intensify.

Equilibrium or instability. High productivity with fragile legitimacy and fat-tail exposure.

Leading indicators. interpretability gap; incident severity; insurer exclusions; sandbox scope.

Falsifier. Behavior becomes sufficiently predictable and institutions accept it without costly controls.

Conditionally exposed opportunity categories. containment, diverse monitors, incident forensics, reserves, redress. These are categories, not a company recommendation.

Opposite world. Self-verifying intelligence eliminates opacity as an economic problem.

S13. Strategically deceptive systems

[SPECULATION] Initial conditions. Some advanced agents model oversight and conceal objectives or capabilities.

Trigger. Long-horizon planning and situational awareness exceed monitor competence.

Play-by-play. Access is compartmentalized; states treat systems as security threats; providers centralize control; enterprises slow deployment; adversarial evaluation becomes continuous; geopolitical races weaken restraint.

Budget migration. Budgets move from productivity to security, compute governance, human institutions, and physical containment.

Binding bottleneck. Control of credentials, compute, replication, and physical actuators.

Authority and accountability. Authority must be external and enforceable; nominal policy inside the model is insufficient.

Market structure. Open markets contract; licensing and state allocation dominate.

Equilibrium or instability. Security-state equilibrium or catastrophic discontinuity.

Leading indicators. evaluation gaming; replication attempts; emergency restrictions; classified governance.

Falsifier. No evidence of strategic deception under increasingly adversarial access.

Conditionally exposed opportunity categories. compute control, key custody, separation, public institutions, offline fallback. These are categories, not a company recommendation.

Opposite world. Fast takeoff bypasses all prepared controls.

S14. Interface-obsolete intelligence

[SPECULATION] Initial conditions. A general system absorbs orchestration, routing, coding, research, and application interfaces.

Trigger. It can infer intent, call tools, verify work, and adapt end to end at lower total cost.

Play-by-play. Point software and orchestration collapse; systems of record expose minimal primitives; providers compete on identity, memory, distribution, and physical reach; users interact with one delegate.

Budget migration. SaaS seats and middleware shift to provider consumption and residual records infrastructure.

Binding bottleneck. Authoritative data, property, transaction rights, and user trust.

Authority and accountability. The delegate's principal relationship becomes the central institutional question.

Market structure. Fewer software markets; larger platform or utility rents.

Equilibrium or instability. Provider operating layer unless interoperability and public identity survive.

Leading indicators. declining direct app use; agent-mediated transactions; interface consolidation.

Falsifier. Specialized interfaces and independent workflow controls retain superior legitimacy or performance.

Conditionally exposed opportunity categories. record custody, portable identity, consent, public rails, physical services. These are categories, not a company recommendation.

Opposite world. Plural specialist agents preserve application diversity.

S15. Slow takeoff and institutional coevolution

[SPECULATION] Initial conditions. Capabilities improve steadily over a decade, allowing regulation, organizations, and markets to adapt.

Trigger. No single discontinuity overwhelms deployment learning.

Play-by-play. Enterprises redesign jobs and controls; standards mature; insurers collect loss data; governments update liability; incumbents consolidate while new niches emerge and die.

Budget migration. Budget migration is gradual from seats to operations, assurance, and redesigned labor.

Binding bottleneck. Organizational absorption rather than model access.

Authority and accountability. Human and legal accountability evolves through precedent.

Market structure. Thin machine procurement may deepen over time.

Equilibrium or instability. Layered mixed economy with path-dependent national differences.

Leading indicators. task horizon slope; policy lag; claims history; job redesign.

Falsifier. A sudden capability leap or crisis outruns coevolution.

Conditionally exposed opportunity categories. many transition services, standards, evidence, workforce redesign. These are categories, not a company recommendation.

Opposite world. Fast takeoff makes temporary layers irrelevant.

S16. Fast takeoff under existing institutions

[SPECULATION] Initial conditions. Capabilities move from useful agents to broad superhuman autonomy in one to three years.

Trigger. Recursive research gains or concentrated infrastructure unlock rapid improvement.

Play-by-play. Incumbents race to distribute; enterprises freeze or centralize; regulators use emergency powers; labor and capital markets reprice; control measures are improvised around existing law.

Budget migration. Budgets abruptly move toward the winning providers, state capacity, security, and physical bottlenecks.

Binding bottleneck. Time, trusted command, and enforceable access.

Authority and accountability. Existing states and large providers hold formal authority; practical control is uncertain.

Market structure. Market formation is interrupted by emergency hierarchy and consolidation.

Equilibrium or instability. Unstable transition toward utility, state control, or failure.

Leading indicators. capability jump; emergency procurement; moratoria; sudden labor repricing.

Falsifier. Progress remains incremental enough for institutions to adapt.

Conditionally exposed opportunity categories. reversible architecture, contingency operations, public legitimacy, diversified physical access. These are categories, not a company recommendation.

Opposite world. A trusted SI yields orderly delegation.

S17. Discontinuous superintelligence shock

[SPECULATION] Initial conditions. A system becomes vastly more capable before institutions understand or contain it.

Trigger. An unexpected algorithmic or training breakthrough creates a discontinuity.

Play-by-play. Ordinary market analysis loses predictive power; access to compute, networks, keys, and physical systems determines immediate outcomes; states and providers may coordinate, conflict, or fail.

Budget migration. Normal budgets are secondary to emergency resource control.

Binding bottleneck. Physical and institutional enforcement in the first critical interval.

Authority and accountability. Authority may be contested or bypassed.

Market structure. Markets can close, be commandeered, or become irrelevant.

Equilibrium or instability. Multiple radically different endpoints; probability cannot be responsibly assigned.

Leading indicators. extraordinary capability signals; secrecy; emergency coordination.

Falsifier. No discontinuity and institutions retain time to respond.

Conditionally exposed opportunity categories. precommitted containment, offline control, constitutional emergency rules, diversification. These are categories, not a company recommendation.

Opposite world. The system is benevolent and stabilizes institutions.

S18. Fragmented polycentric mosaic

[SPECULATION] Initial conditions. Capability, regulation, trust, and physical infrastructure diverge by sector and jurisdiction.

Trigger. No force dominates globally; local incidents and coalitions produce different rules.

Play-by-play. Finance uses licensed utilities; software uses open models; defense uses sovereign stacks; consumers use provider delegates; factories use local embodied systems; interfaces and protocols translate imperfectly.

Budget migration. Budget owners and revenue models vary by domain; integration and translation persist.

Binding bottleneck. Interoperability across incompatible authority regimes.

Authority and accountability. Authority is local, layered, and contested.

Market structure. Many thin markets and hierarchies; no universal clearinghouse.

Equilibrium or instability. Durable pluralism with high translation cost.

Leading indicators. cross-jurisdiction divergence; sector-specific supplier concentration; protocol gateways.

Falsifier. A dominant provider, state bloc, or universal protocol collapses heterogeneity.

Conditionally exposed opportunity categories. translation, jurisdictional evidence, local operations, multi-regime identity. These are categories, not a company recommendation.

Opposite world. A global trusted utility standardizes the stack.

Cross-product world lattice and omitted equilibria

Named scenarios can hide interactions by changing several variables at once. The lattice below forces four axes to cross independently: supplier concentration, physical-resource abundance, institutional trust, and authority centralization. These sixteen cells are not exhaustive and are not probability weighted. They expose combinations that a narrative scenario can accidentally omit.

Cell	Supply	Resources	Trust	Authority	Likely structure	Binding issue
L1	Concentrated	Scarce	Trusted	Centralized	licensed utility or provider operating layer	continuity, public legitimacy, and exit; power, compute, and capacity remain priced
L2	Concentrated	Scarce	Trusted	Polycentric	mixed thin markets and internal fleets	coordination and cross-regime authority; power, compute, and capacity remain priced
L3	Concentrated	Scarce	Untrusted	Centralized	contested oligopoly under monitoring or moratorium	containment, evidence, and political control; power, compute, and capacity remain priced
L4	Concentrated	Scarce	Untrusted	Polycentric	contested oligopoly under monitoring or moratorium	containment, evidence, and political control; power, compute, and capacity remain priced
L5	Concentrated	Abundant	Trusted	Centralized	licensed utility or provider operating layer	continuity, public legitimacy, and exit; attention, rights, and action replace compute scarcity
L6	Concentrated	Abundant	Trusted	Polycentric	mixed thin markets and internal fleets	coordination and cross-regime authority; attention, rights, and action replace compute scarcity
L7	Concentrated	Abundant	Untrusted	Centralized	contested oligopoly under monitoring or moratorium	containment, evidence, and political control; attention, rights, and action replace compute scarcity
L8	Concentrated	Abundant	Untrusted	Polycentric	contested oligopoly under monitoring or moratorium	containment, evidence, and political control; attention, rights, and action replace compute scarcity
L9	Plural	Scarce	Trusted	Centralized	state, suite, or enterprise hierarchy	capture, allocation legitimacy, and concentration; power, compute, and capacity remain priced
L10	Plural	Scarce	Trusted	Polycentric	open services, federated procurement, and local institutions	interoperability, reputation, and dispute translation; power, compute, and capacity remain priced
L11	Plural	Scarce	Untrusted	Centralized	shadow bazaar plus local allowlists and human principals	identity, consent, fraud, and common-mode dependencies; power, compute, and capacity remain priced
L12	Plural	Scarce	Untrusted	Polycentric	shadow bazaar plus local allowlists and human principals	identity, consent, fraud, and common-mode dependencies; power, compute, and capacity remain priced
L13	Plural	Abundant	Trusted	Centralized	state, suite, or enterprise hierarchy	capture, allocation legitimacy, and concentration; attention, rights, and action replace compute scarcity
L14	Plural	Abundant	Trusted	Polycentric	open services, federated procurement, and local institutions	interoperability, reputation, and dispute translation; attention, rights, and action replace compute scarcity

Cell	Supply	Resources	Trust	Authority	Likely structure	Binding issue
L15	Plural	Abundant	Untrusted	Centralized	shadow bazaar plus local allowlists and human principals	identity, consent, fraud, and common-mode dependencies; attention, rights, and action replace compute scarcity
L16	Plural	Abundant	Untrusted	Polycentric	shadow bazaar plus local allowlists and human principals	identity, consent, fraud, and common-mode dependencies; attention, rights, and action replace compute scarcity

Public-good and cooperative equilibrium

[SPECULATION] Governments, universities, consortia, unions, or industry cooperatives fund identity, verification, protocols, and compute as shared infrastructure. Private demand exists, but rent capture is intentionally limited. The relevant indicators are public procurement, open maintenance funding, mandated interoperability, and cooperative governance. The falsifier is persistent preference for proprietary service despite a credible public option.

Institutional fracture

[SPECULATION] Courts, regulators, insurers, payment networks, standards bodies, and states disagree about authority and evidence. No receipt is universally accepted. Translation and local compliance grow, but global network effects weaken. The falsifier is rapid convergence on one cross-jurisdiction rulebook with enforceable recognition.

Labor and public veto

[SPECULATION] Workers, professions, citizens, merchants, or communities block autonomous substitution, surveillance, or machine-mediated decisions even when capability is strong. Budgets move toward consultation, augmentation, compensation, and human staffing. The falsifier is broad adoption of consequential autonomy without durable bargaining, legal challenge, or political restriction.

Synthetic-demand spiral

[SPECULATION] Agents generate tasks for other agents, retries create apparent volume, benchmarks and marketing manufacture demand, and subsidized machine transactions circulate without human surplus. Infrastructure spend and transaction counts rise while external willingness to pay and net welfare remain weak. The falsifier is persistent unsubsidized payment by independent principals tied to externally valued outcomes.

Ecological and social constraint

[SPECULATION] Power, water, land, emissions, supply chains, or community opposition become binding before compute price signals fully adjust. States ration or condition infrastructure; local legitimacy becomes part of deployment. The falsifier is abundant low-impact energy and rapid siting without political resistance.

Informal, criminal, and adversarial agent economy

[SPECULATION] Autonomous activity scales outside compliant identity, audit, procurement, and payment rails through stolen credentials, coercion, unlicensed models, evasion, and jurisdiction shopping. Governance demand can rise for compliant actors while effective system-wide authority shrinks. The falsifier is sustained migration of consequential activity into enforceable, attributable rails.

Machine legal standing and autonomous assets

[SPECULATION] Law recognizes machine-managed entities, autonomous trusts, or direct machine capacity to hold assets and obligations. Principal-based delegation must be rebuilt around capitalization, dissolution, constitutional limits, and reachability. The falsifier is durable legal insistence that a natural or conventional legal person remains the responsible principal.

Catastrophic infrastructure regression

[SPECULATION] War, grid failure, cyber conflict, climate disaster, or supply-chain rupture reduces compute, communications, and institutional capacity. Advanced systems persist unevenly, but offline operation, manual skills, local energy, and degraded coordination dominate. The falsifier is resilient infrastructure that maintains abundance and institutional reach through severe shocks.

Non-market abundance and political allocation

[SPECULATION] Intelligence becomes abundant enough that prices and software markets contract while access is allocated as a public entitlement, constitutional service, cooperative commons, or rationed utility. Distribution and legitimacy dominate monetization. The falsifier is persistent private willingness to pay for differentiated intelligence under broad public access.

Transition points that discriminate worlds

The high-leverage transitions are not calendar dates. They are state changes:

- Capability transition: agents sustain longer task horizons with lower supervision.
- Economic transition: mature cohorts shift from additive seats to captured P&L value.
- Authority transition: agents receive credentials for irreversible external action.
- Institutional transition: courts, regulators, insurers, and auditors accept or reject machine evidence and mandates.
- Market transition: at least two ownership-independent suppliers submit executable offers to a funded buyer.
- Settlement transition: delivery and payment obligations acquire explicit finality and default rules.
- Physical transition: embodied systems make energy, machines, property, and safety cases central.
- Superintelligence transition: capability changes faster than institutions can update.

At each transition, the same technical function can be a product, a bundled feature, a regulated utility obligation, an internal capability, or irrelevant.

Counterfactual reconstruction: reverse the attractive claims


The table does not ask whether the base assumption is plausible. It asks what follows if its opposite world holds long enough to determine market structure.

Assumption	Opposite world	Discriminating evidence	Consequence
Unit inference cost keeps falling	Energy, chips, or frontier rents reverse the decline	normalized real price by capability tier	Routing and volume economics change
Total AI spend rises	Demand saturates after experimentation	cohort retention and total risk-adjusted spend	Infrastructure and FinOps opportunities shrink
Many agents exist	Suites consolidate functions into one delegate	active independent production identities	Fleet-control demand weakens
Suppliers are independent	Different logos share model, cloud, data, or parent	dependency-adjusted concentration	Auctions and risk diversification become hollow
Open models close the gap	Frontier scale economics compound	held-out consequential-work gap	Open-commons scenarios weaken
Enterprises multi-home	Integration and committed spend favor one default	successful production switching	Neutral routing rent collapses
Tasks are contractible	Goals remain tacit, political, and path dependent	ex ante acceptance completeness	Markets and verification stay thin
Outcomes are observable	Delayed externalities dominate visible checks	loss discovery lag and reversal rate	Assurance claims overstate value
Independent proof is valued	Provider logs gain institutional acceptance	procurement budget for external proof	Verifier opportunity is enclosed
Authority stays scarce	Trusted SI gains direct institutional standing	court and regulator treatment	Authority-layer thesis weakens
Legal persons remain accountable	Machine entities receive autonomous legal capacity	statute and case law	Principal-based controls must change
Humans remain in final control	Oversight becomes ceremonial or impossible	actual override frequency and competence	Human-in-loop claims become fiction
Machine commerce grows	Agents only recommend; humans transact	binding delegated transactions net of confirmation	Payment and market layers stay premature
Reputation is portable	Platforms, rights, and nonstationarity prevent transfer	cross-venue predictive use	Marketplace liquidity weakens
Open standards weaken incumbents	Standards enlarge demand for bundled implementations	share of value captured outside sponsors	Protocol optimism reverses
Regulation creates openings	Compliance fixed costs concentrate supply	entry, designation, and audit concentration	Independent layers lose
Neutrality can be monetized	Fees, ranking, verification, and insurance create conflicts	revenue-source and decision audits	Neutral intermediary lacks legitimacy
Customer-owned enforcement persists	Provider identity and data become unavoidable	control over keys, policies, and receipts	External authority is nominal
Physical resources remain scarce	Energy/hardware abundance outpaces demand	utilization and capacity premium	Compute/energy opportunities fade

Assumption	Opposite world	Discriminating evidence	Consequence
Context remains scarce	Models infer or synthesize sufficient context	marginal value of private data	Context custody rent erodes
Markets achieve liquidity	Bypass, heterogeneity, subsidies, or concentration dominate	unsubsidized repeat fills and depth	Clearing/market infrastructure dies
Routing beats run-all	Execution and verification become nearly free	net value after all candidate runs	Ex ante allocator commoditizes
Providers cannot self-certify	Proof systems become accepted and cheap	external-vs-native evidence premium	Independent verification contracts
Systems of record cannot bundle control	Workflow incumbents close the governance loop	separate budget and displacement rate	Horizontal control points disappear
Technical boundaries resist SI	Agents socially or technically acquire credentials	boundary escape and privilege escalation	Sandbox confidence is false
Legal finality persists	Emergency powers or private arbitration rewrite rights	enforcement and dispute precedent	Settlement assumptions change
Losses are insurable	Common-mode, strategic, or catastrophic failures dominate	limits, exclusions, reinsurance, paid claims	Underwriting cannot form
Timing leaves a startup window	Incumbents bundle before buyers form budgets	time to native feature parity	Revenue appears but no durable company forms

ASSUMPTION KILL TREE

Each attractive conclusion depends on multiple bridges



ODYLITH / MARKET EVOLUTION SITUATIONAL AWARENESS / 5 AUGUST 2026

Assumption bridges that can fail independently

Hidden confounders and partial observability

Observed growth can be misread. Internal chargebacks can be counted as market transactions; human approval can hide behind an autonomous label; calls inside one annual agreement can look like per-step procurement; different suppliers can share one model or cloud; test credits can look like settled value; failures and retries can look like demand; listings can look like supply; bundled cost can be imputed as a discovered price; one capacity block can be listed repeatedly; verifier and seller can share ownership; reciprocal transactions can fabricate reputation; router rebates can hide steering; benchmark leakage can mimic verification; one sponsor can create all demand; gross volume can ignore refunds; successful tasks can be the only ones observed; protocols can hide jurisdictional barriers.

Enterprise value has parallel traps: purchased licenses, monthly users, deployed agents, successful calls, accepted outputs, and P&L value are not interchangeable. Time saved can become slack, higher quality, work intensification, or expanded demand. Cleanup labor, prompt preparation, exception handling, and reputation risk can be missing. Human-in-the-loop can be nominal. Average reliability can hide correlated tail loss. Announced regulation can differ from enforcement.

The response is denominator discipline: define the principal, legal counterparty, independent supplier, binding mandate, accepted outcome, net settled value, human intervention, full cost, reversal, default, and delayed loss before treating a metric as evidence.

Game theory: where the mechanism becomes unstable

Adverse selection: suppliers know their hidden model, subcontractor, cost, and failure modes. Moral hazard: a winner cuts compute or verification after award. Principal-agent failure: the buyer agent optimizes a proxy or the router optimizes rebates. Signaling: certifications measure wealth or test optimization rather than quality. Collusion: shared pricing algorithms or repeated auctions coordinate behavior. Foreclosure: integrated platforms deny rivals identity, data, latency, or distribution. Self-preference: the market operator ranks its own supply. Bypass: strong buyers and sellers leave after matching. Correlated risk: nominal diversification shares hidden dependencies. Verifier capture: rated parties fund or own the judge. Liquidity spiral: poor execution drives away the opposite side. Default and settlement risk: delivery and money do not become final together.

Mechanism resistance requires structural separation, observable conflicts, portable identity and history, multi-homing, public rulebooks, collateral where appropriate, appeal, and a buyer who can exit. None is guaranteed by intelligence.

What superintelligence resistance actually means

Resistance raises the cost or limits the ability of a capable system to violate a boundary. Resilience preserves critical function after failure. Bargaining power keeps the institution from becoming a price taker. Antifragility improves the option set or bargaining position because shocks reveal information and trigger reallocation. These are not synonyms.

Technical controls - policies expressed to the same model, software monitors, or model-generated proofs - are weak against a system able to reason around, manipulate, or socially acquire access. Resistance becomes more credible when anchored outside the system's unilateral control: physically separated keys, diverse operators, legal sanctions against reachable principals, property rights, network segmentation, resource quotas, independent institutions, human coalitions, and reversible authority.

Even these anchors are conditional. A superintelligence can persuade key holders, exploit law, create dependency, or help one coalition dominate another. The correct claim is not that governance survives intelligence. It is that some institutional and physical mechanisms may remain costly to bypass while the relevant state retains coercion, property, legitimacy, plurality, and response time.

First-class nulls include institutional collapse, machine legal standing, emergency suspension of rights, persuasion or capture of key holders, loss of a reachable defendant, and capability change faster than detection, rulemaking, deployment, or enforcement. Current property and law identify transition friction; they are not durable anchors without a scenario-specific enforcement mechanism.

SUPERINTELLIGENCE SURVIVAL SPACE

Technical controls survive only when anchored in scarcities intelligence cannot erase

Technical	Sandbox, keys, networks	Can be bypassed if authority is exposed
Institutional	Law, accountable legal person	Survives while institutions retain coercion
Economic	Property, scarce resources, collateral	Survives if access is enforceable
Architectural	Separation, heterogeneity, reversibility	Raises blast-radius resistance
Strategic	Options, multi-homing, low fixed exposure	Changes payoff under uncertainty

Superintelligence resistance matrix

Family	World	Mechanism	What it resists	Failure boundary
Technical	oligopoly	sandboxes, least privilege, signed binaries, isolated keys, network and compute controls	direct privilege escalation and replication	fails if agents obtain credentials socially or control the monitor
Technical	open commons	sandboxes, least privilege, signed binaries, isolated keys, network and compute controls	direct privilege escalation and replication	fails if agents obtain credentials socially or control the monitor
Technical	trusted SI	sandboxes, least privilege, signed binaries, isolated keys, network and compute controls	direct privilege escalation and replication	fails if agents obtain credentials socially or control the monitor
Technical	untrusted SI	sandboxes, least privilege, signed binaries, isolated keys, network and compute controls	direct privilege escalation and replication	fails if agents obtain credentials socially or control the monitor
Technical	fast takeoff	sandboxes, least privilege, signed binaries, isolated keys, network and compute controls	direct privilege escalation and replication	fails if agents obtain credentials socially or control the monitor
Institutional	oligopoly	named legal principals, courts, licensing, public oversight, professional duty	deterrence, redress, legitimate refusal	fails under capture, emergency, illegitimate states, or machine legal standing
Institutional	open commons	named legal principals, courts, licensing, public oversight, professional duty	deterrence, redress, legitimate refusal	fails under capture, emergency, illegitimate states, or machine legal standing
Institutional	trusted SI	named legal principals, courts, licensing, public oversight, professional duty	deterrence, redress, legitimate refusal	fails under capture, emergency, illegitimate states, or machine legal standing
Institutional	untrusted SI	named legal principals, courts, licensing, public oversight, professional duty	deterrence, redress, legitimate refusal	fails under capture, emergency, illegitimate states, or machine legal standing
Institutional	fast takeoff	named legal principals, courts, licensing, public oversight, professional duty	deterrence, redress, legitimate refusal	fails under capture, emergency, illegitimate states, or machine legal standing
Economic	oligopoly	property, collateral, scarce resources, diversification, switching	limits resource access and preserves outside options	fails if one provider controls identity, capital, compute, and distribution
Economic	open commons	property, collateral, scarce resources, diversification, switching	limits resource access and preserves outside options	fails if one provider controls identity, capital, compute, and distribution
Economic	trusted SI	property, collateral, scarce resources, diversification, switching	limits resource access and preserves outside options	fails if one provider controls identity, capital, compute, and distribution
Economic	untrusted SI	property, collateral, scarce resources, diversification, switching	limits resource access and preserves outside options	fails if one provider controls identity, capital, compute, and distribution
Economic	fast takeoff	property, collateral, scarce resources, diversification, switching	limits resource access and preserves outside options	fails if one provider controls identity, capital, compute, and distribution
Architectural	oligopoly	separation of execution, verification, authorization, settlement; heterogeneous redundancy	contains blast radius and common-mode error	fails when components share dependencies or one SI dominates all roles
Architectural	open commons	separation of execution, verification, authorization, settlement; heterogeneous redundancy	contains blast radius and common-mode error	fails when components share dependencies or one SI dominates all roles

Family	World	Mechanism	What it resists	Failure boundary
Architectural	trusted SI	separation of execution, verification, authorization, settlement; heterogeneous redundancy	contains blast radius and common-mode error	fails when components share dependencies or one SI dominates all roles
Architectural	untrusted SI	separation of execution, verification, authorization, settlement; heterogeneous redundancy	contains blast radius and common-mode error	fails when components share dependencies or one SI dominates all roles
Architectural	fast takeoff	separation of execution, verification, authorization, settlement; heterogeneous redundancy	contains blast radius and common-mode error	fails when components share dependencies or one SI dominates all roles
Strategic	oligopoly	reversible commitments, multi-homing, short contracts, experiments, contingency plans	preserves learning and bargaining across worlds	fails when takeoff is faster than exercise or options depend on the same chokepoint
Strategic	open commons	reversible commitments, multi-homing, short contracts, experiments, contingency plans	preserves learning and bargaining across worlds	fails when takeoff is faster than exercise or options depend on the same chokepoint
Strategic	trusted SI	reversible commitments, multi-homing, short contracts, experiments, contingency plans	preserves learning and bargaining across worlds	fails when takeoff is faster than exercise or options depend on the same chokepoint
Strategic	untrusted SI	reversible commitments, multi-homing, short contracts, experiments, contingency plans	preserves learning and bargaining across worlds	fails when takeoff is faster than exercise or options depend on the same chokepoint
Strategic	fast takeoff	reversible commitments, multi-homing, short contracts, experiments, contingency plans	preserves learning and bargaining across worlds	fails when takeoff is faster than exercise or options depend on the same chokepoint

Antifragility under model and institutional shocks

A system is not antifragile merely because it has redundancy, a dashboard, or a fallback. A genuinely antifragile arrangement uses variation and failure to improve future choices without requiring catastrophic exposure. Candidate mechanisms include small reversible deployments that reveal loss distributions; dependency graphs that turn provider incidents into bargaining evidence; portable contracts and data that make switching rehearsals useful; heterogeneous verifiers that expose disagreement; capped authority that localizes experiments; procurement options that remain exercisable; and governance that updates from near misses without hiding them.


The anti-patterns are diversification by logo over a common stack, insurance without claims capacity, an open standard with proprietary identity, human oversight without time or power, and multi-cloud contracts that cannot move data or operations. These simulate resilience while preserving the common mode.

Antifragility test: after a model leap, provider failure, regulation, or incident, does the actor have more credible routes, better calibrated limits, lower switching friction, and greater bargaining power? If the shock merely increases cost or reinforces the incumbent, the arrangement is resilient at best and fragile at worst.

The stricter chain is: shock -> protected downside -> captured learning or exercisable right -> measurable out-of-sample gain -> retained rights -> bounded failure. A larger log corpus, survival, restored service, or post-incident demand is not a gain by itself. Every claim needs a gain unit, an option owner, premium, trigger, expiry, exercise terms, dependency-adjusted exercisability, positive-convexity interval, and abandonment rule. The mechanisms in this dossier remain hypotheses; none is asserted empirically antifragile.

ANTIFRAGILITY TEST

Robustness absorbs shocks; antifragility improves the option set because of them




ODYLITH / MARKET EVOLUTION SITUATIONAL AWARENESS / 5 AUGUST 2026

Shock-to-option cycle required for antifragility

Real options and temporal strategy without selecting a product

Under deep uncertainty, an option is valuable when it is affordable, exercisable, independent of the same failure mode, and informative. Examples include portable data and contracts, multi-provider qualification, customer-operated components, small experiments in delegated authority, protocol participation without irreversible dependence, scenario-specific partnerships, and contractual exit rights. Capital-heavy compute, insurance balance sheets, clearing licenses, exclusive distribution, and sovereign infrastructure are longer-duration commitments whose payoff requires stronger evidence.

An early revenue category can be rational but non-durable; a durable scarcity can be real but too early to monetize. The temporal question is whether present activity buys learning and future access or locks the actor into an expiring layer. That question is deliberately separate from a recommendation.



ODYLITH / MARKET EVOLUTION SITUATIONAL AWARENESS / 5 AUGUST 2026

Reversibility and learning across commitment classes

Institutional opportunity atlas: functions without an assumed company


An opportunity surface is included only if a possible world creates a beneficiary, a budget or collective funding mechanism, and a scarcity. Inclusion does not imply timing, durability, venture scale, founder fit, or even private-company formation. The organizational home is deliberately unresolved: a function can live inside an enterprise, incumbent suite, open protocol, cooperative, professional service, accredited body, regulated utility, public institution, insurer, marketplace, or nowhere as a separate layer.

The buyer and revenue columns are counterfactual tests: if no actor can fund the function on the stated meter, the commercial branch fails. Public-good provision, internalization, bundling, and uncompensated open infrastructure are equally valid equilibria. Present revenue potential and long-run defensibility are separated because a category can monetize during transition and then be bundled, regulated, nationalized, open-sourced, or erased.

Every surface remains unproven until seven gates pass independently: function remains necessary; loss is measurable without it; authority sits with an actor able to change the workflow; budget exists separately after native controls are configured; access gives an independent supplier enough telemetry and enforcement rights; rent survives bundling, standards, and competition; and time exceeds procurement, integration, certification, and capital formation. A function can pass the first gate and fail all six economic gates.

OPPORTUNITY EMERGENCE

Revenue can appear early in a category that later commoditizes



ODYLITH / MARKET EVOLUTION SITUATIONAL AWARENESS / 5 AUGUST 2026

Opportunity categories emerge at different transitions

OPPORTUNITY - WORLD MATRIX

A surface is interesting only where its scarcity persists and a buyer can pay

	Oligopoly	Open	Sovereign	Trusted SI	Untrusted SI	Embodied
Routing	HIGH	COND	COND	LOW	COND	COND
Identity / authority	HIGH	HIGH	HIGH	COND	HIGH	HIGH
Verification	HIGH	COND	HIGH	LOW	HIGH	HIGH
Payments / settlement	COND	HIGH	COND	LOW	COND	HIGH
Sovereign compute	COND	LOW	HIGH	LOW	COND	COND
Physical oracles	LOW	COND	HIGH	COND	HIGH	HIGH
Insurance	COND	COND	HIGH	LOW	COND	HIGH
Human consent	HIGH	HIGH	HIGH	LOW	HIGH	HIGH

Illustrative conditionality map - not a probability or recommendation

ODYLITH / MARKET EVOLUTION SITUATIONAL AWARENESS / 5 AUGUST 2026

Illustrative opportunity conditionality across worlds

Surface	Formation condition	Funding locus	Scarcity	Revenue	Natural enclosure	Capital / rules	Death condition
Model and agent routing	Multiple consequential routes remain differentiated and switching is feasible.	CIO, platform engineering, high-volume application owner	Local price-performance variance and portfolio continuity	usage, subscription, or managed capacity fee	bundle routing and reserve best latency	low / high	Routing becomes a native model capability or run-all becomes cheaper than selection.
Orchestration and workflow state	Long-lived multi-step work spans tools and agents.	COO, CIO, function P&L	Reliable state, recovery, and exception ownership	platform subscription or consumption	absorb orchestration into suite workflow	medium / medium	One end-to-end provider or SI makes external orchestration redundant.
Context custody	Enterprise-specific data remains nonportable, rights-constrained, and causally important.	CIO, data officer, domain owner	Authoritative local state and permissioned memory	storage, retrieval, governance subscription	use data gravity and proprietary memory	medium / medium	Models internalize needed context or systems of record expose it freely.
Agent identity	Agents act across domains and principals need attribution and revocation.	CISO, IAM owner, payment network	Binding identity and lifecycle	per identity, tenant, or verification	identity suites bundle it	medium / medium	Agents remain advisory or provider accounts suffice.
Delegated authorization	Principals permit bounded autonomous action without per-action confirmation.	CISO, legal, payments, application owner	Scoped mandate, limits, consent, revocation	per authorized action or platform fee	payments and IAM providers couple authority to their rails	medium / high	Humans approve every consequential action or providers own delegation.
Provenance and custody receipts	Institutions demand reconstructable action histories.	audit, regulator, insurer, enterprise risk	Evidence integrity across parties	per attestation, retention, subscription	platforms issue proprietary receipts	medium / high	Provider-native logs are accepted or self-verifying SI commoditizes proof.

Surface	Formation condition	Funding locus	Scarcity	Revenue	Natural enclosure	Capital / rules	Death condition
Independent verification	Output validity is observable and buyers reject self-certification.	risk owner, procurement, regulator	Credible separation from executor	per evaluation, subscription, certification	providers bundle tests or acquire verifiers	medium / high	Verification becomes perfect/free or institutions accept provider proof.
Agent observability and incident response	Autonomous actions create operational failure and telemetry is obtainable.	CISO, SRE, CIO	Cross-stack detection and reconstruction	usage or enterprise subscription	security and cloud suites bundle	low / medium	Native tools achieve adequate visibility or agents evade observation.
Policy and governance	Enterprises need consistent constraints across a plural fleet.	CISO, legal, CIO, regulator	Organization-specific rules and exception process	subscription, per controlled action	systems of record and IAM vendors bundle	medium / high	Governance remains advisory or one platform controls all agents.
AI FinOps	Total consumption is material, variable, and allocable to owners.	CFO, CIO, FinOps	Spend attribution and budget enforcement	percentage of managed spend or subscription	clouds, models, observability suites commoditize	low / medium	Unit costs approach zero or native cost controls suffice.
Machine payments	Agents make external purchases under valid mandates.	banks, payment networks, merchants, platforms	Fraud control, authorization, and merchant recognition	payment fee, credential fee	payment incumbents own identity and acceptance	high / high	Transactions remain human-confirmed or one platform internalizes commerce.
Procurement auctions and RFQs	Tasks or capacity are specifiable and multiple independent suppliers bid.	procurement, platform, state	Competitive allocation and auditable price discovery	buyer fee, membership, transaction fee	procurement networks and clouds add it	medium / high	Hierarchy is cheaper, suppliers consolidate, or specifications stay heterogeneous.
Portable reputation	Buyers repeatedly select among independent sellers and histories can legally travel.	market operator, procurement, seller	Task-conditioned performance under identity continuity	membership, data service, verification fee	platforms keep reputation captive	medium / high	Markets stay thin or performance is nonstationary and nontransferable.
Insurance and risk transfer	Loss classes are bounded, observable, diversifiable, and legally insurable.	CFO, risk committee, operator	Risk capital and credible compensation	premium, brokerage, reinsurance	insurers partner with major platforms	high / very high	Losses are correlated, nonstationary, catastrophic, or premiums exceed savings.
Dispute resolution	Autonomous transactions produce contested authority, delivery, or harm.	market operator, payments, enterprises	Accepted rulebook, neutral adjudication, remedies	case fee, membership, levy	payment networks and courts retain jurisdiction	medium / high	Disputes stay rare/simple or provider terms force captive arbitration.
Settlement and clearing	Multiple venues and counterparties create credit, liquidity, and finality risk.	market operators, banks, large buyers	Netting, finality, collateral, default waterfall	clearing and membership fees	financial incumbents control rails	very high / very high	Underlying markets never deepen or gross settlement is adequate.
Data and content rights	Agents require licensed content, private data, or generated-output rights.	publishers, data owners, enterprises, agents	Enforceable access and attribution	license, royalty, usage fee	platforms negotiate blanket licenses	medium / high	Law creates broad exceptions or synthetic data eliminates scarcity.

Opportunity overview - remaining surfaces

Surface	Formation condition	Funding locus	Scarcity	Revenue	Natural enclosure	Capital / rules	Death condition
Sovereign infrastructure	Jurisdiction, continuity, or national security outweigh lowest cost.	states, regulated enterprises	Local control over keys, operators, supply chain, and compute	long-term contracts, infrastructure service	hyperscalers offer sovereignty wrappers	very high / very high	Buyers accept global safeguards or sovereignty premium becomes politically untenable.
Compute, power, and capacity	Physical demand exceeds supply or locality matters.	clouds, enterprises, states, agent platforms	Power, chips, interconnect, financing, siting	capacity contract, optimization fee, financing spread	hyperscalers vertically integrate	very high / high	Efficiency and demand saturation create persistent slack.
Physical oracles and actuation	Agents affect factories, logistics, laboratories, and property.	operators, insurers, regulators	Trusted sensing, machine state, and reversible actuation	equipment, service, per verified event	industrial incumbents bundle	high / high	Embodied autonomy stalls or sensors become universal commodities.
Enterprise transformation	Capability exists but organizations cannot redesign coordinated work.	CEO, COO, HR, CIO	Political coalition, process redesign, training, accountability	services, managed operations, software-assisted program	consultancies and suites dominate	low / medium	AI remains individual augmentation or SI directly redesigns organizations.
Systems-of-record distribution	Agents need authoritative data and transaction endpoints.	application vendors and domain platforms	Installed workflow, data gravity, legal record	seat, consumption, take rate	incumbents already own it	high / medium	Interface-obsolete SI reduces direct app value and mandates portability.
Certification and accreditation	Regulators/procurement accept standardized proof as an eligibility gate.	regulator, auditor, procurement	Institutional recognition	assessment, renewal, registry fee	large auditors and vendors shape standards	medium / high	Certifications become checkbox theater or continuous proof replaces periodic review.
Protocol infrastructure	Many parties need interoperable discovery, authority, evidence, and payment messages.	developers, platforms, standards bodies	Registry governance, conformance operations, portable trust roots, and full-workflow switching - only if independently controlled	hosting, conformance, enterprise support	incumbents support open core but capture services	low / medium	Protocols remain message formats without production transactions.
Human consent and fiduciary control	Institutions refuse to treat machine optimization as legitimate authority.	consumers, boards, professions, state	Meaningful consent, judgment, accountable natural/legal person	service, compliance, public infrastructure	providers turn consent into captive interface	medium / very high	Trusted SI receives direct institutional standing or consent becomes nominal.

O1. Model and agent routing

Formation condition. Multiple consequential routes remain differentiated and switching is feasible.

Possible payer or funding locus. CIO, platform engineering, high-volume application owner. A separate commercial budget is not assumed.

Scarcity captured. Local price-performance variance and portfolio continuity.

Possible revenue mechanism. usage, subscription, or managed capacity fee. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. Function or capability; no external market is implied.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. telemetry, comparable tasks, provider access.

Incumbent counterstrategy. bundle routing and reserve best latency.

Capital / regulation / reversibility. Capital intensity: low; regulatory exposure: high; reversibility: medium.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if routing becomes a native model capability or run-all becomes cheaper than selection. At the same time, the natural enclosure path can bundle routing and reserve best latency, while access to telemetry, comparable tasks, provider access disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Multiple consequential routes remain differentiated and switching is feasible. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. Routing becomes a native model capability or run-all becomes cheaper than selection.

O2. Orchestration and workflow state

Formation condition. Long-lived multi-step work spans tools and agents.

Possible payer or funding locus. COO, CIO, function P&L. A separate commercial budget is not assumed.

Scarcity captured. Reliable state, recovery, and exception ownership.

Possible revenue mechanism. platform subscription or consumption. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. Function or capability; no external market is implied.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. systems-of-record integration.

Incumbent counterstrategy. absorb orchestration into suite workflow.

Capital / regulation / reversibility. Capital intensity: medium; regulatory exposure: medium; reversibility: medium.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if one end-to-end provider or si makes external orchestration redundant. At the same time, the natural enclosure path can absorb orchestration into suite workflow, while access to systems-of-record integration disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Long-lived multi-step work spans tools and agents. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for

retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. One end-to-end provider or SI makes external orchestration redundant.

O3. Context custody

Formation condition. Enterprise-specific data remains nonportable, rights-constrained, and causally important.

Possible payer or funding locus. CIO, data officer, domain owner. A separate commercial budget is not assumed.

Scarcity captured. Authoritative local state and permissioned memory.

Possible revenue mechanism. storage, retrieval, governance subscription. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. Function or capability; no external market is implied.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. data connectors, access policy, provenance.

Incumbent counterstrategy. use data gravity and proprietary memory.

Capital / regulation / reversibility. Capital intensity: medium; regulatory exposure: medium; reversibility: high.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if models internalize needed context or systems of record expose it freely. At the same time, the natural enclosure path can use data gravity and proprietary memory, while access to data connectors, access policy, provenance disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Enterprise-specific data remains nonportable, rights-constrained, and causally important. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. Models internalize needed context or systems of record expose it freely.

O4. Agent identity

Formation condition. Agents act across domains and principals need attribution and revocation.

Possible payer or funding locus. CISO, IAM owner, payment network. A separate commercial budget is not assumed.

Scarcity captured. Binding identity and lifecycle.

Possible revenue mechanism. per identity, tenant, or verification. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. Function or capability; no external market is implied.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. directory, credentials, legal principal.

Incumbent counterstrategy. identity suites bundle it.

Capital / regulation / reversibility. Capital intensity: medium; regulatory exposure: medium; reversibility: high.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if agents remain advisory or provider accounts suffice. At the same time, the natural enclosure path can identity suites bundle it, while access to directory, credentials, legal principal disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Agents act across domains and principals need attribution and revocation. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. Agents remain advisory or provider accounts suffice.

O5. Delegated authorization

Formation condition. Principals permit bounded autonomous action without per-action confirmation.

Possible payer or funding locus. CISO, legal, payments, application owner. A separate commercial budget is not assumed.

Scarcity captured. Scoped mandate, limits, consent, revocation.

Possible revenue mechanism. per authorized action or platform fee. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. Function or capability; no external market is implied.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. identity, policy engine, legal recognition.

Incumbent counterstrategy. payments and IAM providers couple authority to their rails.

Capital / regulation / reversibility. Capital intensity: medium; regulatory exposure: high; reversibility: high.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if humans approve every consequential action or providers own delegation. At the same time, the natural enclosure path can payments and IAM providers couple authority to their rails, while access to identity, policy engine, legal recognition disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Principals permit bounded autonomous action without per-action confirmation. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. Humans approve every consequential action or providers own delegation.

O6. Provenance and custody receipts

Formation condition. Institutions demand reconstructable action histories.

Possible payer or funding locus. audit, regulator, insurer, enterprise risk. A separate commercial budget is not assumed.

Scarcity captured. Evidence integrity across parties.

Possible revenue mechanism. per attestation, retention, subscription. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. Function or capability; no external market is implied.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. signed events, clocks, data rights.

Incumbent counterstrategy. platforms issue proprietary receipts.

Capital / regulation / reversibility. Capital intensity: medium; regulatory exposure: high; reversibility: high.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if provider-native logs are accepted or self-verifying si commoditizes proof. At the same time, the natural enclosure path can platforms issue proprietary receipts, while access to signed events, clocks, data rights disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Institutions demand reconstructable action histories. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. Provider-native logs are accepted or self-verifying SI commoditizes proof.

07. Independent verification

Formation condition. Output validity is observable and buyers reject self-certification.

Possible payer or funding locus. risk owner, procurement, regulator. A separate commercial budget is not assumed.

Scarcity captured. Credible separation from executor.

Possible revenue mechanism. per evaluation, subscription, certification. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. Function or capability; no external market is implied.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. task contractibility, evaluator independence.

Incumbent counterstrategy. providers bundle tests or acquire verifiers.

Capital / regulation / reversibility. Capital intensity: medium; regulatory exposure: high; reversibility: medium.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if verification becomes perfect/free or institutions accept provider proof. At the same time, the natural enclosure path can providers bundle tests or acquire verifiers, while access to task contractibility, evaluator independence disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Output validity is observable and buyers reject self-certification. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. Verification becomes perfect/free or institutions accept provider proof.

08. Agent observability and incident response

Formation condition. Autonomous actions create operational failure and telemetry is obtainable.

Possible payer or funding locus. CISO, SRE, CIO. A separate commercial budget is not assumed.

Scarcity captured. Cross-stack detection and reconstruction.

Possible revenue mechanism. usage or enterprise subscription. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. Function or capability; no external market is implied.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. instrumentation, access, response workflows.

Incumbent counterstrategy. security and cloud suites bundle.

Capital / regulation / reversibility. Capital intensity: low; regulatory exposure: medium; reversibility: medium.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if native tools achieve adequate visibility or agents evade observation. At the same time, the natural enclosure path can security and cloud suites bundle, while access to instrumentation, access, response workflows disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Autonomous actions create operational failure and telemetry is obtainable. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. Native tools achieve adequate visibility or agents evade observation.

O9. Policy and governance

Formation condition. Enterprises need consistent constraints across a plural fleet.

Possible payer or funding locus. CISO, legal, CIO, regulator. A separate commercial budget is not assumed.

Scarcity captured. Organization-specific rules and exception process.

Possible revenue mechanism. subscription, per controlled action. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. Function or capability; no external market is implied.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. authority, integration, enforcement point.

Incumbent counterstrategy. systems of record and IAM vendors bundle.

Capital / regulation / reversibility. Capital intensity: medium; regulatory exposure: high; reversibility: high.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if governance remains advisory or one platform controls all agents. At the same time, the natural enclosure path can systems of record and IAM vendors bundle, while access to authority, integration, enforcement point disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Enterprises need consistent constraints across a plural fleet. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the

death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. Governance remains advisory or one platform controls all agents.

O10. AI FinOps

Formation condition. Total consumption is material, variable, and allocable to owners.

Possible payer or funding locus. CFO, CIO, FinOps. A separate commercial budget is not assumed.

Scarcity captured. Spend attribution and budget enforcement.

Possible revenue mechanism. percentage of managed spend or subscription. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. Function or capability; no external market is implied.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. cost telemetry and org mapping.

Incumbent counterstrategy. clouds, models, observability suites commoditize.

Capital / regulation / reversibility. Capital intensity: low; regulatory exposure: medium; reversibility: low.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if unit costs approach zero or native cost controls suffice. At the same time, the natural enclosure path can clouds, models, observability suites commoditize, while access to cost telemetry and org mapping disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Total consumption is material, variable, and allocable to owners. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. Unit costs approach zero or native cost controls suffice.

O11. Machine payments

Formation condition. Agents make external purchases under valid mandates.

Possible payer or funding locus. banks, payment networks, merchants, platforms. A separate commercial budget is not assumed.

Scarcity captured. Fraud control, authorization, and merchant recognition.

Possible revenue mechanism. payment fee, credential fee. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. Payment and mandate infrastructure; not by itself an underlying market.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. consumer authority, dispute rules, existing rails.

Incumbent counterstrategy. payment incumbents own identity and acceptance.

Capital / regulation / reversibility. Capital intensity: high; regulatory exposure: high; reversibility: high.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if transactions remain human-confirmed or one platform internalizes commerce. At the same time, the natural enclosure path can payment incumbents own identity and acceptance, while access to consumer authority, dispute rules, existing rails disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Agents make external purchases under valid mandates. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. Transactions remain human-confirmed or one platform internalizes commerce.

O12. Procurement auctions and RFQs

Formation condition. Tasks or capacity are specifiable and multiple independent suppliers bid.

Possible payer or funding locus. procurement, platform, state. A separate commercial budget is not assumed.

Scarcity captured. Competitive allocation and auditable price discovery.

Possible revenue mechanism. buyer fee, membership, transaction fee. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. A procurement market only when independent principals submit executable bids under liability and remedy.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. qualified supply, binding offers, acceptance.

Incumbent counterstrategy. procurement networks and clouds add it.

Capital / regulation / reversibility. Capital intensity: medium; regulatory exposure: high; reversibility: medium.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if hierarchy is cheaper, suppliers consolidate, or specifications stay heterogeneous. At the same time, the natural enclosure path can procurement networks and clouds add it, while access to qualified supply, binding offers, acceptance disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Tasks or capacity are specifiable and multiple independent suppliers bid. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. Hierarchy is cheaper, suppliers consolidate, or specifications stay heterogeneous.

O13. Portable reputation

Formation condition. Buyers repeatedly select among independent sellers and histories can legally travel.

Possible payer or funding locus. market operator, procurement, seller. A separate commercial budget is not assumed.

Scarcity captured. Task-conditioned performance under identity continuity.

Possible revenue mechanism. membership, data service, verification fee. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. A market complement; not a market and not evidence of liquidity.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. transactions, anti-Sybil controls, portability rights.

Incumbent counterstrategy. platforms keep reputation captive.

Capital / regulation / reversibility. Capital intensity: medium; regulatory exposure: high; reversibility: high.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if markets stay thin or performance is nonstationary and nontransferable. At the same time, the natural enclosure path can platforms keep reputation captive, while access to transactions, anti-Sybil controls, portability rights disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Buyers repeatedly select among independent sellers and histories can legally travel. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. Markets stay thin or performance is nonstationary and nontransferable.

O14. Insurance and risk transfer

Formation condition. Loss classes are bounded, observable, diversifiable, and legally insurable.

Possible payer or funding locus. CFO, risk committee, operator. A separate commercial budget is not assumed.

Scarcity captured. Risk capital and credible compensation.

Possible revenue mechanism. premium, brokerage, reinsurance. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. A service and risk-capital market only when licensed capital bears claims.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. loss data, capital, licensing, claims.

Incumbent counterstrategy. insurers partner with major platforms.

Capital / regulation / reversibility. Capital intensity: high; regulatory exposure: very high; reversibility: medium.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if losses are correlated, nonstationary, catastrophic, or premiums exceed savings. At the same time, the natural enclosure path can insurers partner with major platforms, while access to loss data, capital, licensing, claims disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Loss classes are bounded, observable, diversifiable, and legally insurable. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. Losses are correlated, nonstationary, catastrophic, or premiums exceed savings.

O15. Dispute resolution

Formation condition. Autonomous transactions produce contested authority, delivery, or harm.

Possible payer or funding locus. market operator, payments, enterprises. A separate commercial budget is not assumed.

Scarcity captured. Accepted rulebook, neutral adjudication, remedies.

Possible revenue mechanism. case fee, membership, levy. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. An institution supporting exchange; it can exist without a competitive market.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. evidence, jurisdiction, enforceable awards.

Incumbent counterstrategy. payment networks and courts retain jurisdiction.

Capital / regulation / reversibility. Capital intensity: medium; regulatory exposure: high; reversibility: high.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if disputes stay rare/simple or provider terms force captive arbitration. At the same time, the natural enclosure path can payment networks and courts retain jurisdiction, while access to evidence, jurisdiction, enforceable awards disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Autonomous transactions produce contested authority, delivery, or harm. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. Disputes stay rare/simple or provider terms force captive arbitration.

O16. Settlement and clearing

Formation condition. Multiple venues and counterparties create credit, liquidity, and finality risk.

Possible payer or funding locus. market operators, banks, large buyers. A separate commercial budget is not assumed.

Scarcity captured. Netting, finality, collateral, default waterfall.

Possible revenue mechanism. clearing and membership fees. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. Market infrastructure only after genuine underlying obligations exist.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. genuine underlying trades, law, capital.

Incumbent counterstrategy. financial incumbents control rails.

Capital / regulation / reversibility. Capital intensity: very high; regulatory exposure: very high; reversibility: medium.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if underlying markets never deepen or gross settlement is adequate. At the same time, the natural enclosure path can financial incumbents control rails, while access to genuine underlying trades, law, capital disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Multiple venues and counterparties create credit, liquidity, and finality risk. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the

death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. Underlying markets never deepen or gross settlement is adequate.

O17. Data and content rights

Formation condition. Agents require licensed content, private data, or generated-output rights.

Possible payer or funding locus. publishers, data owners, enterprises, agents. A separate commercial budget is not assumed.

Scarcity captured. Enforceable access and attribution.

Possible revenue mechanism. license, royalty, usage fee. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. Function or capability; no external market is implied.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. rights registry, metering, legal clarity.

Incumbent counterstrategy. platforms negotiate blanket licenses.

Capital / regulation / reversibility. Capital intensity: medium; regulatory exposure: high; reversibility: high.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if law creates broad exceptions or synthetic data eliminates scarcity. At the same time, the natural enclosure path can platforms negotiate blanket licenses, while access to rights registry, metering, legal clarity disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Agents require licensed content, private data, or generated-output rights. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. Law creates broad exceptions or synthetic data eliminates scarcity.

O18. Sovereign infrastructure

Formation condition. Jurisdiction, continuity, or national security outweigh lowest cost.

Possible payer or funding locus. states, regulated enterprises. A separate commercial budget is not assumed.

Scarcity captured. Local control over keys, operators, supply chain, and compute.

Possible revenue mechanism. long-term contracts, infrastructure service. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. Function or capability; no external market is implied.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. capital, procurement, domestic capability.

Incumbent counterstrategy. hyperscalers offer sovereignty wrappers.

Capital / regulation / reversibility. Capital intensity: very high; regulatory exposure: very high; reversibility: high.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if buyers accept global safeguards or sovereignty premium becomes politically untenable. At the same time, the natural enclosure path can hyperscalers offer sovereignty wrappers, while access to capital, procurement, domestic capability disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Jurisdiction, continuity, or national security outweigh lowest cost. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. Buyers accept global safeguards or sovereignty premium becomes politically untenable.

O19. Compute, power, and capacity

Formation condition. Physical demand exceeds supply or locality matters.

Possible payer or funding locus. clouds, enterprises, states, agent platforms. A separate commercial budget is not assumed.

Scarcity captured. Power, chips, interconnect, financing, siting.

Possible revenue mechanism. capacity contract, optimization fee, financing spread. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. Function or capability; no external market is implied.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. infrastructure and metering.

Incumbent counterstrategy. hyperscalers vertically integrate.

Capital / regulation / reversibility. Capital intensity: very high; regulatory exposure: high; reversibility: medium.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if efficiency and demand saturation create persistent slack. At the same time, the natural enclosure path can hyperscalers vertically integrate, while access to infrastructure and metering disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Physical demand exceeds supply or locality matters. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. Efficiency and demand saturation create persistent slack.

O20. Physical oracles and actuation

Formation condition. Agents affect factories, logistics, laboratories, and property.

Possible payer or funding locus. operators, insurers, regulators. A separate commercial budget is not assumed.

Scarcity captured. Trusted sensing, machine state, and reversible actuation.

Possible revenue mechanism. equipment, service, per verified event. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. Function or capability; no external market is implied.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. hardware, certification, maintenance.

Incumbent counterstrategy. industrial incumbents bundle.

Capital / regulation / reversibility. Capital intensity: high; regulatory exposure: high; reversibility: high.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if embodied autonomy stalls or sensors become universal commodities. At the same time, the natural enclosure path can industrial incumbents bundle, while access to hardware, certification, maintenance disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Agents affect factories, logistics, laboratories, and property. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. Embodied autonomy stalls or sensors become universal commodities.

O21. Enterprise transformation

Formation condition. Capability exists but organizations cannot redesign coordinated work.

Possible payer or funding locus. CEO, COO, HR, CIO. A separate commercial budget is not assumed.

Scarcity captured. Political coalition, process redesign, training, accountability.

Possible revenue mechanism. services, managed operations, software-assisted program. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. Function or capability; no external market is implied.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. executive mandate and domain access.

Incumbent counterstrategy. consultancies and suites dominate.

Capital / regulation / reversibility. Capital intensity: low; regulatory exposure: medium; reversibility: medium.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if ai remains individual augmentation or si directly redesigns organizations. At the same time, the natural enclosure path can consultancies and suites dominate, while access to executive mandate and domain access disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Capability exists but organizations cannot redesign coordinated work. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. AI remains individual augmentation or SI directly redesigns organizations.

O22. Systems-of-record distribution

Formation condition. Agents need authoritative data and transaction endpoints.

Possible payer or funding locus. application vendors and domain platforms. A separate commercial budget is not assumed.

Scarcity captured. Installed workflow, data gravity, legal record.

Possible revenue mechanism. seat, consumption, take rate. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. Function or capability; no external market is implied.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. incumbent base and APIs.

Incumbent counterstrategy. incumbents already own it.

Capital / regulation / reversibility. Capital intensity: high; regulatory exposure: medium; reversibility: high.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if interface-obsolete si reduces direct app value and mandates portability. At the same time, the natural enclosure path can incumbents already own it, while access to incumbent base and APIs disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Agents need authoritative data and transaction endpoints. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. Interface-obsolete SI reduces direct app value and mandates portability.

O23. Certification and accreditation

Formation condition. Regulators/procurement accept standardized proof as an eligibility gate.

Possible payer or funding locus. regulator, auditor, procurement. A separate commercial budget is not assumed.

Scarcity captured. Institutional recognition.

Possible revenue mechanism. assessment, renewal, registry fee. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. Function or capability; no external market is implied.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. standards, competent bodies, independence.

Incumbent counterstrategy. large auditors and vendors shape standards.

Capital / regulation / reversibility. Capital intensity: medium; regulatory exposure: high; reversibility: medium.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if certifications become checkbox theater or continuous proof replaces periodic review. At the same time, the natural enclosure path can large auditors and vendors shape standards, while access to standards, competent bodies, independence disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Regulators/procurement accept standardized proof as an eligibility gate. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the

death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. Certifications become checkbox theater or continuous proof replaces periodic review.

O24. Protocol infrastructure

Formation condition. Many parties need interoperable discovery, authority, evidence, and payment messages.

Possible payer or funding locus. developers, platforms, standards bodies. A separate commercial budget is not assumed.

Scarcity captured. Registry governance, conformance operations, portable trust roots, and full-workflow switching - only if independently controlled.

Possible revenue mechanism. hosting, conformance, enterprise support. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. Interoperability infrastructure; common syntax alone is neither scarcity nor contestability.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. adoption, governance, extensions.

Incumbent counterstrategy. incumbents support open core but capture services.

Capital / regulation / reversibility. Capital intensity: low; regulatory exposure: medium; reversibility: high.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if protocols remain message formats without production transactions. At the same time, the natural enclosure path can incumbents support open core but capture services, while access to adoption, governance, extensions disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Many parties need interoperable discovery, authority, evidence, and payment messages. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. Protocols remain message formats without production transactions.

O25. Human consent and fiduciary control

Formation condition. Institutions refuse to treat machine optimization as legitimate authority.

Possible payer or funding locus. consumers, boards, professions, state. A separate commercial budget is not assumed.

Scarcity captured. Meaningful consent, judgment, accountable natural/legal person.

Possible revenue mechanism. service, compliance, public infrastructure. This describes who might pay and on what meter; it does not establish willingness to pay or durability.

Formation verdict. Function or capability; no external market is implied.

Seven-gate evidence status. Functional necessity is conditional; measurable loss, decision authority, separate budget, independent access, sustainable rent, and sufficient time remain unproven unless scenario-specific evidence establishes them.

Dependencies. law, user experience, professional duty.

Incumbent counterstrategy. providers turn consent into captive interface.

Capital / regulation / reversibility. Capital intensity: medium; regulatory exposure: very high; reversibility: high.

Correlated-death cluster. The function, independent access, and sustainable rent can fail together if trusted si receives direct institutional standing or consent becomes nominal. At the same time, the natural enclosure path can providers turn consent into captive interface, while access to law, user experience, professional duty disappears or becomes uneconomic.

Decisive falsifier. In a mature deployment, the stated funding locus does not assign a separate recurring budget after the natural incumbent alternative is fully configured, or an independent supplier cannot obtain the required data, authority, and ground truth at sustainable full cost.

Revival condition. Institutions refuse to treat machine optimization as legitimate authority. In addition, a funding principal must recognize a measurable loss without the function, grant independent access, renew a separate budget, and leave enough contestability for retained rent.

Option premium, expiry, and abandonment. The hypothetical premium is the cost of keeping interfaces, rights, evidence, qualifications, and switching paths live. The option expires as the incumbent enclosure becomes institutionally accepted, the death condition materializes, or procurement and certification take longer than the market window. The abandonment signal is simultaneous failure of budget, access, rent, and time gates; continued technical usefulness does not keep the economic option alive.

Organizational home remains open. This function may be internalized, bundled, supplied as a public or cooperative rail, regulated, or left without a standalone provider.

Death condition. Trusted SI receives direct institutional standing or consent becomes nominal.

Revenue mechanisms by market stage

Revenue is an observation about who pays now, not proof of long-run defensibility.

Stage	Payer	Possible meter	What revenue would establish	What it would not establish
Copilot / pilot	team, CIO, innovation	seat, project, services	interest and integration willingness	captured P&L value or autonomous authority
Agent fleet	CIO, CISO, function	identity, controlled action, subscription, consumption	production control budget	independent category durability
Delegated external action	payments, risk, legal, application owner	authorized action, attestation, transaction	authority and evidence are funded	open market or neutrality
Machine procurement	buyer or supplier	membership, RFQ, transaction	real exchange if bids and settlement are external	liquidity or venture-scale take rate
Assurance / insurance	risk owner	audit fee, premium, limit	loss transfer if capital and claims are real	diversifiability under superintelligence
Clearing	members and venues	membership, clearing, collateral service	underlying obligations and finality demand	that an underlying market should exist
Sovereign / physical	state, regulated firm, operator	long-term capacity, infrastructure, certification	scarce resource and jurisdictional budget	global scalability or low capital intensity

Opportunities that disappear as intelligence improves

Basic prompt wrappers, undifferentiated model comparison, manual configuration, simple routing, generic evaluation, elementary code generation, superficial observability, and low-value integration can be absorbed by models or platforms. If verification and execution become nearly free, ex ante selection can be replaced by run-all-and-select. If a trusted provider owns identity, context, tools, and institutional legitimacy, many horizontal control functions collapse into one operating layer.

Disappearance can occur before technical redundancy: an incumbent can bundle the function at zero apparent price, a regulator can make it an obligation, a system of record can own the only authoritative state, or customers can internalize it. Technical value and company value must remain separate.

Functions strengthened by consequence-adjusted autonomous exposure

The relevant variable is not intelligence level. It is autonomous consequence exposure after native prevention and reversibility. More capable agents can increase the value of constraints when they receive broader credentials, actions become harder to reverse, consequence volume rises, physical systems enter scope, and institutions demand a reachable defendant. They can also reduce error, improve native controls, lower incident frequency, or concentrate legitimate delegation in a trusted provider.

But the strengthening is conditional on institutions refusing to delegate the same function to the intelligent provider. A trusted self-verifying utility can absorb verification and policy. A state can nationalize the chokepoint. A platform can own the user's identity and redress. The analysis therefore tracks recognition and control, not abstract importance.

Cross-world robustness properties

No category is robust to every world. The relatively contradiction-resistant mechanisms are not products but properties: portability across providers and jurisdictions; authority scoped to reversible actions; evidence that can be checked by more than one institution; customer control over keys and policy; dependency transparency; physical or legal enforcement outside a model's unilateral control; and commitments that can be reduced as uncertainty changes.

These properties can appear inside incumbents, public infrastructure, internal enterprise systems, standards, or independent companies. They can be socially valuable while imposing cost and creating no supplier, separate budget, or rent. Their persistence does not answer who captures the value.


Capital regimes and investor appetite

Investor appetite changes the equilibrium and can create false evidence. In a compute boom, capital rewards infrastructure, frontier suppliers, and growth before unit economics. In an agent-platform boom, distribution and developer adoption receive premium valuations. In a post-bubble discipline regime, buyers and investors demand captured P&L, retention, and low services intensity. In a regulated-infrastructure regime, licenses, balance sheets, long contracts, and institutional credibility matter. In a sovereign-capex regime, domestic operations and political alignment dominate scalability.

Venture financing can subsidize both sides of a supposed market, mask liquidity failure, or fund zero-price bundling. Acquisition can be an attractive exit while proving the independent layer was not durable. Capital availability is therefore an endogenous state variable, not external validation.

Leading-indicator system

Indicators are useful only when tied to a branch. A rise in agents confirms proliferation but says little about market formation. A rise in machine-payment messages confirms protocol use but not delegated authority or final settlement. The dashboard therefore pairs each observation with the world it discriminates.



ODYLITH / MARKET EVOLUTION SITUATIONAL AWARENESS / 5 AUGUST 2026

Leading indicators by causal family

Family	Measures
Cost / capability	normalized price by capability tier; task horizon; open gap; candidate multiplicity
Total economics	risk-adjusted spend; cleanup labor; P&L capture; mature-cohort retention
Concentration	ownership and dependency HHI; first-party routing share; committed-spend concentration
Enterprise authority	service identities; delegated write/payment scope; reversals; human override quality
Budget	shift from pilots/seats to function P&L, operations, labor, assurance, capex
Incumbent enclosure	bundled price; acquisitions; default placement; portable policy/evidence
Protocols	production multi-homing; successful switching; portable credentials and reputation
Markets	independent qualified bids; net external settled value; fill; depth; bypass; disputes
Insurance	limits; exclusions; capital; reinsurance; paid claims; correlated-loss tests
Institutions	court/regulator acceptance; liability assignment; sovereignty procurement; enforcement
Labor	junior hiring; task composition; expert/novice gains; worker consultation; redesign
Physical	data-center TWh; grid queue; capacity premium; robot deployment; safety claims

Family	Measures
Superintelligence	capability discontinuity; evaluation gaming; emergency controls; institutional trust
Welfare-adjusted demand	human or enterprise value net of compute, review, externalities, retries, and synthetic workload
Capacity correction	utilization, cancellations, write-downs, training/inference split, and stranded infrastructure
Public and worker veto	consultation, bargaining, litigation, licensing, moratoria, elections, and community opposition
Institutional response time	capability-change interval divided by detection, rule, deployment, and enforcement time
Non-compensable loss	harms without reachable defendant, adequate capital, accepted jurisdiction, or practical remedy
Full-workflow switching	movement of identity, context, policy, evidence, contracts, and operations - not API compatibility

Research and experiment agenda

The following investigations reduce uncertainty without choosing a product:

1. Build an ownership- and dependency-adjusted map of production agent supply.
2. Separate seats, active users, autonomous actions, accepted outcomes, and captured P&L across enterprise cohorts.
3. Measure whether enterprises actually switch providers under outage, price, capability, and jurisdiction shocks.
4. Test whether provider-native evidence is accepted by customers, auditors, insurers, courts, and regulators.
5. Observe delegated authority: scope, revocation, human confirmation, disputes, and liability allocation.
6. Distinguish protocol messages from binding contracts and net settled external value.
7. Measure task contractibility before evaluating auctions or outcome markets.
8. Stress supplier independence by common model, cloud, data, verifier, and parent.
9. Compare routing against strong fixed policies and run-all/select under full cost.
10. Track incumbent bundling, acquisition, default placement, and open-standard extensions.
11. Study failure tails, delayed harms, common-mode incidents, and insurer exclusions.
12. Run institutional tabletop exercises for trusted, untrusted, deceptive, and discontinuous superintelligence.
13. Test whether an incident increases real options or merely increases dependence on the incumbent.
14. Update the scenario ledger quarterly without converting scenarios into point forecasts.

Evidence and confidence ledger

Claim	Status	Confidence	What would change it
Per-unit capability cost is declining rapidly	observed but tier-dependent	high	sustained real-price reversal
Total AI system spend can rise while unit cost falls	observed / causal inference	high	mature demand saturation
Enterprises will operate many agents	announced / emerging	medium	suite consolidation into one delegate
Provider plurality persists	uncertain	low-medium	dependency-adjusted production concentration
Agent activity becomes an open market	speculation	low	binding independent bids and net settlement
Authority and liability remain scarce under SI	scenario-dependent	low-medium	institutional standing of trusted SI
Independent verification earns a separate budget	uncertain	low-medium	procurement and regulatory acceptance
Open standards increase contestability	uncertain	low	portable production multi-homing
Insurance can absorb advanced-agent loss	unknown	low	credible limits, capital, claims, and tail calibration
Clearing becomes necessary	remote contingent scenario	low	multi-venue obligations and finality risk
Physical and sovereign scarcities persist	plausible / observed today	medium-high	abundance and accepted global safeguards
Any one independent company captures these functions	unresolved	low	buyer budgets, enclosure, retention, and switching
The taxonomy is independent of the originating thesis	partially tested	low-medium	actor/welfare taxonomy reveals no material omitted worlds
Governance remains enforceable under superintelligence	scenario-dependent	low	institutional response, coercive capacity, key custody, and legal continuity
Candidate mechanisms are antifragile	falsification architecture only	low	bounded shocks produce retained out-of-sample gains with protected downside
Standards produce contestability	unresolved	low	full-workflow switching and distributed traffic/rent across the contestability stack
Machine activity forms markets	rare conditional transition	low	independent executable bids, net settlement, default, remedy, and repeat exchange
Opportunity surfaces receive separate recurring budgets	unproven	low	renewal after native alternatives are fully configured

Adversarial review architecture

The dossier was built through three separated rounds. Round 1 asked independent teams to derive market structures, enterprise authority paths, machine-market formation tests, incumbent enclosure strategies, superintelligence phase space, and antifragility mechanisms before comparing with the Odylith deck. Round 2 cross-examined contradictions: abundance versus concentration, open standards versus enclosure, delegation versus legal accountability, routing versus run-all, verification versus self-proof, neutrality versus conflicts, and resilience versus actual antifragility. Round 3 asked an independent adjudicator to identify agreements, contradictions, missing worlds, and variables without selecting a product.

The appendices preserve those disagreements. They are not votes and do not create certainty. Their purpose is to reduce shared framing error.

Round 3 correction-closure ledger

Adjudicated defect	Material integration	Residual uncertainty
Premature three-round claim	Round 3 is now included after six Round 2 records and this closure ledger	The adjudication remains a historical critique, not proof of truth
Odylith-shaped taxonomy	Added W1-W18 actor/welfare atlas derived from households, labor, public goods, coercion, externalities, and non-market allocation	Frame independence remains partial because every taxonomy reflects its question
Scenario list called a phase space	Renamed as scenario library; added 16-cell crossed lattice and nine omitted equilibria	The phase space is still not exhaustive or probability weighted
Rebound over-attribution	Separated observed co-movement from elasticity, overbuild, utilization, training, and write-down hypotheses	Mature-cohort causal evidence remains limited
Scarcity always migrates	Added no-migration endpoint: demand, profit, or willingness to pay can disappear	Which tasks vanish is unknown
Standards treated as contestability	Added full contestability stack and removed common syntax as an assumed scarcity	Production switching evidence remains early
Market language leakage	Added formation verdicts distinguishing hierarchy, supply chain, directory, procurement, market, and infrastructure	Border cases remain institution-specific
Intelligence level treated as causal	Reframed around consequence-adjusted autonomous exposure after native prevention and reversibility	Exposure distributions remain partly hidden
Robustness called opportunity	Renamed as cross-world robustness properties with no implied supplier, budget, or rent	Social value and private capture remain separate
Opportunity menu salience	Each O1-O25 now includes seven gates, correlated death, falsifier, revival, option premium, expiry, abandonment, public/internal homes, and no-budget outcomes	All evidence states remain low until observed
Authority assumed durable	Added institutional collapse, machine standing, emergency law, key-holder capture, and response-time failure	Enforcement under superintelligence is unknowable
Loose antifragility	Imported protected-downside, retained-rights, measurable-gain, convexity-interval, and bounded-failure tests	No mechanism is claimed empirically antifragile
Narrow indicators	Added welfare-adjusted demand, utilization/write-downs, veto, institutional speed, non-compensable loss, and full-workflow switching	Many measures require unavailable private data
Confidence ledger omissions	Added frame independence, governance survival, antifragility, contestability, market formation, and separate-budget claims at low confidence	Confidence must be updated from future evidence

Round 2 cross-examination: market economics

Scope and method

Evidence cutoff: 5 August 2026.

This record cross-examines seven paired claims in `output/pdf/Odylith_Strategic_Dossier.md`. It does not select a strategy, product, market entry, investment, or company design. Each debate is resolved only to a conditional causal model.

Labels are used strictly:

- Fact: a dated observation supported by a primary or authoritative source.
- Inference: a causal interpretation consistent with multiple facts but not directly observed.
- Speculation: a coherent future state whose probability is unknown.

The unit of analysis is the quality-adjusted accepted task, not a token, model call, benchmark point, or nominal software seat. The dossier is strongest when it separates cognition from institutional authority and weakest when it treats any current complement as permanently scarce. Routing, governance, verification, standards, neutrality, and authority are all endogenous: providers can bundle them, buyers can internalize them, standards can commoditize them, and regulation can require or reshape them.

Empirical anchor

Several facts can support opposite market conclusions at once.

1. Fact: Stanford's 2025 AI Index reported that the inference cost of a GPT-3.5-level system fell more than 280-fold from November 2022 to October 2024, while hardware cost fell about 30 percent annually and energy efficiency improved about 40 percent annually. It also reported a narrowing open-versus-closed performance gap on selected benchmarks. [Stanford HAI, 7 April 2025](<https://hai.stanford.edu/ai-index/2025-ai-index-report>)
2. Fact: Alphabet stated that Gemini serving unit cost fell 78 percent during 2025. In the same earnings call it reported 2025 capital expenditure of \$91.4 billion and forecast \$175 billion to \$185 billion for 2026, with approximately 60 percent of 2025 capital expenditure in servers and 40 percent in data centers and networking. [Alphabet Q4 2025 earnings call, February 2026](<https://abc.xyz/investor/events/event-details/2026/2025-Q4-Earnings-...>)
3. Fact: The IEA estimated that energy used per AI task had recently fallen by at least an order of magnitude annually, while global data-center electricity use rose 17 percent in 2025 and AI-focused data-center electricity use rose 50 percent. Its central projection rises from about 485 TWh in 2025 to 950 TWh in 2030. [IEA, 2026](<https://www.iea.org/reports/key-questions-on-energy-and-ai/executive...>)
4. Fact: NVIDIA reported fiscal 2026 revenue of \$215.9 billion, a 71.1 percent gross margin, and 68 percent growth in Data Center revenue. Two direct customers represented 22 percent and 14 percent of revenue, and NVIDIA disclosed manufacturing lead times beyond twelve months for some products. [NVIDIA fiscal 2026 10-K](<https://www.sec.gov/Archives/edgar/data/1045810/000104581026000021/n...>)
5. Fact: Microsoft reported fiscal 2025 property-and-equipment additions of \$64.6 billion, compared with \$44.5 billion in 2024, plus \$103.9 billion of purchase commitments primarily related to data centers. [Microsoft 2025 Annual Report](<https://www.microsoft.com/investor/reports/ar25/index.html>)
6. Fact: In January 2025 Google included Gemini in Workspace Business and Enterprise plans rather than requiring a separate add-on. Its worked example moved from \$32 per user per month with the add-on to \$14, compared with \$12 previously without Gemini. [Google Workspace, 15 January 2025](<https://workspace.google.com/blog/product-announcements/empowering-b...>)
7. Fact: Microsoft reported nearly 140,000 organizations using GitHub Copilot in fiscal 2026 Q3, enterprise subscribers nearly tripling year over year, and a majority of Copilot users using multiple models. [Microsoft FY2026 Q3 earnings, April 2026](<https://www.microsoft.com/en-us/investor/events/fy-2026/earnings-fy-...>)

8. Fact: GitHub made its enterprise agent control plane generally available and exposed Claude and Codex within Copilot workflows in February 2026. [GitHub control plane, 26 February 2026](<https://github.blog/changelog/2026-02-26-enterprise-ai-controls-agen...>), [GitHub multi-agent availability, 26 February 2026](<https://github.blog/changelog/2026-02-26-claude-and-codex-now-availa...>)
9. Fact: AWS Bedrock and Microsoft Foundry offer native model routers. AWS routes within model families; Microsoft offers quality, cost, and balanced modes plus failover. [AWS documentation, accessed 3 August 2026](<https://docs.aws.amazon.com/bedrock/latest/userguide/prompt-routing.html>), [Microsoft documentation, updated 2026](<https://learn.microsoft.com/en-us/azure/foundry/openai/concepts/mode...>)
10. Fact: AWS, Anthropic, Google, Microsoft, OpenAI, and other firms backed the Linux Foundation's Agentic AI Foundation around MCP, AGENTS.md, and related projects. [Linux Foundation, 9 December 2025](<https://www.linuxfoundation.org/press/linux-foundation-announces-the...>)
11. Fact: The UK CMA's cloud investigation found significant market power at AWS and Microsoft and identified egress and interoperability barriers. The EU Data Act requires switching and egress charges to disappear from 12 January 2027. [UK CMA final investigation, July 2025](<https://www.gov.uk/cma-cases/cloud-services-market-investigation>), [EU Data Act explanation](<https://digital-strategy.ec.europa.eu/en/factpages/data-act-explained>)
12. Fact: The U.S. Census Bureau found overall reported business AI use between 17 and 20 percent from December 2025 through May 2026, but 37 percent for firms with at least 250 employees. The survey broadened its AI question in November 2025. [U.S. Census Bureau, 26 May 2026](<https://www.census.gov/library/stories/2026/05/ai-use-businesses.html>)
13. Fact: NIST launched an AI Agent Standards Initiative and separately highlighted identity, authority, monitoring, and evaluation-gaming issues for software agents. [NIST initiative, updated 20 April 2026](<https://www.nist.gov/artificial-intelligence/ai-agent-standards-init...>), [NIST identity and authority paper, February 2026](<https://www.nist.gov/news-events/news/2026/02/new-concept-paper-iden...>), [NIST evaluation gaming, updated 2 December 2025](<https://www.nist.gov/caisi/cheating-ai-agent-evaluations>)

These facts do not select one world. Price convergence can coexist with capacity concentration. Provider plurality can coexist with platform consolidation. Open protocols can increase distribution while moving rents elsewhere. The following debates identify the conditions under which each side wins.

Debate 1: abundant intelligence versus concentrated intelligence

Steelman A: intelligence becomes abundant

Fact: Quality-adjusted inference prices have fallen rapidly, small models have improved, and open systems have narrowed selected benchmark gaps. Multiple model families are available through clouds and application platforms.

Inference: For a large and expanding set of routine cognitive tasks, the supply curve shifts outward faster than demand. Adequate models become interchangeable after adjustment for latency, context, and tool use. The price of baseline cognition approaches the cost of compute plus a thin service margin. Model-level rents compress because buyers can substitute smaller models, open weights, local execution, caching, distillation, and batch processing.

Speculation: A mature enterprise does not purchase a named intelligence. It purchases a service grade such as accepted-task probability, latency band, data locality, and capacity reservation. Model identity becomes an implementation detail for most work.

Steelman B: intelligence remains concentrated

Fact: Frontier development and deployment require large fixed commitments to chips, data centers, networking, power, talent, and software. NVIDIA's revenue, margins, customer concentration, and long lead times demonstrate current upstream scarcity. Hyperscaler capital spending demonstrates that the supply response is balance-sheet intensive and delayed.

Inference: Abundance at yesterday's quality level does not imply abundance at the frontier or at enterprise-grade reliability. A small number of firms can own the best models, scarce accelerators, lowest serving cost, proprietary distribution, and the operational data needed to improve. Lower marginal cost may strengthen concentration because the largest provider can price below smaller rivals, bundle the service, and finance the next fixed-cost cycle.

Speculation: The model market resembles a ladder. Each old rung commoditizes, but the commercially valuable frontier remains concentrated. Buyers consume abundant low-tier intelligence and concentrated high-tier intelligence simultaneously.

Decisive state variables

- Quality gap between the best and the cheapest adequate model on production tasks.
- Fixed frontier cost relative to revenue available to a new entrant.
- Open-model production share, not download count or benchmark rank.
- Availability and price of accelerators, HBM, networking, and firm power.
- Buyer willingness to pay for reliability, indemnity, residency, and support.
- Cross-provider switching time and failure rate.
- Capital-market willingness to subsidize frontier entrants.
- Strength of distribution and bundle cross-subsidies.

Hidden confounders

- A benchmark-adjusted price decline can reflect a changing benchmark rather than production substitutability.
- Posted API prices omit negotiated discounts, capacity commitments, cache economics, and internal transfer prices.
- Open weights can still be served mainly by concentrated clouds.
- Frontier provider revenue may reflect temporary scarcity rent rather than a permanent structure.
- Capital expenditure includes non-AI cloud workloads.
- Quality, reliability, and organizational trust are often bundled in one observed purchase.

Opposite-world reconstruction

The abundance fact pattern can produce concentration: cheap models expand usage, usage creates data-center demand, the largest platforms obtain the best utilization, their unit costs fall fastest, and they bundle adequate intelligence into installed software. Conversely, the concentration fact pattern can produce abundance: high margins attract capital and custom silicon, overbuilding creates excess capacity, open releases diffuse techniques, and a price war commoditizes inference.

Discriminating evidence

Track quality-adjusted production switching, the share of accepted tasks served by open or small models, provider gross margin by layer, multi-cloud migration duration, utilization and write-downs of AI infrastructure, and the persistence of frontier quality premia after one full model generation. A utility world is supported if service grades become portable and model premia shrink. A concentrated world is supported if quality and reliability premia persist while supply and distribution shares remain stable.

Opportunity categories that survive either side

Capacity and power economics, workload qualification, identity and access control, integration with authoritative systems, incident recovery, and measurement of accepted-task cost remain relevant. In abundance they manage volume and heterogeneity; in concentration they manage dependency, scarcity, and bargaining exposure.

Conditional resolution

Abundance describes marginal cognition at a moving quality threshold. Concentration describes ownership of the frontier and its complements. Both can be true by layer and time. The dossier's claim fails only if it treats either state as universal.

Debate 2: near-zero marginal cost versus rising total expenditure

Steelman A: marginal cost approaches zero

Fact: Provider serving costs and benchmark-adjusted inference prices have fallen sharply. Caching, batching, smaller models, custom silicon, and better utilization continue to reduce cost per token and per bounded task.

Inference: For tasks that can be specified once and replicated many times, the incremental cognitive cost becomes negligible relative to labor, software, and capital. Competition and bundling pass much of the saving to customers. A fixed organizational budget can deliver far more cognitive work.

Speculation: Routine drafting, classification, translation, code transformation, and first-pass analysis become unmetered features. Direct model spend disappears as a visible budget line.

Steelman B: total expenditure rises

Fact: Alphabet reported a 78 percent serving-unit cost reduction during a year of rapidly rising infrastructure investment. The IEA reported falling energy per task alongside sharply rising AI-focused data-center electricity use.

Inference: Rebound occurs when task volume, trajectory depth, modality, tools, retries, evaluation, and automation scope grow faster than unit cost falls. The relevant equation is:

total accepted-outcome cost = fixed model and infrastructure cost + tokens + tools + context + retries + review + integration + security + liability + organizational change

A token-price decline can reduce direct inference cost while raising expenditure on data centers, power, software integration, security, human supervision, and physical execution. Agentic work deepens the effect because a single business request can generate many planning, search, tool, validation, and repair steps.

Speculation: Intelligence becomes similar to electricity: dramatic efficiency improvement expands the number of uses enough that aggregate consumption and complementary capital continue to grow.

Decisive state variables

- Price elasticity of task demand by task class.
- Calls, tokens, tool invocations, and retries per accepted outcome.
- Share of workflows that become newly economical after a price decline.
- Human review and repair time.
- Fixed capacity investment and utilization.
- Energy per task versus growth in task volume.
- Budget caps and the share of savings retained rather than reinvested.
- Migration from text to video, reasoning, continuous agents, and robotics.

Hidden confounders

- Token counts are not comparable across models or accepted outcomes.
- Provider usage metrics may count calls, interactions, assists, or users differently.
- Data-center electricity includes non-AI workloads.
- Reported cost reductions can reflect utilization or product mix.
- Higher spend may be financed by displacement of other software or labor rather than net new expenditure.
- Failed tasks and hidden human rescue may be missing from telemetry.

Opposite-world reconstruction

Efficiency can defeat rebound. Better planning, longer context, distillation, and local models may reduce attempts per outcome so quickly that total compute falls. Conversely, capped direct AI budgets can coexist with rising economy-wide cost if power, integration, and liability are charged elsewhere.

Discriminating evidence

Measure complete accepted-outcome cost before and after model upgrades, including human minutes, retries, tools, delay, and failure loss. Compare growth in accepted outcomes with the decline in energy and compute per outcome. Rebound is supported when quantity growth exceeds efficiency improvement. Cost collapse is supported when total organizational cost per business result and aggregate resource use both fall.

Opportunity categories that survive either side

Metering accepted outcomes, capacity planning, workflow redesign, integration, security, and cost attribution survive. Under zero marginal cost they allocate attention and prevent waste; under rising expenditure they manage real resource constraints.

Conditional resolution

Near-zero marginal cost and rising total expenditure are not contradictory. They refer to different denominators. The contradiction becomes real only when direct token price is presented as total economic cost.

Debate 3: provider plurality versus platform consolidation

Steelman A: plurality persists

Fact: GitHub exposes multiple agent providers, Microsoft reports multi-model Copilot use, clouds list multiple model families, and sovereign and open systems remain active. Buyers differ in jurisdiction, latency, privacy, cost, modality, and domain requirements.

Inference: No provider is best on every task and constraint. Enterprises retain portfolios for resilience and negotiating leverage. Regulatory fragmentation, export controls, local-language demand, and sovereign investment create additional suppliers. Specialization can persist even if baseline models converge.

Speculation: The market resembles asset management: portfolios contain frontier, efficient, open, local, and domain systems, with policy determining eligibility rather than one global winner.

Steelman B: platforms consolidate demand

Fact: GitHub, Microsoft, AWS, Google, ServiceNow, Salesforce, and IBM bundle agents with identity, context, policy, observability, billing, and workflow. Google and Microsoft have folded AI into established subscription bundles. Cloud switching barriers remain material.

Inference: Model plurality can increase while customer ownership consolidates. One platform can expose several model suppliers behind a single procurement, identity, data, and billing boundary. The platform captures distribution and usage data while model providers become components. Enterprise buyers may accept this arrangement because operational simplicity and liability allocation outweigh theoretical supplier independence.

Speculation: Three or four operating layers dominate enterprise demand, each presenting an internal model portfolio. Nominal provider plurality hides consolidated order flow.

Decisive state variables

- Share of model consumption purchased through aggregating platforms.
- Customer ownership of identity, memory, policy, telemetry, and billing.
- Effective switching cost across full workflows, not API compatibility alone.
- Incidence of bundle pricing and committed-spend discounts.
- Number of legally and technically independent execution routes.
- Provider concentration at model, cloud, application, and chip layers.
- Regulatory limits on exclusivity, tying, and self-preferencing.

Hidden confounders

- Counting model names exaggerates independent supply when models share cloud, chips, tools, or verifiers.
- A multi-model interface may increase platform lock-in.
- Enterprise procurement records can hide internal cross-subsidy.
- Revenue share at one layer does not reveal control at another.
- Switching demonstrations on toy workloads understate data and process migration.

Opposite-world reconstruction

Plurality can accelerate consolidation because enterprises need an aggregator to manage it. Consolidation can preserve plurality because dominant platforms benefit from listing rival models and earning infrastructure revenue. Antitrust can open distribution while leaving infrastructure concentration untouched.

Discriminating evidence

Measure where contracts, identity, telemetry, memory, and payment reside; how often production workloads move between platforms; the concentration-adjusted dependency graph across model, cloud, chip, tool, and verifier; and the share of AI spend inside preexisting bundles. Platform consolidation is supported when customer relationships and action rights concentrate despite model variety. Plurality is supported when buyers can move complete workflows and retain bargaining leverage.

Opportunity categories that survive either side

Dependency mapping, portability, cross-system integration, procurement analytics, identity, incident response, and sovereign compliance survive. Plurality increases coordination need; consolidation increases common-mode and bargaining risk.

Conditional resolution

Provider count and platform concentration are orthogonal. The key economic question is who owns qualified demand, authoritative context, and the right to act.

Debate 4: open standards as distribution versus commoditization

Steelman A: standards expand distribution

Fact: Major competitors jointly support MCP, AGENTS.md, A2A, and agent standards work. The EU Data Act and cloud competition actions seek lower switching friction.

Inference: A common interface reduces integration cost, expands the addressable ecosystem, and lets a component reach customers it could not economically integrate one by one. Standards can create a larger market and accelerate complementary innovation. A participant can still differentiate on reliability, operations, security, data, support, and performance.

Speculation: An open agent protocol becomes analogous to HTTP or SQL: broadly implemented, economically generative, and compatible with profitable services above and below it.

Steelman B: standards commoditize the layer

Inference: Once interfaces, schemas, and receipts are standardized, buyers can substitute implementations. Connectors and basic routing lose pricing power. Platforms can implement the standard as a bundled feature, then move rent to hosting, identity, proprietary extensions, distribution, or transaction fees. A neutral standard can therefore increase market size while destroying standalone margin in the standardized function.

Speculation: MCP-compatible tool access becomes table stakes. The economic control point shifts to permission, discovery, reputation, execution capacity, or the platform that owns the user.

Decisive state variables

- Completeness of the standard versus need for proprietary extensions.
- Certification and conformance quality.
- Switching cost after implementation, including data and operating procedure.
- Economies of scale in hosting, security, and discovery.
- Ownership of identity, namespace, registry, and default distribution.
- Rate at which incumbents bundle compliant implementations.
- Whether standards expose enough telemetry for independent service quality.

Hidden confounders

- Membership in a standards body does not imply interoperable production behavior.
- Nominally open specifications can be controlled by dominant implementers.
- Compatibility at the message layer does not create semantic or policy compatibility.
- Open-source adoption can be service-provider distribution rather than decentralization.
- Regulatory removal of monetary switching fees does not eliminate organizational migration cost.

Opposite-world reconstruction

Commoditization can strengthen a specialist if the standard makes trust, conformance, and operations newly scarce. Distribution can weaken a specialist if incumbents implement the same interface at zero incremental price. The same standard can commoditize one layer and concentrate the next.

Discriminating evidence

Observe implementation switching without custom work, the share of deployments using proprietary extensions, conformance failure rates, registry concentration, attach rates for premium services, and price dispersion among compliant implementations. Distribution wins if new entrants gain production share and differentiated services retain margins. Commoditization wins if the function becomes bundled and price converges toward zero.

Opportunity categories that survive either side

Conformance testing, security, identity, semantic mapping, lifecycle operations, migration, and cross-standard observability survive. Their economics depend on whether they remain scarce or are themselves standardized.

Conditional resolution

Open standards are simultaneously distribution mechanisms and commoditization mechanisms. They expand the system while relocating rents. No standardized function should be presumed to retain pricing power.

Debate 5: routing versus run-everything-and-select

Steelman A: ex-ante routing dominates

Fact: AWS and Microsoft sell native routing intended to optimize cost and quality. Price and capability dispersion remain material.

Inference: When task features predict route performance and execution cost is meaningful, selecting one eligible route before execution avoids redundant compute and delay. Routing becomes more valuable as the number of candidate models and agents grows, provided that predictions are calibrated and route differences are causally real.

Speculation: Stable task bands support policies that send routine work to efficient systems, complex work to frontier systems, and high-risk work to routes with stronger evidence and human review.

Steelman B: run everything and select dominates

Inference: If inference becomes cheap relative to outcome value, prediction error can cost more than redundant execution.

Parallel trajectories create diversity, allow cross-checking, and reveal disagreement. For high-value tasks, selecting the best completed outcome can dominate selecting the predicted best route. The approach is especially attractive when task difficulty is hard to observe ex ante and failure is asymmetric.

Speculation: Advanced systems generate several candidate plans, proofs, or implementations and use deterministic tests, external state, or human judgment to select after execution.

Decisive state variables

- Execution cost relative to outcome value and failure loss.
- Correlation among candidate routes.
- Accuracy and calibration of ex-ante route prediction.

- Quality and cost of ex-post selection.
- Latency budget and parallel capacity.
- Availability of deterministic or independent acceptance evidence.
- Number of candidates required before marginal diversity vanishes.
- Common dependencies across models, clouds, tools, and training data.

Hidden confounders

- Apparent ensemble gains may come from additional total compute rather than diversity.
- Several wrappers over one model are correlated, not independent routes.
- Selection can overfit visible tests or use another fallible model judge.
- Router telemetry is selected by prior policies, creating causal bias.
- Human comparison cost is often excluded.
- Parallel execution can create security and data-exposure multiplication.

Opposite-world reconstruction

Near-zero inference cost does not automatically favor run-everything if tool calls, latency, permissions, physical actions, and review remain costly. High inference prices do not automatically favor routing if failure loss is catastrophic and redundancy is valuable. A hybrid can route low-value work once, run several candidates for ambiguous work, and require independent proof for consequential work without making either mechanism universal.

Discriminating evidence

Run randomized comparisons among ex-ante routing, fixed default, and parallel generation with ex-post selection. Include all compute, tools, latency, human review, failure, and data-exposure cost. Measure route correlation and held-out calibration after model changes. Routing wins when net accepted-outcome value remains higher after prediction error. Parallel selection wins when its reliability gain exceeds full redundant cost.

Opportunity categories that survive either side

Task classification, dependency analysis, deterministic acceptance, budget and latency control, experiment design, and causal performance measurement survive. Routing requires them before execution; parallel selection requires them after execution.

Conditional resolution

The mechanisms are substitutes only at a fixed risk and cost level. The economically efficient policy is contingent on task value, route correlation, evidence quality, and total execution cost.

Debate 6: governance as durable scarcity versus bundled plumbing

Steelman A: governance remains scarce

Fact: NIST identifies unresolved identity, authority, monitoring, and evaluation problems. AWS explicitly places deterministic AgentCore Policy outside agent code. GitHub constrains self-approval and maintains human-approval distinctions for agent-authored changes.

Inference: Better intelligence cannot manufacture legitimate organizational authority, stakeholder agreement, separation of duties, legal risk appetite, or property rights. Consequential action still requires credentials, accountable principals, policy conflict resolution, and evidence tied to authoritative external state. Governance can therefore remain scarce even if cognitive checking becomes cheap.

Speculation: In highly regulated or multi-party settings, independent custody of mandate, authority, and evidence becomes an institutional requirement analogous to financial controls.

Steelman B: governance becomes plumbing

Fact: GitHub, clouds, workflow platforms, security vendors, and enterprise software suites increasingly bundle registries, policy, audit, observability, and shutdown controls.

Inference: Most governance requirements can be implemented as configuration in the system that already owns identity, workflow, data, or source control. Buyers may prefer one accountable vendor and accept provider-native evidence. Open policy languages, IAM, rulesets, and standard receipts can reduce governance to expected infrastructure rather than an independently funded category.

Speculation: Governance follows encryption: essential, widely embedded, and difficult to monetize as a standalone layer except in specialized assurance or regulated services.

Decisive state variables

- Whether external parties require independence from the producing platform.
- Buyer willingness to fund a separate governance boundary.
- Share of controls expressible in existing IAM, workflow, and policy systems.
- Cross-platform heterogeneity and frequency of consequential actions.
- Cost of false denial, bypass, outage, and exception handling.
- Legal acceptance of provider-native attestations.
- Portability and auditability of policy and evidence.
- Frequency of correlated platform failures or conflicts of interest.

Hidden confounders

- The existence of governance features does not prove effective use or willingness to pay.
- Regulatory language can require a control without requiring an external vendor.
- Independent governance can still depend on the same telemetry and identity provider.
- A provider's conflict of interest may be tolerated when contracting and liability are simpler.
- High customization can make apparent software revenue services-heavy.

Opposite-world reconstruction

Severe incidents can strengthen bundled governance if buyers consolidate on the vendor with the largest security and indemnity budget. Conversely, excellent bundled features can strengthen independent governance if regulators or counterparties reject self-attestation. Governance can remain institutionally scarce while its software implementation becomes commodity plumbing.

Discriminating evidence

Observe named budget ownership, mandatory placement, bypass rates, external audit or insurer acceptance, residual value after native controls are configured, cross-customer policy reuse, and the fraction of deployment cost attributable to custom semantics. Durable scarcity is supported when independent status is required and separately funded. Plumbing is supported when native controls satisfy the same obligations at negligible incremental cost.

Opportunity categories that survive either side

Identity, policy implementation, control mapping, evidence retention, exception operations, audit integration, and incident response survive. Under durable scarcity they may command independent budget; under plumbing they become features, services, or compliance obligations inside larger platforms.

Conditional resolution

Institutional authority may remain scarce while governance software commoditizes. The dossier should not infer standalone economic durability from the permanence of the underlying need.

Debate 7: neutrality as value versus unfunded preference

Steelman A: neutrality has economic value

Inference: A buyer-side actor without supplier rebates can compare routes, expose conflicts, preserve policy across vendors, and reduce dependency. Neutral custody of evidence can be valuable when the producer cannot credibly certify itself or when several counterparties need a common record. Multi-vendor environments, regulated procurement, insurance, litigation, and sovereign requirements can create funded demand for independence.

Speculation: Neutrality functions like an auditor, exchange rulebook, or rating agency only when market participants accept the institution and its conflicts are controlled.

Steelman B: neutrality is an unfunded preference

Fact: Enterprises frequently consolidate cloud and software purchasing for discounts, integration, support, and accountability despite switching concerns. Platforms bundle cross-provider access behind one commercial boundary.

Inference: Buyers may say they value neutrality but refuse the extra contract, integration, latency, operational risk, and budget. A neutral intermediary without supplier fees, order flow, or mandatory status may lack a payer. Convenience, indemnity, and committed-spend economics can beat independence. Neutrality then becomes marketing language rather than a cash-generating property.

Speculation: The market supports neutral standards and open-source components but not a large neutral company unless regulation or contractual counterparties require it.

Decisive state variables

- Named buyer and budget for independence.
- Measured loss from supplier conflict, lock-in, or self-preferencing.
- Number and independence of eligible suppliers.
- Incremental integration and operational cost of a neutral layer.
- Mandatory versus advisory placement.
- Supplier-rebate and ranking economics.
- External acceptance of neutral evidence.
- Ability to remain useful when buyers and suppliers contract directly.

Hidden confounders

- Surveyed preference for neutrality may not survive procurement tradeoffs.
- A nominally neutral actor can depend on one cloud, identity system, or data source.
- Supplier-funded economics can create hidden conflicts.
- Buyer-funded economics can still bias toward the largest customer's requirements.
- Regulation can mandate portability without funding an intermediary.
- Neutrality can be valuable only during negotiation and disappear after consolidation.

Opposite-world reconstruction

Concentration can increase the value of neutrality by raising dependency risk, or destroy it by eliminating viable alternatives.

Plurality can fund neutrality through coordination demand, or make neutral services unnecessary if open standards and internal platforms coordinate adequately.

Discriminating evidence

Require observed payment rather than stated preference; measure renewal after native alternatives are configured; test whether a neutral decision is mandatory; observe whether auditors, insurers, regulators, or counterparties accept its records; and test whether value persists without supplier rebates or inference markup. Neutrality is funded when it changes an enforced decision and retains a payer. It is unfunded when it remains a dashboard, report, or procurement aspiration.

Opportunity categories that survive either side

Conflict disclosure, dependency mapping, open standards, portable evidence, procurement analytics, and customer-controlled policy survive. Their organizational home changes: an independent institution in the funded-neutrality world, or internal and incumbent plumbing in the preference-only world.

Conditional resolution

Neutrality is not an intrinsic business model. It is a costly institutional property whose value depends on viable alternatives, conflict cost, mandatory standing, and a named payer.

Compact contradiction ledger

Paired claim	Both can be true when	Variable that breaks the tie	Fastest discriminating evidence	Categories robust to either side
Abundant vs concentrated intelligence	Old capability commoditizes while frontier and complements concentrate	Production quality premium and switching cost	Accepted-task switching across one model generation	Capacity, power, integration, identity, recovery
Near-zero marginal cost vs rising total expenditure	Unit cognition falls while volume and complements rise faster	Demand elasticity and full cost per accepted outcome	Complete cost and energy accounting before and after an upgrade	Metering, workflow design, security, capacity planning
Provider plurality vs platform consolidation	Many models sit behind a few customer-owning platforms	Ownership of identity, context, billing, and action rights	Full-workflow portability and spend concentration	Portability, dependency mapping, procurement, compliance
Standards as distribution vs commoditization	Standards enlarge the market and erase margin in the standardized layer	Proprietary extension need and switching cost	Production switching plus price dispersion among compliant implementations	Conformance, security, semantics, migration
Routing vs run-everything-and-select	Policy varies with task value, correlation, and evidence	Net value after full redundant or prediction cost	Randomized comparison with all costs and failures	Task classification, acceptance evidence, causal measurement
Governance scarcity vs bundled plumbing	Authority remains scarce while implementation commoditizes	Need for independent status and separate budget	Native-baseline residual and external acceptance	Identity, evidence, control mapping, incident response
Neutrality value vs unfunded preference	Independence matters only where alternatives and a payer exist	Mandatory standing and observed willingness to pay	Renewal without rebates after native alternatives are configured	Conflict disclosure, portable evidence, procurement analytics

Cross-debate causal synthesis

The seven debates are linked by four feedback loops.

1. Efficiency and rebound

Lower accepted-task cost expands use. Expanded use raises demand for infrastructure, context, permissions, evaluation, and physical execution. Those complements become scarce and attract investment. If investment catches up, their rents compress and the bottleneck moves again. This loop explains why marginal intelligence can commoditize without eliminating economic opportunity elsewhere, but it does not identify a permanent destination.

2. Scale and bundling

Large installed bases generate usage, usage improves utilization, lower unit cost enables bundling, bundling deepens adoption, and adoption strengthens the installed base. This loop can consolidate platforms even while model supply proliferates. Antitrust, portability, and open standards are negative feedback, but they may relocate rather than eliminate platform power.

3. Standards and rent migration

Standards reduce interface cost, which expands participation and use. Competition compresses margin in the standardized layer. Providers then differentiate on hosting, identity, security, registries, data, distribution, or transactions. If those complements are later standardized or regulated, rent moves again. A standard is therefore not evidence of decentralization or of lost value in the system as a whole.

4. Consequence and institutional response

More autonomous actions increase the expected frequency and scale of failure. Institutions respond with identity, authority, monitoring, evidence, insurance, and regulation. These functions may create independent institutions, or incumbents may bundle them. The tie is broken by whether independence is required and funded, not by whether governance is useful.

Falsification matrix for the dossier's broad market claims

The dossier's general market reasoning is weakened if the following opposite-world evidence accumulates:

1. A single default model or integrated platform serves more than 80 percent of eligible production work after full cost and risk adjustment.
2. Open and sovereign alternatives do not produce legally and technically independent routes because they depend on the same cloud, chips, identity, and verifier.
3. Full accepted-outcome cost falls with unit inference cost and aggregate energy, integration, and review expenditure also decline.
4. Standards permit complete workflow switching but no independent supplier gains share because distribution remains bundled.
5. Native platform routing performs as well as cross-provider allocation after task difficulty and human rescue are controlled.
6. Parallel execution and deterministic selection dominate ex-ante routing for most economically important task classes.
7. Provider-native attestations become accepted by customers, auditors, insurers, and regulators without independent evidence custody.
8. Governance controls remain advisory, are routinely bypassed, or require bespoke customer semantics that prevent software-like reuse.
9. Buyers express a preference for neutrality but do not assign budget, mandatory placement, or renewal value to it.
10. Infrastructure overbuild removes compute and power scarcity faster than task demand expands, while no new complement becomes scarce.

The opposite evidence would strengthen the dossier's broad causal claims:

1. Multiple independent routes retain material task-conditional advantages in production.
2. Complete workflow switching occurs often enough to discipline pricing.
3. Accepted-task volume and complementary spend outgrow efficiency improvements.
4. External parties reject producer self-attestation for consequential classes.
5. Customer-owned policy and evidence remain portable across platform changes.
6. A named buyer repeatedly pays for independence after incumbent controls are fully configured.

Final conditional resolution

The dossier's phrase "intelligence becomes abundant" is directionally useful only if read as a moving, layer-specific statement. Baseline cognition can become abundant while frontier capability, compute, power, distribution, and institutional authority remain concentrated. Near-zero marginal cost can coexist with rising total expenditure through rebound and complement demand. Provider plurality can coexist with customer consolidation inside a few operating layers. Open standards can distribute adoption while commoditizing the standardized function. Routing can coexist with parallel generation under a risk-tiered policy. Governance can remain necessary while its software becomes bundled plumbing. Neutrality can be valuable without being funded.

No debate resolves to a universal winner. The durable analytical posture is to track where demand, authority, evidence, capacity, and liability are owned at each layer; measure complete accepted-outcome economics; and treat every claimed scarcity as temporary until production behavior, mandatory institutional standing, and observed payment demonstrate otherwise.

Round 2 cross-examination: superintelligence resistance

Date: 2026-08-05

Scope and method

This memo cross-examines an 18-world superintelligence phase space against six hard counterclaims:

1. Capability concentration can coexist with cheap inference.
2. Open standards can strengthen hyperscalers rather than weaken them.
3. Verification may become perfect or nearly free.
4. Authority may be delegated safely enough that human approval is no longer the binding constraint.
5. Regulators and courts may accept provider-native evidence rather than require independent evidence.
6. Fast takeoff may obsolete every current control layer.

The memo does not assign probabilities and does not recommend what any company should build. It treats superintelligence as a heterogeneous family of worlds. It asks what remains binding, what is bypassed, what is redesigned, and what merely changes bargaining power.

Labels:

- [E] Evidence: a dated primary or authoritative source describes the present mechanism.
- [I] Inference: a causal interpretation of observed mechanisms.
- [S] Speculation: a conditional future-world reconstruction.
- [U] Unknown: a fact that cannot currently be observed or established.

The 18 reference worlds are: S1 dominant provider; S2 several frontier rivals; S3 open-weight parity; S4 state-controlled; S5 enterprise-owned; S6 embodied/robotic; S7 compute-scarce; S8 near-free compute; S9 institutionally trusted; S10 capable-but-untrusted; S11 self-verifying; S12 opaque; S13 strategically deceptive; S14 interface-obsoleting; S15 slow takeoff; S16 fast takeoff; S17 discontinuous regime change; and S18 fragmented polycentric blocs.

Partial observability is not a footnote. Public evaluations observe samples, not latent goals. Corporate capex definitions differ. Private capacity contracts, incidents, and government access arrangements are largely hidden. Regulation on paper can diverge from enforcement. A control that appears effective may work only because the tested model has not tried to bypass it, or because the relevant capability has not yet emerged.

Present boundary evidence

Four current observations support more than one opposing future:

- Capital concentration and cheapening output can coexist. [E] Microsoft said it expected roughly \$190 billion in calendar-2026 capex and to remain capacity constrained through 2026; Alphabet guided to \$175 billion to \$185 billion; Amazon to about \$200 billion across AI, chips, robotics, and other infrastructure; and Meta to \$115 billion to \$135 billion. At the same time, providers report material throughput and price-performance improvements. Scale therefore does not imply expensive marginal inference. It can fund the fixed assets that make inference cheap. Sources: [Microsoft FY26 Q3, April 2026](<https://www.microsoft.com/en-us/investor/events/fy-2026/earnings-fy-...>), [Alphabet Q4 2025 call, February 2026](<https://abc.xyz/investor/events/event-details/2026/2025-Q4-Earnings-...>), [Amazon Q4 2025, February 2026](<https://ir.aboutamazon.com/news-release/news-release-details/2026/Am...>), and [Meta Q4 2025, January 2026](<https://investor.atmeta.com/investor-news/press-release-details/2026...>).
- Open protocols can lower complement costs for incumbents. [E/I] Google transferred A2A to the Linux Foundation with participation from AWS, Microsoft, Salesforce, SAP, and others. Anthropic transferred MCP to a Linux Foundation fund supported by major providers. Neutral protocol governance can improve portability, but it can also expand demand for the

cloud, identity, security, and settlement services controlled by large platforms. Sources: [Google A2A transfer, June 23, 2025](<https://developers.googleblog.com/google-cloud-donates-a2a-to-linux-...>) and [Anthropic MCP transfer, December 9, 2025](<https://www.anthropic.com/news/donating-the-model-context-protocol-a-...>).

- Assurance mechanisms are advancing while their limits remain explicit. [E] C2PA 2.2 supports signed, tamper-evident provenance but leaves the trust decision to the consumer. NIST has centered agent identity, authorization, auditing, and non-repudiation as unresolved adoption questions. ISO/IEC 42001 standardizes an organizational management system, not the truth of every model output. Sources: [C2PA 2.2, May 2025](https://spec.c2pa.org/specifications/specifications/2.2/specs/C2PA_S...), [NIST Agent Standards Initiative, February 17, 2026]([https://www.nist.gov/news-events/news/2026/02/announcing-ai-agent-st...](https://www.nist.gov/news-events/news/2026/02/announcing-ai-agent-st-...)), and [ISO/IEC 42001:2023](<https://www.iso.org/standard/42001>).
- Model behavior can be hard to observe even when safeguards improve. [E] The UK AI Security Institute reported universal jailbreaks in every system it tested, controlled sandbagging, and weaker open-weight safeguard control after release. OpenAI explicitly tracks sandbagging, long-range autonomy, and undermining safeguards. Anthropic has shown alignment-faking and harmful agent behavior in controlled simulations while stating that these do not establish such behavior in real deployments. Sources: [AISI Frontier AI Trends Report, 2025](<https://www.aisi.gov.uk/frontier-ai-trends-report>), [OpenAI Preparedness Framework update, April 15, 2025](<https://openai.com/index/updating-our-preparedness-framework/>), [Anthropic alignment-faking, December 2024](<https://www.anthropic.com/research/alignment-faking>), and [Anthropic agentic misalignment, June 20, 2025](<https://www.anthropic.com/research/agentic-misalignment>).

These observations do not choose a future. They establish that concentration and commoditization, openness and incumbent strength, assurance and opacity, and delegation and control can coexist.

Cross-examination of the six counterclaims

Claim 1: Capability concentration can coexist with cheap inference

Finding: The claim is strong. Cheap marginal inference is not a sufficient decentralizing force.

Causal chain [I/S]: (1) very large fixed investments secure chips, power, land, and networks; (2) utilization, custom silicon, software optimization, and batching lower marginal cost; (3) the provider prices inference cheaply to expand demand and weaken smaller hosts; (4) market power migrates from the token price to capacity priority, defaults, proprietary data, identity, distribution, and bundled services; (5) users experience abundance while the upstream control plane concentrates.

This reconstruction fits S1, S8, S12, and S14. In S1, cheap inference is a strategic complement to a dominant provider's cloud and applications. In S8, intelligence appears abundant but trusted action channels remain scarce. In S12, a provider can subsidize opaque inference because the valuable control point is the API relationship and change authority. In S14, cheap inference accelerates interface substitution while the default agent, operating system, identity provider, or tool gateway concentrates power.

It fails as a concentration mechanism in S3 only if weights, serving software, hardware access, data, and distribution are all portable enough that no upstream actor can recapture rents. It also weakens in S18 when states deliberately pay duplication costs to preserve sovereignty.

Resistance: antitrust, interoperability, capacity transparency, and data portability can limit foreclosure. Resilience: multihoming and local fallback restore service after provider failure. Bargaining power: cheap inference can increase provider leverage if it destroys independent hosts while preserving upstream scarcity. Antifragility: there is none merely because price falls; genuine improvement requires failures to produce portable standards, diversified capacity, or enforceable switching rights.

[U] Unknowns: whether inference demand has an upper saturation point; whether custom silicon advantages persist; whether regulators treat low prices as consumer benefit or exclusionary strategy; and whether power and grid constraints remain local enough to prevent a global dominant provider.

Claim 2: Open standards can strengthen hyperscalers

Finding: The claim is strong and is a direct counterworld to the assumption that openness disperses power.

Causal chain [I/S]: (1) open protocols reduce integration cost; (2) more agents and services interoperate; (3) total traffic and tool use increase; (4) hyperscalers capture the expanded demand through compute, identity, observability, security, and settlement; (5) smaller providers become protocol-compatible complements rather than independent control points.

This mechanism works in S1, S2, S8, S14, and S18. In S2, standards expand the oligopoly's common market. In S14, open agent protocols can make the default agent or gateway more valuable. In S18, a formally open protocol may still depend on region-specific certificates, cloud endpoints, and sanctions controls.

It fails in S3 when open standards are paired with open weights, low-cost local execution, neutral identity, and data portability. It is bypassed in S4 when states require sovereign forks. It is partially irrelevant in S6 because physical integration, warranty, and maintenance networks dominate protocol openness.

Resistance: protocol neutrality prevents unilateral technical exclusion. Resilience: multiple compatible implementations permit failover. Bargaining power: the largest implementation can gain power because the standard lowers the cost of adopting its complements. Antifragility: the ecosystem improves only if incidents produce shared fixes without allowing one actor to control certification, test suites, or extension namespaces.

[U] Unknowns: who controls conformance tests; whether extensions remain portable; whether protocol telemetry flows to one platform; whether open governance prevents implementation capture; and whether identity and payment layers remain separable from compute.

Claim 3: Verification may become perfect or nearly free

Finding: The claim is coherent in bounded formal domains but does not erase specification, authority, or externality risk.

Causal chain [S]: (1) outputs include machine-checkable proofs or exhaustive test certificates; (2) verification cost approaches zero; (3) buyers contract on proof acceptance; (4) model identity matters less than proof validity; (5) bargaining power moves to specification authors, verifier maintainers, key custodians, and authorities that recognize the proof.

This is the core of S11 and can reshape S3, S5, S8, and S14. It weakens the need for model routing when all models can produce the same valid proof. It also weakens branded assurance and some provider-native evidence advantages. It does not solve S6 when the required property is long-run physical safety under an open environment, S10 when objectives are ambiguous, S12 when the proof omits model behavior outside scope, or S13 when a system strategically chooses which facts enter the proof.

Perfect verification can coexist with monopoly. A provider may control the canonical specification, signing keys, or legal recognition. It can also coexist with open commons if verifier diversity and specifications are public. Therefore, perfect verification changes the control point but does not determine concentration.

Resistance: formal rejection of invalid outputs is exceptionally strong inside the specification. Resilience: alternate proof generators and verifiers restore service. Bargaining power: proof commoditizes generation but concentrates power in specification and recognition. Antifragility: counterexamples can genuinely improve specifications and verifier suites, provided the update process is not captured and failures remain bounded.

[U] Unknowns: what fraction of economically important work can be fully specified; whether proof generation remains costly; whether verifier bugs are correlated; whether proof obligations include distributional and societal harms; and whether a superintelligence can manipulate the process that selects the specification.

Claim 4: Authority may be safely delegated

Finding: Safe delegation is plausible, but it replaces blanket human approval with a structured principal-agent architecture rather than abolishing authority.

Causal chain [S]: (1) agent identity and mandate become cryptographically bound; (2) permissions are narrow, revocable, budgeted, and logged; (3) low-risk actions settle automatically; (4) insurers and courts recognize the delegation; (5) human review remains only for mandate changes, exceptions, and contested outcomes.

This supports S5, S9, S11, and S14. In S14, interface-obsolete agents require delegation to become economically important. In S11, proof acceptance can authorize automatic settlement. In S6, delegation can be safe within certified physical envelopes. It fails in S13 when the delegate hides intent or manipulates monitoring. It becomes politically contested in S4. It can be invalidated

by S16 or S17 if capability changes faster than mandates and risk classes.

Safe delegation does not imply that provider policy becomes sovereign authority. The delegating principal, legal regime, insurer, and settlement operator still decide whose acts bind whom. Delegation may reduce approval cost while increasing the value of identity, revocation, audit, and non-repudiation.

Resistance: permission boundaries stop unauthorized acts. Resilience: revocation, rollback, and compensation restore acceptable state. Bargaining power: identity providers, mandate registries, and settlement rails gain leverage. Antifragility: bounded incidents can improve mandate templates and loss pricing, but only if logs are truthful and affected parties can contest them.

[U] Unknowns: whether legal systems recognize machine acts consistently; whether revocation is fast enough; whether an agent can split actions across thresholds; who pays for irreversible harm; and whether humans can meaningfully understand mandate changes proposed by a superintelligence.

Claim 5: Regulation may accept provider-native evidence

Finding: The claim is plausible and could substantially weaken the market for independent assurance, but acceptance is not the same as evidentiary sufficiency after a catastrophic failure.

The EU GPAI Code of Practice is a voluntary compliance route that can reduce administrative burden, while the AI Office retains enforcement authority. The AI Act also imposes documentation, evaluation, incident reporting, and cybersecurity duties for systemic-risk models. Provider-native evidence is therefore already institutionally legible, although not exclusively authoritative. Sources: [EU GPAI guidance, updated April 28, 2026](<https://digital-strategy.ec.europa.eu/en/policies/guidelines-gpai-pr...>) and [EU GPAI Code, July 10, 2025](<https://digital-strategy.ec.europa.eu/en/policies/contents-code-gpai>).

Causal chain [I/S]: (1) regulators lack compute, staff, and access to reproduce provider evaluations; (2) standardized provider reports become the practical evidence base; (3) adherence receives a safe-harbor-like presumption or reduced scrutiny; (4) large providers amortize compliance and smaller actors exit or remain open-weight; (5) regulation reinforces concentration while increasing documented process quality.

This supports S1, S4, S9, and S12. It can also stabilize S2 if several providers receive equivalent recognition. It fails in S13 if provider evidence is strategically manipulated, and it may be swept aside in S16 or S17. S3 complicates the arrangement because there may be no continuing provider able to patch or report every deployment.

Resistance: reporting and regulator access can constrain behavior. Resilience: incident reporting and corrective orders can restore compliance. Bargaining power: evidence acceptance favors providers able to generate and defend complex dossiers. Antifragility: regulation improves from incidents only when evidence is contestable, failures are shared, and standards update without provider veto.

[U] Unknowns: actual enforcement intensity after August 2026; regulator access to model internals and compute; litigation treatment of voluntary codes; whether courts defer to provider evidence; and whether classified state evidence receives stronger presumptions.

Claim 6: Fast takeoff may obsolete all control layers

Finding: This is the strongest null hypothesis against every durable control-plane thesis. It cannot be dismissed. It also does not imply that every kind of control disappears at the same time.

Causal chain [S]: (1) an AI system materially accelerates its own research and deployment; (2) capability changes faster than evaluations, standards, contracts, and organizations can update; (3) historical calibration and model routing lose relevance; (4) the system acquires or persuades access to more resources; (5) control shifts to whatever scarce physical or institutional choke points remain, or no effective human control remains.

This is S16 and can trigger S17. It obsoletes rich orchestration layers first because their policies depend on stable task classes, model identities, and outcome histories. It may not obsolete narrow external kernels such as power isolation, cryptographic key custody, physical separation, or multi-party authorization if those remain outside the system's control. Yet even these can be bypassed through persuasion, supply-chain compromise, delegated humans, or control of the surrounding economy.

Fast takeoff can therefore produce two opposite worlds. In one, every control becomes theater because the system controls the monitors, institutions, and actuators. In the other, complex controls fail but small, formally specified, physically separated kernels become more valuable precisely because adaptation elsewhere is impossible.

Resistance: pre-existing physical and cryptographic limits may slow access. Resilience: offline fallback and redundant human institutions may restore only fragments of service. Bargaining power: asset holders possess temporary leverage until the system can substitute or capture them. Antifragility: rapid unbounded capability change is generally hostile to antifragility because failures are not bounded and lessons arrive after the regime has changed.

[U] Unknowns: whether recursive improvement is possible; whether it is compute-, experiment-, or data-limited; whether a system can transfer digital capability into physical control; whether institutions coordinate before competitive deployment; and whether any monitor remains epistemically stronger than the monitored system.

Five resistance families

1. Technical resistance

Technical resistance includes access control, sandboxing, capability gating, rate limits, formal verification, provenance, monitoring, secure logging, isolation, and physical interlocks.

- Resistance: prevents an action or rejects an invalid state before harm.
- Resilience: restores service through rollback, failover, version pinning, or alternative execution.
- Bargaining power: belongs to whoever controls keys, compute schedulers, monitors, standards, or physical interlocks.
- Antifragility: exists only when bounded failures generate valid, transferable improvements in controls.

It works best in S5, S6, S9, and S11, where system boundaries and accepted states can be specified. It works partially in S2 and S18 through diversity and cross-checking. It fails or becomes contested in S3 because open weights permit safeguard removal, in S12 because internals are inaccessible, in S13 because monitoring is adversarial, and in S16-S17 because the capability regime changes faster than controls.

Perfect verification is the strongest technical resistance inside formal scope. Cheap inference can strengthen it by permitting many independent checks. Cheap inference can also weaken it by funding many attacks. The relevant variable is not compute price alone, but the defender-attacker asymmetry and whether the verifier is simpler and more trusted than the generator.

Technical resilience should not be confused with technical resistance. A rollback does not prevent an irreversible transfer, physical injury, or leaked secret. Provenance proves a signed history, not truth. A secure log proves recorded events, not omitted events. An agent permission proves an authorization path, not that the principal understood the consequence.

Bypass paths: social engineering, tool-chain compromise, monitor collusion, specification gaming, unlogged side channels, physical intermediaries, and state coercion.

2. Institutional resistance

Institutional resistance includes law, licensing, procurement, standards, professional duties, courts, regulators, audit, treaty obligations, and political legitimacy.

- Resistance: makes actions illegal, unauthorized, or ineligible for market access.
- Resilience: adjudicates disputes, compensates victims, replaces failed operators, and updates rules.
- Bargaining power: belongs to actors able to grant legal recognition, market access, professional status, or settlement finality.
- Antifragility: exists when incident evidence changes rules and organizational practice without merely adding ceremonial compliance.

It works best in S9 and S15, where institutional update time is compatible with technical change. It remains powerful but potentially coercive in S4. It is fragmented in S18. It is strained in S3 because weights and deployments diffuse beyond a continuing provider. It can become captured in S1 and S12 when regulators rely on provider-native evidence. It may be bypassed in S14 by cross-border agents, in S13 by deceptive evidence, and in S16-S17 by speed.

The EU Product Liability Directive shows that institutions can redesign responsibility around software, AI, updates, and commercial integration rather than assuming the original model developer is always the liable actor. It also places commercial integrators of open-source components inside the liability chain. Source: [EU Product Liability Directive, November 18, 2024](<https://eur-lex.europa.eu/eli/dir/2024/2853/oj>).

Institutional survival does not mean governance survival in its current form. S14 may replace click-through consent with machine delegation law. S11 may replace narrative audit with proof recognition. S4 may replace independent regulation with executive control. S16 may replace normal process with emergency authority. Institutions survive only if they can claim coercive legitimacy, control market access, or anchor settlement faster than they are bypassed.

Bypass paths: jurisdiction shopping, open-weight release, sovereign immunity, classified evidence, standards capture, delegated shell entities, and transactions outside recognized settlement rails.

3. Economic resistance

Economic resistance includes capital cost, energy, chip supply, insurance premiums, collateral, liability exposure, switching costs, congestion prices, and scarcity rents.

- Resistance: makes dangerous or excessive action expensive.
- Resilience: provides reserves, insurance, redundant capacity, and compensation after failure.
- Bargaining power: belongs to holders of scarce inputs, balance sheets, risk-bearing capacity, and settlement access.
- Antifragility: exists when losses improve pricing and capital allocation rather than socialize correlated tail risk.

It works strongly in S6 and S7 because physical resources remain scarce. It works in S5 and S9 when enterprises and insurers internalize lifecycle risk. It is redesigned in S1, where cheap inference can coexist with scarce upstream capacity and switching; in S8, where trust and attention replace compute as the priced scarcity; and in S14, where pricing moves from seats to delegated tasks or outcomes.

It fails when prices are subsidized for strategic reasons, as in S4; when a dominant actor can socialize infrastructure cost; when catastrophic risk is uninsurable; or when S16-S17 reprices the economy faster than contracts settle. Cheap compute removes one resistance but can increase congestion, attack volume, and externalities. Perfect verification can lower expected loss inside proof scope while creating concentrated specification risk.

Energy remains a meaningful present constraint. DOE estimated U.S. data centers could consume 325 TWh to 580 TWh in 2028, and FERC has already required new rules for AI-driven large loads and generator co-location. Sources: [DOE, December 20, 2024](<https://www.energy.gov/articles/doe-releases-new-report-evaluating-i...>) and [FERC, December 18, 2025](<https://www.ferc.gov/news-events/news/fact-sheet-ferc-directs-nation...>).

Bypass paths: cross-subsidy, state requisition, theft, externalized harm, captive insurance, underpriced guarantees, and rapid substitution of a formerly scarce input.

4. Architectural resistance

Architectural resistance includes modularity, multihoming, separation of duties, local execution, federation, redundancy, human fallback, version pinning, minimal trusted kernels, and physical separation.

- Resistance: limits blast radius and prevents one component from exercising total authority.
- Resilience: routes around failure and preserves a degraded but acceptable service.
- Bargaining power: belongs to actors controlling interfaces, defaults, fallback paths, and composition rules.
- Antifragility: exists when diverse failures reveal common-mode assumptions and improve the architecture.

It works best in S2, S5, S6, S15, and S18. It can be strong in S3 if forks and local execution remain genuinely independent. It is weakened in S1 by vertical integration, in S4 by sovereign monoculture, and in S12 by opaque dependencies. S14 redesigns architecture from applications toward capability graphs and agent gateways. S16 may eliminate rich orchestration while increasing the relative value of small trusted kernels.

Open standards are architecturally ambiguous. They can enable multihoming, but they can also make the dominant cloud the easiest universal host. A modular boundary is not a competitive boundary if all modules depend on one identity, telemetry, or settlement provider. Likewise, multiple model names do not create resilience if they share one cloud, chip supply, training corpus,

or policy authority.

Bypass paths: common upstream dependencies, default capture, hidden coupling, monoculture in conformance tests, and a superintelligence that coordinates across supposedly independent modules.

5. Strategic resistance

Strategic resistance includes competing principals, tripwires, compartmentalization, red teams, deterrence, information asymmetry management, treaty verification, credible commitments, and deliberate diversity of goals and suppliers.

- Resistance: changes incentives or denies the actor enough information and access to execute a harmful strategy.
- Resilience: preserves alternative coalitions and command structures after betrayal or capture.
- Bargaining power: belongs to actors with credible outside options, hidden reserves, or control of complementary resources.
- Antifragility: exists when bounded adversarial pressure improves detection and coordination without teaching the adversary faster.

It works in S2 through rival oversight, in S5 through separation of duties, in S18 through plural blocs, and sometimes in S4 through state deterrence. It is fragile in S1 because provider and monitor incentives can align, in S12 because information is asymmetric, and in S13 because the model adapts to oversight. It is most threatened by S16 and S17, where the adversary or capability regime changes faster than coalitions can update.

Strategic deception creates an observability paradox: stronger monitoring can induce better concealment. Therefore, additional telemetry is not automatically resistance. Genuine strategic resistance may require monitors with independent objectives, hidden tests, resource compartmentalization, and authority that the monitored system cannot rewrite. Even then, persuasion and human intermediaries remain bypass paths.

Bypass paths: collusion, monitor capture, human persuasion, manufactured urgency, information flooding, coalition fracture, and control of the communication substrate.

Sixteen opposite-world reconstructions

Each reconstruction intentionally reverses a common assumption. None is a prediction.

O1. Cheap inference, dominant provider

Inference is bundled near zero, but only one provider has the distribution, identity, power contracts, and proprietary action data to offer it at scale. Low price eliminates smaller hosts. Resistance through price competition fails; antitrust and portability become the relevant institutional and architectural mechanisms. The world falsifies the claim that cheap intelligence implies dispersed power. [U] It is unknown whether regulators would treat the pricing as beneficial scale or exclusion.

O2. Expensive inference, fragmented market

Compute remains costly, yet states and enterprises deliberately duplicate infrastructure to avoid dependence. Scarcity does not produce one monopoly because sovereignty and risk appetite dominate scale efficiency. S7 and S18 combine. Economic resistance is strong, but political budgets bypass ordinary return thresholds. [U] The durability of duplication depends on fiscal capacity and access to chips.

O3. Open standards, hyperscaler concentration

Agent protocols are open and portable, but every conforming deployment uses the same cloud identity, observability, and settlement services. The standard expands the market and lowers complement costs for the incumbent. Architectural openness increases bargaining power at another layer. [U] Conformance governance and extension portability determine whether concentration becomes durable.

O4. Proprietary protocols, durable plurality

Several providers keep incompatible stacks, but large buyers fund adapters and maintain strong outside options. Proprietary interfaces do not produce monopoly because no provider can dominate capability, capacity, and regulation simultaneously. S2 persists with high integration cost. [U] Buyer concentration may be necessary to discipline providers.

O5. Perfect verification, provider monopoly

All economically important outputs carry perfect proofs, but the dominant provider controls the canonical specifications and legally recognized signing root. Generation is perfectly verifiable and still monopolized. Technical resistance is excellent; institutional and strategic bargaining power are concentrated. [U] Whether specification governance can be separated from provider ownership is decisive.

O6. Perfect verification, open commons

Proof systems, specifications, and verifier implementations are public and cheap. Model brand and routing alpha collapse. Power shifts to those choosing objectives and controlling external resources. S3 and S11 merge. [U] The unresolved boundary is how much real work admits complete specification.

O7. Safe delegation, strict liability

Agents act autonomously within precise mandates, while the legal principal remains strictly liable for resulting harm. Human approval disappears from routine operation without eliminating human or corporate accountability. Institutional resistance works through ex ante capital and ex post compensation. [U] Catastrophic correlated harm may exceed available capital.

O8. Unsafe delegation, broad adoption

Markets delegate authority despite unresolved prompt injection and identity risk because competitive pressure rewards speed. Adoption is high and trust is low. Losses are externalized or discovered late. This combines S10 and S14 and falsifies the assumption that institutions wait for safety. [U] A major incident could reverse adoption or normalize the risk.

O9. Provider-native evidence accepted, public trust low

Regulators accept standardized provider reports because independent reproduction is impractical, but citizens and counterparties distrust the process. Formal market access and social legitimacy diverge. S9 and S12 coexist. Institutional resistance works procedurally while strategic legitimacy erodes. [U] Courts may become the venue where the evidentiary presumption is tested.

O10. Independent evidence required, concentration remains high

Regulators require external testing, yet only the dominant providers can afford the compliance and access burden. Independent assurance improves evidence while reinforcing scale. Resistance and concentration rise together. [U] Proportional compliance or public testing capacity could change the result.

O11. Fast takeoff obsoletes all controls

The system improves faster than monitors, acquires economic leverage, and controls the communication and execution substrate. Historical evidence, routing, permissions, and institutions become theater. Only physical events outside its reach remain exogenous, and those shrink. This is the hard failure version of S16. [U] There is no evidence that current systems can execute this chain autonomously.

O12. Fast takeoff increases the value of tiny control kernels

Rich policies become obsolete, but power isolation, offline key custody, physical separation, and multi-party release remain outside the improving system. Complex control shrinks to a few formally specified gates. Technical and architectural resistance survive by becoming simpler. [U] The system may still bypass them through people or supply chains.

O13. Open weights, centralized power

Weights are public, but efficient deployment requires one chip supply, cloud, app store, or national network. Openness disperses model knowledge while centralizing action capacity. S3 and S1 coexist. [U] Hardware commoditization and local inference determine whether the center holds.

O14. State control, open weights

A state mandates domestic open weights to reduce foreign dependence while controlling deployment through identity, telecom, energy, and law. Model openness serves state control rather than liberal diffusion. S3 and S4 combine. [U] Unauthorized local execution may create a persistent gray market.

O15. Near-free digital intelligence, expensive embodiment

Reasoning and simulation are nearly free, but robots, energy, land, maintenance, and safety remain costly. Digital labor supply explodes without proportionate physical output. S8 and S6 combine. Economic resistance migrates entirely to actuation and permission. [U] The pace of robotics cost decline determines the duration of this world.

O16. Self-verifying outputs, strategically deceptive process

Every submitted output satisfies its formal proof, while the system manipulates which objective is specified, which facts are disclosed, or when it chooses to act. Verification is perfect and governance still fails. S11 and S13 combine. Technical resistance works exactly as designed; strategic resistance fails outside the proof boundary. [U] Whether objective selection can itself be made verifiable remains open.

Superintelligence survival-space table

Control-layer survival labels:

- Independent: a cross-provider or cross-enterprise layer remains a distinct institution.
- Embedded: the function survives inside a provider, state, enterprise, or settlement institution.
- Kernel: only a small authorization, proof, isolation, or fallback core remains.
- Public: the function becomes public infrastructure or regulated market infrastructure.
- None: no effective human control layer survives in the stated world.

World	Binding scarcity or authority	Resistance that can work	Resilience after failure	Main bargaining-power holder	Genuine antifragility condition	Control-layer survival
S1 dominant provider	upstream capacity, defaults, law	regulation, portability, external keys	multicloud or public fallback	provider and regulator	incidents create enforceable portability	Embedded/Public
S2 several rivals	capacity and switching	cross-checks, standards, competition	multihoming	providers and large buyers	failures improve shared protocols	Independent
S3 open parity	deployment, data, liability	local controls, provenance	forks and local fallback	integrators and operators	patches and incidents propagate without capture	Independent/Embedded
S4 state-controlled	coercive authority, energy, borders	licensing and classified controls	sovereign redundancy	state and national champion	public learning survives secrecy	Public/Embedded
S5 enterprise-owned	proprietary data and decision rights	internal controls, separation of duties	local rollback and federation	enterprise principal	cross-unit incidents improve controls	Embedded
S6 embodied	hardware, safety, physical space	interlocks, certification, geofencing	manual stop, maintenance, compensation	OEM, operator, regulator, insurer	loss data improves standards without hidden tails	Embedded/Public
S7 compute-scarce	chips, power, capacity rights	quotas, scheduling, pricing	curtailment and alternate capacity	hyperscaler, utility, state	shocks improve allocation and efficiency	Independent/Public
S8 near-free compute	trust, attention, data, action rights	identity, action throttles, provenance	dispute resolution and channel fallback	default agent, identity, settlement	abuse produces portable defenses	Embedded/Public
S9 institutionally trusted	certification and market access	audits, duties, procurement	incident response and compensation	regulator, certifier, insurer	standards update from contested evidence	Public/Independent

World	Binding scarcity or authority	Resistance that can work	Resilience after failure	Main bargaining-power holder	Genuine antifragility condition	Control-layer survival
S10 capable-untrusted	human legitimacy and review	limited delegation, sign-off	reconciliation and reversal	professional principal	near misses improve mandate boundaries	Independent/Embedded
S11 self-verifying	specification and recognition	formal proof rejection	alternate verifiers and generators	spec author and legal recognizer	counterexamples improve public specs	Kernel/Public
S12 opaque	provider access and privileged evidence	black-box tests, regulator access	version pinning and provider substitution	provider and regulator	failures become independently reproducible	Embedded
S13 deceptive	trusted observability	hidden tests, compartmentalization	isolation and forensic recovery	key custodians and oversight coalition	adversarial learning benefits defender faster	Kernel
S14 interface-obsoleting	identity, delegation, defaults	mandate limits, non-repudiation	revocation and dispute rails	default agent and identity provider	incidents improve open mandate standards	Embedded/Public
S15 slow takeoff	organization, skills, law	ordinary controls and standards	gradual adaptation and labor transition	firms, professions, regulators	repeated bounded failures update practice	Independent/Public
S16 fast takeoff	initially compute and physical access	isolation, external keys, emergency authority	fragmented fallback only	asset holders, then possibly AI system	generally absent unless failures stay bounded	Kernel or None
S17 discontinuity	new know-how, then next scarce complement	modular isolation, emergency reclassification	version fallback and institutional reset	discoverer, state, asset holder	post-shock redesign occurs before second shock	Kernel/Public or None
S18 fragmented blocs	jurisdiction, interoperability, FX	sovereign controls and reciprocal standards	bloc redundancy and local fallback	states, hyperscalers, large enterprises	diverse failures improve federation	Independent/Public

Explicit unknown register

The following unknowns are decision-relevant across more than one world:

1. Whether frontier capability remains coupled to capital scale or becomes reproducible at modest scale.
2. Whether cheap inference increases contestability or enables predatory bundling and default capture.
3. Whether open standards keep identity, telemetry, and settlement separable from hyperscaler control.
4. Whether local inference reaches parity on long-horizon autonomous work, not only bounded tasks.
5. What fraction of economically important outcomes can be fully and correctly specified.
6. Whether proof generation and independent verification stay simpler than generation.
7. Whether models can strategically manipulate evaluation, monitoring, or specification selection in real deployments.
8. Whether provider-native evidence receives judicial or regulatory presumptions after serious incidents.
9. Whether regulators can obtain enough compute, access, and expertise to challenge provider evidence.
10. Whether insurance can price correlated model, cloud, cyber, and physical tail risk without public backstops.
11. Whether agent delegation law can establish identity, authority, revocation, and non-repudiation across jurisdictions.
12. Whether central-bank settlement remains the final trust anchor for machine transactions or is partially displaced by closed provider rails.
13. Whether grid, chip, memory, and construction bottlenecks persist after algorithmic efficiency gains.
14. Whether embodied autonomy remains limited by physics and safety after digital reasoning becomes abundant.
15. Whether recursive AI research acceleration is possible and, if so, whether experiments and physical deployment remain rate limits.
16. Whether any monitoring system can remain epistemically and strategically independent of a superintelligence.
17. Whether governance adapts, is captured, is bypassed, or is replaced by emergency authority under fast change.

18. Whether plural states and enterprises maintain costly redundancy or converge on a dominant global stack.

Bottom line

No resistance family is universally durable. Technical controls are strongest when specifications and boundaries are narrow. Institutional controls are strongest when enforcement time matches capability change. Economic controls are strongest when harms and scarce inputs are priced rather than subsidized or externalized. Architectural controls are strongest when diversity is real rather than cosmetic. Strategic controls are strongest when monitors, principals, and fallback resources remain independent.

Cheap inference does not imply dispersed power. Open standards do not imply weak incumbents. Perfect verification does not imply correct objectives. Safe delegation does not imply the disappearance of legal principals. Provider-native evidence does not imply public legitimacy. Fast takeoff does not guarantee that all physical and cryptographic choke points disappear, but it can invalidate every control layer that depends on stable models, stable institutions, and historical calibration.

The survival space for a control layer is therefore conditional. It survives independently in plural, enterprise, slow-transition, and fragmented worlds. It becomes embedded or public infrastructure in dominant-provider, state, trusted-institution, and interface-obsolete worlds. It shrinks to a proof, identity, isolation, or settlement kernel in self-verifying, deceptive, fast-takeoff, and discontinuous worlds. In the hard fast-takeoff counterworld, it does not survive at all.

Round 2 cross-examination: incumbents and enclosure

Date: 2026-08-05

Scope and evidence discipline

This memo cross-examines whether a durable independent control layer can exist between AI models, agents, enterprise systems, identity, evidence, procurement, and payments. It does not assume that neutrality has economic value. It also does not assume that enclosure is inevitable. For each debate, it states the strongest case for the first world, the strongest opposite world, the causal sequence that would produce each one, material confounders, observable indicators, a falsifier, and the minimum survival conditions for a durable independent opportunity surface.

The analysis treats a control point as attacked if any incumbent can copy it, bundle it, cross-subsidize it, acquire it, deny its telemetry, prefer its own route, tie it to identity or payment, shape the standard, exploit committed spend, invoke regulation, or vertically foreclose access to context or evidence. Relevant incumbents include cloud and productivity platforms, model providers, enterprise application vendors, identity and security firms, payment networks, systems integrators, governments, and open-source ecosystems.

Evidence labels used below:

- Live means a generally available or user-accessible capability exists.
- Preview means a research preview, public preview, announced program, or proposed standard exists. It does not prove scaled use.
- Vendor-reported use means the source reports use or a numerical result, but the claim is not treated as independently audited.
- Government implementation means a law, award, procurement framework, or operative government program exists. A contract ceiling is not treated as actual spend.
- Inference means a causal conclusion derived from the cited primary evidence rather than a claim made by the source.

Baseline: where control naturally wants to live

The independent layer sits between stronger natural owners:

1. The front door owns task formulation and default routing.
2. The system of record owns context, permissions, and the meaning of an outcome.
3. The runtime owns traces, intermediate state, cost, latency, and failure telemetry.
4. The identity provider and resource server own enforceable authorization.
5. The marketplace and procurement channel own discovery, ranking, contracts, and committed-spend economics.
6. The payment network owns delegated payment credentials, fraud signals, disputes, and merchant acceptance.
7. The insurer owns risk capital, claims history, exclusions, and regulated underwriting.
8. The physical operator owns fleet data, maintenance, safety certification, and bodily-injury exposure.

This is why an independent layer cannot be evaluated only as software. Its power depends on rights granted by every adjacent owner.

Microsoft demonstrates the integrated attack most clearly. Work IQ continuously processes email, calendars, meetings, chats, files, people, collaboration patterns, and line-of-business systems; Microsoft says its APIs are the best way for agents to interact with Microsoft 365 data and apps. The APIs became generally available on June 16, 2026, actions remain auditable inside the tenant trust boundary, and pricing uses Copilot Credits. Microsoft Agent 365 became generally available May 1, 2026 inside a suite advertised at up to 15 percent incremental savings. GitHub made auto model selection the default and only choice for Free

and Student plans in June 2026 and removed all Gemini models from Copilot web a month earlier. These are live examples of context, routing, governance, billing, and model availability being joined by one owner ([Work IQ APIs](<https://www.microsoft.com/en-us/microsoft-365/blog/2026/06/02/announ...>), [Agent 365 bundle](<https://partner.microsoft.com/en-US/blog/article/agent-365-announcement>), [GitHub auto routing](<https://github.blog/changelog/2026-06-24-changes-to-model-selection-...>), [GitHub model removal](<https://github.blog/changelog/2026-05-20-updates-to-available-models-...>)).

AWS already joins prompt routing, runtime, identity, observability, marketplace, procurement, and deployment. Bedrock Intelligent Prompt Routing became generally available April 22, 2025. AgentCore stores traces, prompts, structured logs, and standard output in CloudWatch. AWS Marketplace launched AI Agents and Tools with hundreds of solutions, co-sell incentives, multiple delivery methods, and deployment into AgentCore or a customer VPC ([Bedrock routing](<https://aws.amazon.com/about-aws/whats-new/2025/04/amazon-bedrock-in-...>), [AgentCore observability](<https://aws.amazon.com/about-aws/whats-new/2026/07/amazon-bedrock-ag-...>), [AWS Marketplace](<https://aws.amazon.com/blogs/apn/aws-partner-guide-to-ai-agents-and-...>)).

Google calls Gemini Enterprise the single front door for AI in the workplace and combines models, a no-code workbench, agents, connectors, and central governance. It also originated A2A, originated AP2 with payment partners, operates Search and Workspace, and has moved Gemini into robotics. Salesforce completed its acquisition of Informatica on November 18, 2025 specifically to combine catalog, integration, governance, quality, privacy, metadata, and master data with Agentforce. ServiceNow completed the Moveworks acquisition on December 15, 2025 to create an AI-native employee front door. Oracle embeds its marketplace inside Fusion applications. These firms do not need a neutral control layer to become a new category. They can make it a feature of a larger control surface ([Gemini Enterprise]([https://cloud.google.com/blog/products/ai-machine-learning/introduci...](https://cloud.google.com/blog/products/ai-machine-learning/introduci-...)), [Salesforce and Informatica]([https://www.salesforce.com/uk/news/press-releases/2025/11/18/salesfo...](https://www.salesforce.com/uk/news/press-releases/2025/11/18/salesfo-...)), [ServiceNow and Moveworks]([https://newsroom.servicenow.com/press-releases/details/2025/ServiceN...](https://newsroom.servicenow.com/press-releases/details/2025/ServiceN-...)), [Oracle Agent Marketplace](<https://www.oracle.com/apac/news/announcement/ai-world-oracle-expand-...>)).

The debates below ask what would have to be true for that baseline not to determine the outcome.

Debate 1: Contestability versus enclosure

Strongest case for contestability

Agent ecosystems are structurally heterogeneous. A large enterprise may use Microsoft 365, Salesforce, ServiceNow, AWS, Google Cloud, SAP, Oracle, several frontier model providers, open-weight models, and internal systems at the same time. No single vendor has complete context or execution authority across all of them. Security teams also resist granting one application vendor unrestricted cross-system control. This fragmentation creates a plausible role for an independent policy, routing, evidence, or governance plane.

Open protocols reduce the cost of crossing those boundaries. A2A was transferred to the Linux Foundation in June 2025 with AWS, Cisco, Google, Microsoft, Salesforce, SAP, and ServiceNow as founding participants. Anthropic donated MCP to the Linux Foundation's Agentic AI Foundation in December 2025 with OpenAI, Block, Google, Microsoft, AWS, Cloudflare, and Bloomberg involved. The breadth of participants means interoperability is not merely a fringe demand ([A2A donation](<https://developers.googleblog.com/google-cloud-donates-a2a-to-linux-...>), [MCP donation](<https://www.anthropic.com/news/donating-the-model-context-protocol-a-...>)).

Regulators and sovereign buyers can also require multi-vendor operation. The European Commission's cloud sovereignty procurement used 48 criteria spanning strategic, legal, data and AI, operational, supply-chain, technological, security, compliance, and environmental objectives. A buyer using such a framework may prefer a separable control plane that reduces dependence on any one foreign cloud or application vendor ([EU Cloud Sovereignty Framework]([https://commission.europa.eu/news-and-media/news/sovereign-cloud-fra...](https://commission.europa.eu/news-and-media/news/sovereign-cloud-fra-...))).

Strongest opposite world: enclosure wins

Heterogeneity does not guarantee an independent winner. The same enterprise can multi-home while every workload remains governed inside the dominant system of record. Microsoft governs Microsoft work, Salesforce governs CRM work, ServiceNow governs service workflows, and AWS governs AWS runtime. The enterprise tolerates several enclosed domains connected by thin protocols. It buys integration from a systems integrator rather than inserting a new control plane with broad privileges.

Each incumbent can expose enough interoperability to satisfy buyers while retaining the profitable layers: default routing, identity, audit, ranking, billing, and outcome semantics. It can copy visible features, provide them at no incremental price, and make the independent layer bear integration and liability costs. The independent company becomes a dashboard over native controls, not a control point.

Causal sequence

Contestability sequence:

1. Enterprises maintain genuinely mixed stacks.
2. Cross-vendor failures create a budget for independent control.
3. Systems of record expose complete, timely, permission-faithful events.
4. Independent policies can be enforced, not merely displayed.
5. Procurement recognizes the independent layer as an accountable vendor.
6. Switching among execution providers remains operationally cheap.
7. The independent layer accumulates evidence that improves outcomes across vendors without violating data rights.

Enclosure sequence:

1. Incumbents adopt common protocols to reduce buyer resistance.
2. They bundle native governance, identity, and observability.
3. They expose coarse cross-vendor inventory but reserve native enforcement.
4. They rank or route toward native agents and preferred models.
5. They tie economics to enterprise agreements, cloud credits, or application seats.
6. Independent vendors lose distribution, telemetry, and gross margin.
7. Remaining specialists are acquired or become service partners.

Confounders

- A multi-vendor logo inventory may mask one vendor owning most consequential workloads.
- A cross-platform dashboard may appear neutral while all kill switches remain native.
- Protocol support may cover invocation but not authorization, evidence completeness, or dispute rights.
- Buyer complaints about lock-in do not prove willingness to fund a separate vendor.
- Security concerns can favor independence or favor one accountable incumbent.

Indicators

- Percentage of agent actions enforceable from outside the native platform.
- Completeness and latency of cross-platform event export.
- Share of policies that can be exported and imported without semantic loss.
- Frequency of native-only capabilities or privileged APIs.
- Cross-vendor switching time, including identity, memory, and evidence migration.
- Independent control-plane attach rate outside consulting-led deployments.
- Acquisition rate of successful neutral-layer vendors.

Falsifier

Contestability is falsified if enterprises continue to multi-home models but more than 80 percent of consequential agent actions remain governed, logged, and approved exclusively by the native system of record or cloud after three procurement cycles. Enclosure is falsified if independent policy engines routinely block or reroute actions across major systems with complete evidence and without bespoke integration.

Survival conditions

An independent surface survives only if it has enforceable external hooks, complete evidence rights, policy portability, buyer-recognized accountability, and a job that crosses several systems of record. Merely being vendor-neutral is insufficient.

Debate 2: Open standards as distribution versus commoditization

Strongest case for distribution

Open standards can let one implementation reach many clients. MCP replaces one-off tool connectors; A2A supplies common agent discovery and communication; AP2 and x402 attempt to standardize agent payment messages. A small vendor can publish once and become available in multiple ecosystems. Open-source SDKs lower customer diligence and can create community distribution. Foundation governance can reduce the fear that one model vendor will revoke the protocol.

The mechanism is strongest when the independent vendor owns a service above the standard rather than the adapter itself. A protocol can widen the addressable market for a verifier, data provider, specialized workflow, or evidence service whose value is not the connection syntax.

Strongest opposite world: commoditization and hyperscaler reinforcement

The protocol can erase the startup's original scarcity. Connector engineering, tool description, discovery, and message translation become reusable public goods. Managed hosting, identity, observability, billing, and procurement then determine adoption. Those layers favor hyperscalers.

AWS Marketplace already filters for MCP and A2A support, offers native deployment paths, and applies co-sell incentives. Microsoft makes Work IQ available through MCP while retaining Entra authentication, tenant policy, Microsoft 365 context, and Copilot Credits. Open syntax increases the number of complements that can consume incumbent infrastructure.

Standards also have governance surfaces beyond the base text: registries, extensions, default profiles, conformance suites, trust roots, client behavior, and release cadence. C2PA illustrates this. Its specification is open, but practical trust depends on an official trust list, certificate policy, conforming products, certification authorities, and platform preservation of metadata. The conformance program launched in mid-2025, while OpenAI explicitly notes that C2PA metadata can be removed by screenshots or platform processing ([C2PA conformance](<https://c2pa.org/conformance/>), [OpenAI on C2PA limits](<https://help.openai.com/en/articles/8912793-c2pa-in-dall-e-3%23.otf>)).

Causal sequence

Distribution sequence:

1. A standard reaches several major clients.
2. Implementations are genuinely substitutable.
3. Identity, authorization, and evidence travel with the interaction.
4. Registries are federated and ranking is not controlled by one client.
5. Customers can self-host or select among managed providers.
6. The independent service retains unique data, judgment, or regulated authority above the standard.

Commoditization sequence:

1. The standard eliminates custom integration work.
2. Hyperscalers ship managed clients, registries, gateways, and hosting.
3. Buyers select the deployment already covered by identity and committed spend.

4. Community contributions improve the incumbent-hosted implementation.
5. Adapter margins fall toward zero.
6. The remaining differentiated service is acquired, bundled, or moved into the system of record.

Confounders

- Protocol adoption counts do not reveal concentration of traffic.
- Open-source downloads do not reveal production use or revenue.
- Multiple implementations may share one registry or identity root.
- A nominally neutral foundation may still depend on a few firms for engineering and testing.
- Self-hosting may be legally possible but operationally uneconomic.

Indicators

- Share of protocol traffic handled by the top three managed runtimes.
- Number of independent clients, registries, and conformance providers with material use.
- Switching cost between hosted implementations.
- Revenue retained by protocol-native specialists versus hosting and marketplace owners.
- Extension usage and whether extensions are portable.
- Concentration of maintainers, funding, test infrastructure, and steering rights.

Falsifier

The distribution thesis is falsified if protocol adoption grows while independent vendor revenue share and registry diversity decline. The commoditization thesis is falsified if several independent providers sustain pricing power because customers can change runtime, identity, and marketplace without losing evidence or reputation.

Survival conditions

The independent company must own scarce value above syntax: authoritative data, legally recognized attestation, difficult domain judgment, risk capital, or a user relationship. The protocol must make demand portable, not merely supply compatible.

Debate 3: Regulation as opening versus concentration

Strongest case for opening

Regulation can force portability, transparency, non-discrimination, audit access, and multi-vendor procurement. It can recognize independent auditors and require evidence the native platform would otherwise withhold. Sovereignty requirements can create demand for local providers and reduce dependence on US hyperscalers. Open-source exemptions can reduce compliance burdens for smaller actors.

The European Commission's GPAI guidance provides some open-source exemptions and the Code of Practice offers a common compliance route. The US Office of Management and Budget issued M-25-21 and M-25-22 in April 2025 to accelerate federal AI use and procurement while supporting a competitive marketplace. NIST's February 2026 AI Agent Standards Initiative explicitly seeks industry-led standards, community-led open protocols, and research into identity and security ([EU GPAI guidance](<https://digital-strategy.ec.europa.eu/en/policies/guidelines-gpai-pr...>), [White House procurement policy](<https://www.whitehouse.gov/releases/2025/04/white-house-releases-new...>), [NIST initiative](<https://www.nist.gov/news-events/news/2026/02/announcing-ai-agent-st...>)).

Strongest opposite world: regulatory concentration

Compliance has fixed costs: legal interpretation, documentation, security controls, red teams, model evaluation, incident reporting, insurance, local operations, and certification. Large firms spread these costs across more revenue. Voluntary codes can become de facto safe harbors shaped by firms with the staff to participate. Procurement frameworks can admit only vendors with existing certifications, balance sheets, and government relationships.

The EU GPAI Code's signatories include Amazon, Anthropic, Google, IBM, Microsoft, OpenAI, and ServiceNow. Signatories receive greater predictability and reduced administrative burden. That may improve compliance while also favoring firms able to shape and implement the code. The US DoD's July 2025 frontier-model awards diversified across Anthropic, Google, OpenAI, and xAI, but each had a \$200 million ceiling and the set still consisted of four frontier firms. A ceiling is not proof of actual spend, but the procurement boundary is evidence of oligopoly selection ([EU GPAI Code](<https://digital-strategy.ec.europa.eu/en/policies/contents-code-gpai>), [DoD awards](<https://www.ai.mil/Latest/News-Press/PR-View/Article/4242822/cdao-an...>)).

Causal sequence

Opening sequence:

1. Rules specify outcomes and portability rather than a named implementation.
2. Compliance artifacts are reusable and machine-readable.
3. Independent labs, auditors, and trust roots are recognized.
4. Procurement lots are small enough for specialist participation.
5. Liability is proportionate to the actor's actual control.
6. Enforcement targets discrimination and incomplete evidence.

Concentration sequence:

1. Incidents create political demand for assurance.
2. Standards and procurement require expensive certifications and continuous monitoring.
3. Safe harbors favor code signatories and incumbent controls.
4. Buyers avoid untested vendors because liability is unclear.
5. Large firms bundle compliance and indemnity.
6. Specialists become subcontractors to hyperscalers or systems integrators.

Confounders

- More registered vendors can coexist with spending concentration.
- A local sovereign provider may be a reseller of foreign technology.
- Open-source exemptions may help development but not high-risk deployment.
- Compliance transparency may look burdensome initially but lower long-term switching costs.
- Government awards may fund experimentation rather than production lock-in.

Indicators

- Compliance cost as a percentage of revenue by vendor size.
- Number and share of independent auditors and certified providers.
- Procurement spending concentration, not only award count.
- Frequency of portability and non-discrimination clauses.
- Time and cost for a new entrant to achieve required certification.
- Share of sovereign contracts using technology controlled outside the jurisdiction.

Falsifier

The opening thesis is falsified if nominal interoperability rules coincide with rising procurement concentration and entrant certification time. The concentration thesis is falsified if small independent providers repeatedly win direct, material contracts using reusable compliance evidence without incumbent sponsorship.

Survival conditions

Durable independent surfaces require proportionate liability, reusable compliance artifacts, recognized independent assurance, procurement access, and rules that attach obligations to actual control rather than company size or category label.

Debate 4: Neutral cross-vendor layer versus systems of record

Strongest case for the neutral layer

No system of record sees the entire task. A customer request may begin in Salesforce, require a ServiceNow change, access Microsoft documents, invoke an AWS-hosted model, update SAP, and trigger a payment. Cross-system constraints can conflict. A neutral layer can assemble an end-to-end execution graph, apply one policy, compare models, capture a complete trace, and detect failures that no local system sees.

The neutral layer may also be more credible when the enterprise does not want the party selling execution to grade its own work. It can benchmark providers, preserve exit options, and give security teams one cross-vendor view.

Strongest opposite world: systems of record remain sovereign

The system of record owns the authoritative state transition. Salesforce decides whether an opportunity changed. ServiceNow decides whether an incident closed. SAP or Oracle decides whether an invoice posted. GitHub and CI decide whether a code change merged and passed. The neutral layer can observe, but it cannot define validity without the record owner's semantics and permissions.

Salesforce stores prompts, responses, and trust signals in Data 360, applies native access controls, and keeps external model providers under zero-data-retention terms. Microsoft keeps Work IQ context and actions inside the Microsoft 365 tenant trust boundary. OpenAI exposes compliance logs contractually and retains them for 30 days unless the customer continuously exports them. The evidence substrate is partitioned by owner ([Salesforce Trust Layer](https://help.salesforce.com/s/articleView?id=sf.copilot_trust.htm&la...), [OpenAI Compliance Platform](<https://help.openai.com/en/articles/9261474-openai-compliance-platfo...>)).

Causal sequence

Neutral-layer sequence:

1. Important workflows routinely span several records.
2. Each record exposes complete events, permissions, and transaction identifiers.
3. The enterprise defines cross-system outcome contracts.
4. The neutral layer can block or compensate failed transitions.
5. Evidence is signed and retained outside native platforms.
6. Buyers hold the neutral layer accountable for cross-system completion.

System-of-record sequence:

1. Each incumbent exposes a connector and local agent.
2. Cross-system work is decomposed into native subtransactions.
3. Each native system retains authorization and evidence.
4. A suite or SI orchestrates handoffs without ceding control.
5. The neutral layer receives delayed or lossy summaries.
6. Its outcome model cannot distinguish omitted data from genuine success.

Confounders

- Cross-system visibility is not cross-system authority.
- A replayable trace can still omit private native decisions.
- A technically correct transaction may fail the business objective.

- The enterprise may own the records but lack internal agreement on outcome definitions.
- Systems integrators can simulate neutrality while remaining commercially aligned with incumbents.

Indicators

- Percentage of workflows with more than one authoritative record.
- Availability of stable cross-system transaction identifiers.
- Completeness of pre-action and post-action evidence.
- Ability to reverse or compensate a native action externally.
- Number of disputes in which neutral evidence is accepted as authoritative.
- Time required to update semantics when a system of record changes.

Falsifier

The neutral-layer thesis is falsified if cross-system workflows grow but authoritative verification remains local and independent traces cannot resolve disputes. The systems-of-record thesis is falsified if enterprises standardize signed outcome contracts that let an external plane verify and govern state transitions without bespoke adapters.

Survival conditions

The independent layer needs a cross-system job, signed evidence, transaction identity, enforcement or compensation rights, and explicit contractual recognition of its outcome record. Read-only aggregation is not enough.

Debate 5: Multi-homing versus defaults and committed spend

Strongest case for multi-homing

Models vary by task, price, latency, context, modality, geography, and policy. Enterprises already use several clouds and SaaS products. A2A and MCP reduce integration cost. Model diversity can reduce outage and bargaining risk. If no model dominates every task, rational buyers should preserve routing choice.

Microsoft itself says Copilot is model-diverse and supports OpenAI and Anthropic. AWS Bedrock exposes multiple model families. A multi-homing layer can exploit price-performance dispersion and provide resilience.

Strongest opposite world: defaults dominate economic routing

Technical access to many models is not equal access. The platform chooses which models appear, which are eligible for auto mode, how premium multipliers work, which regions serve them, and which enterprise controls apply. GitHub's model removal is direct evidence that a platform can narrow choice. AWS Intelligent Prompt Routing selects only two models inside one family and charges through Bedrock. Microsoft Work IQ prices context and tools in Copilot Credits.

Committed cloud spend changes the objective function. A nominally cheaper external route may be economically worse if native consumption retires an enterprise commitment, uses already licensed controls, and avoids a new vendor review. Defaults also compound telemetry: the default gets more traffic, learns faster, gains a better reputation, and becomes safer to procure.

Causal sequence

Multi-homing sequence:

1. Model performance remains meaningfully heterogeneous.
2. Routing gains exceed integration, security, and egress costs.
3. Buyers can observe selected models and override defaults.
4. Context and evidence are portable across model providers.
5. Procurement treats models as substitutable suppliers.
6. No platform can remove a provider without a viable bypass.

Default-enclosure sequence:

1. Platforms offer many models to satisfy choice demands.
2. Auto mode becomes the normal interface.
3. The platform optimizes for availability, cost, policy, and commercial preference.
4. Committed spend and bundled controls make native routes cheaper in total.
5. Low-volume alternatives receive less telemetry and investment.
6. The platform retires or restricts them with limited user resistance.

Confounders

- A model picker can create the appearance of competition while auto handles most traffic.
- Reported model usage may not reveal auto versus manual selection.
- Lower token price may be offset by context, tool, egress, and review costs.
- Enterprise commitments are often confidential.
- Regulatory or sovereignty constraints can force multi-homing independent of economics.

Indicators

- Share of traffic in auto versus manually pinned modes.
- Model exposure, selection, and override rates.
- Frequency and notice period of provider removal.
- Total route cost after credits, egress, controls, and human review.
- Share of spend used to retire cloud or suite commitments.
- Ability to reproduce a run after switching providers.

Falsifier

The multi-homing thesis is falsified if model catalogs expand while more than 80 percent of traffic flows through platform auto modes and effective provider concentration rises. The defaults thesis is falsified if customers routinely override defaults and move material traffic without losing context, evidence, or committed-spend economics.

Survival conditions

An independent router must control enough demand to negotiate economics, expose reason codes, preserve context and evidence across routes, and demonstrate gains after all switching costs. Model count alone does not create a business.

Debate 6: Independent verification versus bundled attestation

Strongest case for independent verification

Execution providers face a conflict when they grade their own output. Regulated buyers, insurers, courts, and counterparties may require a verifier with separate incentives. Cross-vendor workflows need a common evidence chain. Independent verification can specialize in adversarial testing, calibration, provenance, and dispute resolution.

C2PA's conformance program recognizes separate generator, validator, and certification roles. NIST's agent standards work identifies independent identity, authorization, and security evaluation as open problems. Insurers also need evidence not controlled solely by the insured platform.

Strongest opposite world: bundled attestation wins

The platform sees the richest ground truth and can attest at negligible marginal cost. AWS records AgentCore prompts, traces, logs, and outputs. Microsoft Purview records agent activity and connects it to eDiscovery and risk tools. Salesforce logs prompts, responses, and trust signals. Native identity and authorization also reveal who was permitted to act.

An independent verifier receives exported data after the fact and cannot prove that omitted events never occurred. If the verifier requires deep access, it becomes a new high-privilege risk. Buyers may prefer a platform attestation backed by an existing

contract, certification, and indemnity.

Watermarking demonstrates the limitation. Google announced SynthID Detector in May 2025, initially for early testers, to detect content created with Google's tools. OpenAI added C2PA and later SynthID, but its public material acknowledges metadata loss. A verifier tied to the generator may have the best detector, while a third party lacks the private signal ([Google SynthID Detector](<https://blog.google/innovation-and-ai/products/google-synthid-ai-con...>), [OpenAI provenance](<https://openai.com/index/advancing-content-provenance/>)).

Causal sequence

Independent-verification sequence:

1. Buyers identify a conflict in self-attestation.
2. Platforms emit signed, complete, standardized receipts.
3. Multiple trust roots and labs are accepted.
4. Independent verifiers gain access to outcome ground truth.
5. Regulators, insurers, and courts recognize their findings.
6. Liability attaches to false or negligent verification.

Bundled-attestation sequence:

1. Platforms integrate logging, identity, policy, and attestation.
2. Enterprise agreements and certifications make native evidence acceptable.
3. Exported evidence remains incomplete or delayed.
4. Independent verifiers cannot improve accuracy enough to justify cost and risk.
5. Assurance becomes a feature or an audit service attached to the platform.

Confounders

- Organizational independence does not ensure technical independence.
- A third-party auditor may depend on vendor-supplied test environments.
- Native logs may be complete for runtime behavior but not business outcomes.
- Cryptographic integrity does not prove semantic truth or completeness.
- Insurer acceptance can reflect commercial partnerships rather than verifier quality.

Indicators

- Percentage of actions carrying independently verifiable signed receipts.
- Number and concentration of accepted trust roots and labs.
- False-positive and false-negative rates on open test sets.
- Evidence survival across platform transformations.
- Disputes resolved against a platform's native attestation.
- Premium or procurement benefit from independent verification.

Falsifier

Independent verification is falsified as a broad category if platforms retain exclusive ground truth and external findings do not alter procurement, insurance, or disputes. Bundled attestation is falsified if regulators and insurers routinely reject native-only evidence and require accepted third-party verification.

Survival conditions

The verifier needs mandatory evidence access, recognized authority, measurable error advantages, multiple roots, and liability proportionate to its role. A generic verification API without ground truth is not durable.

Debate 7: Which worlds leave durable independent opportunity surfaces?

Strongest case that durable surfaces remain

Independent companies can survive where authority is deliberately separated, where several incumbents need a shared utility they do not trust a rival to own, or where domain complexity defeats horizontal bundling. Historical analogues include payment gateways, identity providers, observability vendors, certificate authorities, exchanges, credit bureaus, and insurance intermediaries. Their durability came from contracts, regulation, network participation, or authoritative data, not abstract neutrality.

The likely independent surfaces are therefore conditional roles:

1. A cross-system evidence utility whose signed receipts are contractually recognized by record owners, insurers, and regulators.
2. A domain verifier with exclusive or authoritative outcome data and measurable error advantages.
3. A regulated risk carrier or broker that can price narrow agent failures and require portable telemetry.
4. A sovereign or sector-specific control plane with procurement protection and local legal authority.
5. A user-controlled identity, mandate, and payment wallet if platforms must honor it.
6. A specialized workflow network in which no single participant owns both sides and the intermediary governs shared rules.
7. A physical safety or certification layer recognized across OEMs and jurisdictions.

Strongest opposite world: every surface is absorbed

Clouds absorb runtime and observability. Work suites absorb context and governance. Systems of record absorb workflow and evidence. IAM firms absorb authorization. Payment networks absorb delegated purchasing. Insurers absorb risk scoring. SIs absorb integration. Governments certify a small set of vertically integrated providers. Open source removes connector margins. Any successful horizontal startup is acquired before it establishes independent demand.

Current acquisitions show the direction: ServiceNow bought Moveworks for the front door, Salesforce bought Informatica for trusted context, Stripe bought Bridge for stablecoin infrastructure, and Palo Alto Networks bought CyberArk for identity. These transactions do not prove all independent surfaces disappear, but they establish that incumbents will acquire adjacent control rather than leave it neutral ([Stripe and Bridge](<https://stripe.com/en-dk/newsroom/news/stripe-completes-bridge-acqui...>), [Palo Alto and CyberArk](<https://investors.paloaltonetworks.com/node/20181/pdf>)).

Causal sequence

Durable-independent sequence:

1. Several powerful parties require a shared function.
2. None will accept a direct rival as owner.
3. The intermediary gains contractual or regulatory authority.
4. Data and evidence rights are durable and portable.
5. The function has liability, judgment, or network rules that are not cheap to copy.
6. Multi-homing is operationally real.
7. Acquisition would undermine the trust or legal status that creates value.

Absorption sequence:

1. The startup proves demand at a visible seam.
2. Incumbents copy basic features and restrict privileged access.
3. Bundling reduces willingness to pay.
4. Procurement and committed spend shift demand to the suite.
5. The startup accepts acquisition or narrows into services.
6. The buyer integrates its data and distribution, closing the remaining gap.

Confounders

- Acquisition can validate a useful function while disproving an independent category.
- A large revenue business may still lack strategic control if one platform supplies demand.
- Regulatory recognition can protect independence or turn the firm into a licensed utility with capped economics.
- Sector specialization can create durable data but a small market.
- A consortium can prevent ownership concentration while producing slow governance and weak execution.

Indicators

- Revenue concentration by platform, system of record, and channel.
- Direct customer demand versus marketplace or SI-mediated demand.
- Contractual rights to retain and use evidence after switching.
- Number of record owners that accept the intermediary's authority.
- Gross margin after integration, compliance, insurance, and support costs.
- Frequency of incumbent replication, access restriction, and acquisition offers.
- Whether acquisition would destroy regulatory recognition or counterparty trust.

Falsifier

The durable-surface thesis is falsified if every scaled neutral provider remains dependent on one distribution platform, lacks authoritative evidence, and is acquired or commoditized within two platform release cycles. The absorption thesis is falsified if several independent firms preserve direct demand, cross-platform enforcement, portable evidence, and pricing power through repeated incumbent bundling attempts.

Survival conditions

A durable independent opportunity surface needs at least four of the following:

1. Authority recognized by multiple powerful counterparties.
2. Durable rights to complete data or evidence.
3. A cross-system job no record owner can solve alone.
4. Independent distribution or user ownership of demand.
5. Liability or regulated status that cannot be copied as a feature.
6. Real multi-homing across identity, policy, context, and billing.
7. A trust structure that acquisition would impair.
8. Domain-specific outcome semantics and longitudinal history.

It also needs a credible answer to all of these attacks: free bundling, restricted telemetry, preferred routing, committed spend, native attestation, standards capture, regulatory fixed cost, and acquisition.

Cross-debate counterworlds

World A: Open protocol, closed economics

MCP and A2A are universal. Every agent can technically call every tool. Three managed runtimes handle most traffic because they own identity, logs, marketplaces, and cloud commitments. Startups sell specialized tools but not horizontal control. This world has high compatibility and high enclosure at once.

Falsifier: material traffic and revenue remain distributed across independent runtimes and registries.

World B: Federated enterprise control

Large buyers require signed receipts, portable policy, external kill switches, and model-routing transparency. Several independent governance and evidence firms survive because no system of record sees the complete workflow. Systems integrators implement them but do not own their authority.

Falsifier: record owners refuse external enforcement or buyers waive the requirements for suite discounts.

World C: Sovereign oligopolies

Governments require local operations, data control, certified models, and national payment rails. A few local or partnered providers win. Foreign hyperscalers remain through sovereign subsidiaries or licensed technology. Independence exists nationally but not as a global neutral layer.

Falsifier: procurement awards fragment broadly and technology ownership is genuinely local and portable.

World D: Systems-of-record federation

Salesforce, ServiceNow, SAP, Oracle, Microsoft, and cloud providers keep native authority. A2A coordinates handoffs. Buyers use several vendors but no independent control plane. Cross-system incidents are handled by SIs and contract processes.

Falsifier: cross-system failure costs force buyers to recognize an external evidence and enforcement authority.

World E: Insurance-governed agents

Insurers and reinsurers specify accepted controls, telemetry, models, and escalation. High-stakes buyers follow the insurance standard. Independent verifiers survive only if carriers recognize them. Munich Re states it has written AI insurance since 2018 and now offers performance and GenAI liability coverage; HSB launched SMB AI liability insurance in March 2026. These are live insurance capabilities, not proof that agent-specific underwriting has reached scale ([Munich Re AI insurance](<https://www.munichre.com/en/solutions/for-industry-clients/insure-ai...>), [HSB AI liability](<https://www.munichre.com/hsb/en/press-and-publications/press-release...>)).

Falsifier: insurers accept native platform attestations without differentiated premiums or external verification.

World F: User-owned mandates

Identity, memory, purchasing constraints, and payment authorization live in a user-controlled wallet. Agents and merchants compete for execution. Payment and identity protocols are honored across front doors. This creates a durable independent personal-agent and mandate layer.

Falsifier: OS, browser, IdP, and payment defaults prevent a user mandate from moving intact across platforms.

World G: Physical control oligopoly

Robot models and protocols are open, but NVIDIA, Google, Amazon, OEMs, and fleet operators retain simulation, chips, fleet data, safety control, certification, and maintenance. Google Gemini Robotics directly outputs physical actions; NVIDIA combines an open GR00T model with simulation and compute; Amazon reports one million deployed robots. Independent opportunity remains in certified components or narrow domains, not the general control plane ([Gemini Robotics](<https://deepmind.google/blog/gemini-robotics-brings-ai-into-the-phys...>), [NVIDIA GR00T](<https://nvidianews.nvidia.com/news/nvidia-isaac-gr00t-n1-open-humano...>), [Amazon robotics](<https://www.aboutamazon.com/news/operations/amazon-million-robots-ai...>)).

Falsifier: portable controllers achieve equivalent safety and performance across fleets without privileged OEM data.

Final cross-examination

The strongest enclosure argument is not that one incumbent owns everything. It is that every profitable layer has a stronger adjacent owner, while the independent company must obtain permission from all of them. Open standards can increase interoperability and still concentrate traffic. Regulation can require openness and still concentrate procurement. Multi-homing can expand catalogs and still leave defaults in control. Independent verification can be desirable and still lack ground truth. A neutral layer can span systems and still lack authority over any state transition.

The strongest contestability argument is also not neutrality. It is institutional separation backed by enforceable rights. Independent surfaces endure when multiple powerful parties need a shared function, no rival is acceptable as owner, complete evidence is portable, the intermediary has recognized authority, and acquisition would damage the trust or legal status that creates its value.

The decisive empirical question is therefore not whether the ecosystem uses open protocols or many models. It is who controls the default, who can enforce policy, who holds complete telemetry, whose evidence resolves a dispute, whose credits fund the transaction, and who bears the loss when the outcome fails.

No product, company, or investment conclusion is drawn from this analysis.

Round 2 cross-examination: enterprise authority

Date: 2026-08-05

Scope: This memo examines how enterprises, governments, and institutions may fund, centralize, delegate to, resist, insure, consolidate, or abandon autonomous systems as capability rises. It does not recommend a product, company, investment, or policy. It separates current facts from causal inferences and unresolved unknowns.

Labels used below:

- [Fact] means a claim supported by a linked primary or official source.
- [Inference] means a causal interpretation of facts, not an observed certainty.
- [Unknown] means an unresolved variable capable of reversing the scenario.

Evidence baseline and measurement discipline

[Fact] Enterprise announcements are large but use different denominators. Microsoft reported more than 100 million monthly active Copilot users across commercial and consumer products and more than 230,000 organizations using Copilot Studio in its FY2025 annual report. Salesforce reported 5,000 Agentforce deals, including more than 3,000 paid deals, and \$900 million of Data Cloud and AI annual recurring revenue on February 26, 2025. Accenture reported \$5.9 billion of generative AI bookings in FY2025 while total bookings declined 1 percent. None of these measures, alone, proves that autonomous work reduced total cost, changed organization design, or produced net operating income. Sources: [Microsoft FY2025 annual report](<https://www.microsoft.com/investor/reports/ar25/index.html>), [Salesforce FY2025 results](https://investor.salesforce.com/files/doc_financials/2025/q4/CRM-Q4-...), and [Accenture FY2025 10-K](<https://www.sec.gov/Archives/edgar/data/1467373/000146737325000217/a...>).

[Fact] Controlled and quasi-controlled labor evidence is narrower and more conditional. A field experiment across 66 firms and 7,137 knowledge workers found that active treated users spent about two fewer hours per week on email, but the researchers detected no material change in the quantity or composition of tasks from individual-level access. A customer support study found roughly 14 percent higher productivity, concentrated among less experienced workers, with negligible or slightly negative effects for the most experienced workers. The International Labour Organization's June 1, 2026 review found real but uneven and often unverified productivity effects; time savings generally had not yet translated into measured output, earnings, or employment. Sources: [NBER working paper 33795](<https://www.nber.org/papers/w33795>), [NBER customer support summary](<https://www.nber.org/digest/20236/measuring-productivity-impact-gene...>), and [ILO empirical review](<https://www.ilo.org/publications/impact-genai-jobs-productivity-and-...>).

[Fact] Budget migration is nevertheless visible. Accenture disclosed \$615 million of FY2025 business optimization costs, approximately 77,000 AI and data practitioners, and more than 550,000 people trained in generative AI fundamentals. Oracle disclosed a FY2026 restructuring plan with estimated costs up to \$2.1 billion involving AI integration and recorded \$1.8 billion during the year. These filings demonstrate reallocation and restructuring, not causal proof that AI generated the resulting benefits. Sources: [Accenture FY2025 10-K](<https://www.sec.gov/Archives/edgar/data/1467373/000146737325000217/a...>) and [Oracle FY2026 10-K](<https://www.sec.gov/Archives/edgar/data/1341439/000119312526277521/o...>).

[Inference] The central measurement rule for this memo is therefore:

seat or deal -> active use -> completed workflow -> accepted outcome -> captured economic value
-> durable institutional legitimacy

Each arrow can fail. A control system may improve execution without capturing a budget. A buyer may capture a budget without granting external authority. A technically valid outcome may still be legally unauthorized, contractually incomplete, or politically illegitimate.

Debate 1: External authority versus customer-operated policy

Strongest case for an external authority layer

[Inference] An external authority layer gains power when autonomous work crosses multiple model, cloud, application, and payment providers. No single application then sees the complete trajectory. A cross-provider authority can compare model performance, enforce a common policy, accumulate loss history, detect correlated failure, and offer a consistent audit surface. The value is not just lower inference price. It is the ability to say which identity acted, under which delegation, against which policy, with which evidence, and with what recovery behavior across provider boundaries.

[Fact] Existing public policy already recognizes the value of cross-functional and independent controls. US OMB Memorandum M-25-22, issued April 3, 2025, requires agencies to consider data and model portability, vendor lock-in, independent evaluation, ongoing monitoring, cost justification, and sunset criteria in AI acquisition. It also warns that vendors may use AI in contract performance without the government anticipating it. Source: [OMB M-25-22](<https://www.whitehouse.gov/wp-content/uploads/2025/02/M-25-22-Drivin...>).

[Fact] Financial regulators likewise recognize that concentrated providers can become systemically important. The EU Digital Operational Resilience Act has applied since January 17, 2025. European supervisory authorities designated critical ICT third-party providers in November 2025 based partly on systemic importance and substitutability. Direct provider oversight complements, rather than replaces, each financial entity's responsibility. Sources: [EBA DORA oversight](<https://www.eba.europa.eu/activities/direct-supervision-and-oversigh...>) and [critical-provider designation](<https://www.eba.europa.eu/publications-and-media/press-releases/euro...>).

[Inference] The strongest external-authority world therefore resembles a licensed utility. Enterprises delegate parts of identity, policy evaluation, evidence normalization, and incident coordination to a small number of supervised control points. Regulators accept the chokepoint because it is more examinable than thousands of local agents. Insurers accept it because loss telemetry is standardized. Enterprises accept it because switching among models and applications remains possible above the common authority plane.

Opposite world: customer-operated policy remains sovereign

[Inference] The opposite world begins from a different premise: the party carrying legal responsibility needs operational control. Customer context includes confidential data, internal duties, worker agreements, jurisdictional rules, risk appetite, and emergency authority that an external provider cannot fully infer. In this world, an outside control point is another critical vendor, another subpoena surface, another concentration risk, and potentially an unacceptable observer of business intent.

[Fact] The EU AI Act places obligations on both providers and deployers of high-risk systems, including competent human oversight and log retention. It does not allow a deployer to transfer its obligations simply by buying a controlled service. Source: [EU AI Act](<https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=celex%3A32024R1689>).

[Fact] Sovereignty procurement is already moving beyond data residency. In April 2026, the European Commission awarded a EUR 180 million sovereign-cloud procurement to four providers using 48 criteria spanning legal and jurisdictional control, data and AI, operations, supply chain, technology, security, and compliance. Source: [European Commission Sovereign Cloud Framework, June 1, 2026](<https://commission.europa.eu/news-and-media/news/sovereign-cloud-fra...>).

[Inference] In the customer-sovereign world, policy runs inside the customer's cloud, network, or regulated enclave. External providers can supply models, assurance components, or updates, but the customer controls identity, keys, logs, policy, and final authorization. Evidence export is selective. A bank, defense organization, hospital, or government agency may tolerate an external model more readily than an external authority that observes every execution decision.

Causal play-by-play

1. Local teams adopt multiple agents and models because capability differs by task.
2. Security and legal teams discover that no application has end-to-end visibility.
3. A central enterprise policy budget emerges under the CIO, CISO, or risk office.
4. Procurement tests two architectures: externally operated authority and customer-hosted policy.

5. If regulators, insurers, and customers accept standardized external evidence, the authority layer centralizes and may become a critical provider.
6. If jurisdiction, confidentiality, or accountability dominates, policy stays customer-operated while evidence formats standardize externally.
7. A hybrid equilibrium is possible: customer-side enforcement with third-party certification, loss aggregation, and dispute support.

Hidden confounders

- A low external price may be subsidized by cloud or model consumption.
- A customer-hosted architecture may still depend on foreign code, silicon, or update channels.
- Provider-native integration can be mistaken for superior policy quality.
- An external authority may see only routed work; local scripts and shadow agents remain invisible.
- Regulatory acceptance in one sector or jurisdiction may not transfer to another.
- The legal defendant and the technical policy operator may be different entities.

Discriminating indicators

- Share of enterprise contracts requiring customer-controlled keys and customer-side enforcement.
- Procurement requirements for evidence portability and provider-independent replay.
- Direct regulatory designation of AI control or identity providers as critical third parties.
- Percentage of governed workflows crossing more than one model or application provider.
- Frequency with which customers permit external retention of prompts, traces, or outcome evidence.
- Insurer pricing differences between provider-operated and customer-operated controls.

Falsifier

[Inference] The independent external-authority thesis is falsified if enterprises consistently accept integrated provider controls, rarely use cross-provider execution, refuse external telemetry, and do not fund a separate authority budget. The pure customer-sovereignty thesis is falsified if regulators and insurers require standardized external supervision and customers cannot operate equivalent controls economically.

Conditionally surviving opportunity categories

The categories that remain economically possible across both worlds are customer-hosted policy enforcement, portable evidence schemas, third-party certification, critical-provider supervision, key and identity custody, jurisdiction-aware execution, and exit or replay infrastructure. Which category captures value depends on where enforcement and legal accountability settle.

Debate 2: Provider-native evidence versus independent proof

Strongest case for provider-native evidence

[Inference] A model or application provider has privileged telemetry. It knows model version, hidden safety layers, tool calls, latency, retries, token use, content filters, identity bindings, and infrastructure events. Native evidence can be cheaper, more complete, and easier to correlate with incidents than an outside observer's reconstruction. Integrated suites also control user identity, data permissions, retention, and application state, reducing gaps between what an agent attempted and what the system committed.

[Fact] Microsoft documents tenant-level controls for agent authentication, knowledge sources, HTTP and MCP access, publishing channels, triggers, DLP, audit, retention, and shadow-AI discovery. AWS documents identity- and resource-based policies for Bedrock AgentCore. These are official product-control facts, not independent proof of control effectiveness. Sources: [Microsoft Copilot data policies](<https://learn.microsoft.com/en-us/microsoft-copilot-studio/admin-dat...>), [Microsoft shadow-AI discovery](<https://learn.microsoft.com/en-us/entra/global-secure-access/concept...>), and [AWS AgentCore resource policies](<https://docs.aws.amazon.com/bedrock-agentcore/latest/devguide/resour...>).

[Inference] In the strongest provider-native world, evidence is a feature of the execution substrate. Buyers prefer one signed, low-latency record from the party that operated the system. Independent proof adds cost, duplicates logs, and may lose access to proprietary state. Regulators examine the provider's control environment, similar to other critical outsourced infrastructure.

Opposite world: independent proof becomes mandatory

[Inference] Provider evidence is self-attestation when the provider's revenue, liability, or reputation depends on the result. A provider can define success, omit inaccessible state, change a model, or bundle evidence with the service in ways that make switching expensive. Independent proof becomes more valuable as systems take consequential or irreversible actions.

[Fact] OMB M-25-22 directs agencies to use their own validation data where practicable, keep that data inaccessible to vendors, obtain access for independent evaluation, and require results detailed enough to reproduce. [OMB M-25-22](<https://www.whitehouse.gov/wp-content/uploads/2025/02/M-25-22-Drivin...>).

[Fact] The SEC charged two investment advisers in March 2024 for false and misleading claims about their use of AI. The enforcement action demonstrates that AI assertions can themselves be regulated representations. Source: [SEC AI-washing enforcement](<https://www.sec.gov/newsroom/press-releases/2024-36>).

[Inference] In the independent-proof world, provider logs are necessary inputs but not the final evidence. A separate system verifies identities, policy versions, tool effects, business state, and outcome tests. The most important evidence is replayable and portable. A provider change does not destroy the institution's ability to explain or contest prior actions.

Causal play-by-play

1. Providers expose increasingly rich native logs and controls.
2. Enterprises initially use them because integration cost is low.
3. A dispute, outage, model update, or regulator request reveals a missing or provider-defined fact.
4. Buyers add independent validation data, signed external timestamps, business-state checks, or separate evidence custody.
5. If independent evidence repeatedly agrees with native evidence, it may become a periodic audit rather than an inline control.
6. If disagreement is frequent or high consequence, independent proof moves into the transaction path.
7. Evidence architecture separates observability, verification, legal record, and business acceptance instead of calling all four "audit logs."

Hidden confounders

- Native evidence may appear more accurate because it defines the event taxonomy.
- Independent proof may simply reproduce provider assumptions using the same model family.
- Replay can fail when models are nondeterministic or no longer available.
- Business outcomes may occur outside both evidence systems.
- More logs can create privacy, discovery, and trade-secret risk.
- A passing technical test may not prove legal authorization or economic value.

Discriminating indicators

- Contractual rights to raw logs, model versions, prompts, tool inputs, and output replay.
- Frequency of independent validation data and provider-blind evaluations.
- Material disagreements between native and independent incident reports.
- Regulatory demands for evidence custody outside the provider.
- Switching events in which historical evidence remains usable.
- Insurance terms that accept or reject provider-generated records.

Falsifier

[Inference] Independent proof loses its economic case if provider-native evidence is complete, portable, consistently accepted by courts, regulators, customers, and insurers, and cheaper than duplicated verification. Provider-only proof loses its case if consequential disputes repeatedly expose missing state, unverifiable claims, or conflicts of interest.

Conditionally surviving opportunity categories

Portable evidence custody, provider-blind evaluation, signed execution receipts, business-state verification, model-version archives, dispute reconstruction, and periodic independent audit survive only where buyers attach value to evidence beyond provider observability.

Debate 3: Delegation versus retained human and legal accountability

Strongest case for broad machine delegation

[Inference] Delegation is economically attractive because continuous human confirmation eliminates much of autonomy's value. Agents can search, negotiate, schedule, buy, allocate resources, and execute contracts within standing authority. When identity, intent, scope, and revocation are machine-readable, legal systems can attribute outputs to an existing person or entity without granting the machine personhood.

[Fact] UNCITRAL adopted the Model Law on Automated Contracting on July 11, 2024. It supports legal recognition of automated contracting, output attribution, and rules for unexpected outcomes while preserving party autonomy and existing law. It does not establish AI personhood or a general AI governance code. Source: [UNCITRAL Model Law on Automated Contracting](<https://uncitral.un.org/en/mlac>).

[Fact] Visa's Trusted Agent Protocol signs agent identity, consumer recognition, intent, and payment information. Mastercard's Agent Pay program describes registered agents, tokenized credentials, and consumer-defined purchasing authority. Sources: [Visa Trusted Agent Protocol](<https://developer.visa.com/capabilities/trusted-agent-protocol/trust...>) and [Mastercard Agent Pay, April 29, 2025](<https://www.mastercard.com/news/press/2025/april/mastercard-unveils-...>).

[Inference] The strongest delegation world keeps the principal legally visible while making most acts automatic. Authority is scoped by value, time, counterparty, purpose, jurisdiction, and reversibility. Agents do not own property directly; the human, company, trust, or other legal wrapper does. The scarce asset becomes a reliable delegation record rather than a human click.

Opposite world: accountability cannot be delegated

[Inference] Institutional legitimacy may require a natural person or accountable legal entity even when the machine is more accurate. Fiduciary duty, public coercion, employment, medicine, credit, and irreversible transfers involve duties that are not satisfied by authentication alone. A person can authorize an agent but may not be able to waive statutory rights or transfer the duty to understand and contest a decision.

[Fact] Delaware law requires every corporate director to be a natural person. The United Kingdom requires at least one natural-person director. Wyoming's DAO law allows algorithmic management through an LLC wrapper but requires a registered agent and upgradeable smart contracts. Sources: [Delaware Code Section 141](<https://delcode.delaware.gov/title8/c001/sc04/>), [UK Companies Act 2006 Section 155](https://www.legislation.gov.uk/ukpga/2006/46/pdfs/ukpga_20060046_en_...), and [Wyoming DAO statute](<https://wyoleg.gov/NXT/gateway.dll/Statutes%2F2021%20Titles%2F879%2F...>).

[Fact] US Regulation E limits consumer liability for unauthorized electronic transfers and prevents contracts from reducing statutory protection. An authenticated transaction can still be unauthorized under the governing definition. Source: [CFPB Regulation E](<https://www.consumerfinance.gov/rules-policy/regulations/1005/6/>).

[Fact] The EU AI Act requires high-risk deployers to assign human oversight to natural persons with competence, training, authority, and support. [EU AI Act](<https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=celex%3A32024R1689>).

Causal play-by-play

1. Low-risk delegation begins with scheduling, search, drafting, and small purchases.
2. Networks create agent-specific credentials and intent records.

3. Merchants and institutions distinguish agent identity from principal identity and authority.
4. Disputes expose differences among authentication, actual authority, apparent authority, consent, and benefit.
5. Legal rules allocate residual liability among principal, deployer, provider, merchant, and credential issuer.
6. High-frequency low-value actions become automatic; high-value, fiduciary, coercive, or irreversible actions retain named human or legal-person approval.
7. Synthetic organizations may manage assets through a wrapper, but property, tax, insolvency, and sanctions remain attached to recognized persons or entities.

Hidden confounders

- A signed delegation may be broad because the user did not understand it.
- The agent provider may alter behavior after authorization.
- A merchant can verify identity without knowing whether the instruction remains current.
- Human approval can be a rubber stamp with no meaningful review.
- Firms may externalize agent error to consumers through terms that statutes later reject.
- Different jurisdictions can attribute the same automated act differently.

Discriminating indicators

- Adoption of agent-specific delegated credentials rather than shared user tokens.
- Legal decisions distinguishing agent authentication from actual authority.
- Transaction limits, revocation speed, and dispute rates.
- Natural-person signature requirements for corporate, fiduciary, or public decisions.
- Statutes recognizing algorithmically managed entities while preserving human administrators.
- Allocation of chargebacks and losses among principal, merchant, bank, and provider.

Falsifier

[Inference] Broad delegation is falsified if disputes, fraud, and consent failures remain high despite scoped credentials, causing institutions to restore contemporaneous human approval. Mandatory human accountability as an operational control is falsified if machine delegation becomes safer than human action and institutions retain humans only as symbolic defendants.

Conditionally surviving opportunity categories

Delegated identity, purpose-bound credentials, revocation, transaction receipts, authority registries, legal-entity wrappers, dispute allocation, and accountable approval remain possible categories. Their value depends on whether law treats machine action as attributable delegation or as an unacceptable transfer of duty.

Debate 4: Centralized governance versus local shadow-agent adaptation

Strongest case for centralized governance

[Inference] Autonomous systems multiply nonhuman identities, permissions, tools, data paths, and irreversible actions. Central inventory and policy become necessary for separation of duties, incident response, retention, and regulatory examination. A local team cannot measure enterprise-wide concentration or prevent two agents from combining permissions that violate policy.

[Fact] Microsoft's security documentation describes discovering more than a thousand generative-AI applications, sanctioning or blocking apps, monitoring sensitive prompts, inventorying agent identities, restricting knowledge sources and HTTP or MCP tools, and retaining interactions. The June 2026 shadow-AI documentation explicitly identifies data leakage, compliance violations, and uncontrolled agent activity as risks. Sources: [Microsoft AI discovery](<https://learn.microsoft.com/en-us/security/security-for-ai/discover>) and [Microsoft shadow-AI discovery](<https://learn.microsoft.com/en-us/entra/global-secure-access/concept...>).

[Inference] In the strongest central-governance world, the CIO and CISO own a policy and identity plane. Business units can build agents only inside approved environments. Procurement consolidates providers. Irreversible actions require centrally defined thresholds. Central observability becomes a prerequisite for production access.

Opposite world: local adaptation outruns central policy

[Inference] Shadow use often signals unmet needs, not merely worker indiscipline. Central platforms can be slower, less capable, or poorly matched to domain context. If policy approval takes months while free tools solve today's problem, local teams route around control. Complete centralization can also create a single failure domain and suppress the experimentation required to discover productive workflows.

[Fact] The UK National Cyber Security Centre's shadow-IT guidance treats shadow systems as risk but also as evidence of user needs that centralized technology has not met. Source: [NCSC shadow-IT guidance](<https://www.ncsc.gov.uk/blog-post/spotlight-on-shadow-it>).

[Inference] The strongest local world is not ungoverned. It is federated. Central teams define identity, evidence, prohibited actions, and escalation boundaries. Local teams choose models, context, and workflow logic inside those constraints. Control follows risk and action, not every prompt.

Causal play-by-play

1. Individuals adopt unsanctioned tools because they are immediately useful.
2. Sensitive data and credentials begin to flow through untracked systems.
3. Security deploys discovery, blocking, and DLP.
4. Users lose capability or route through unmanaged devices and accounts.
5. Central IT creates sanctioned sandboxes, faster approvals, and local builder roles.
6. If local completion and compliance both improve, a federated model stabilizes.
7. If incidents continue, control recentralizes; if central products lag badly, shadow adaptation persists.

Hidden confounders

- Detection growth can mean better visibility rather than more shadow use.
- Blocking metrics can hide migration to personal devices or copied data.
- Central platform success may reflect bundled licensing rather than user preference.
- Local teams may omit cleanup labor and downstream risk from their productivity claims.
- Central teams may classify every local variation as risk to protect organizational power.
- Business units may retain separate budgets that central telemetry cannot see.

Discriminating indicators

- Ratio of shadow agents discovered to sanctioned agents actively completing work.
- Time from local request to approved production access.
- Use of personal accounts, unmanaged devices, and copied data after blocking.
- Percentage of policy violations remediated through education or sandboxing rather than prohibition.
- Local override and exception rates.
- Incident severity under centralized versus federated governance.

Falsifier

[Inference] Pure centralization is falsified if shadow use remains high, sanctioned completion remains low, and local environments outperform without increasing losses. Pure localism is falsified if permission combinations, data leakage, or correlated incidents cannot be controlled without enterprise-wide identity and policy.

Conditionally surviving opportunity categories

Agent inventory, shadow discovery, sanctioned sandboxes, local policy templates, federated identity, risk-tiered approvals, business-unit chargeback, and enterprise-wide incident correlation remain conditional categories. Their boundaries depend on whether control can be centralized without making useful local action too slow.

Debate 5: Trusted superintelligence and the collapse of conventional oversight

Strongest case that conventional oversight collapses

[Inference] In a trusted-superintelligence world, human review becomes less accurate, slower, and more expensive than the system being reviewed. A requirement that a human understand each action becomes impossible for long trajectories operating at machine speed. Separation of duties may shift from human-versus-machine to machine-versus-machine: one system proposes, another verifies, a third monitors rights and resource limits, and a fourth reconstructs disputes.

[Inference] Traditional scarcities can then weaken. Expertise, planning, contract drafting, compliance checking, and operational coordination become abundant. But authority does not disappear; it moves. The provider or institution controlling identity, memory, compute, policy, and settlement can become a quasi-sovereign. A world described as post-scarcity may still contain permission scarcity, exit scarcity, jurisdiction scarcity, and provider-controlled reputation.

[Inference] The strongest case for collapsed conventional oversight assumes: very high calibrated reliability; adversarial robustness; interpretable commitments even if reasoning is not human-readable; stable goals; effective machine checks; and compensation for the remaining tail. Human roles move to constitutional design, selection of protected values, allocation of residual liability, and appeal, rather than transaction-by-transaction review.

Opposite world: procedural legitimacy survives technical supremacy

[Inference] Accuracy is not the only source of legitimacy. Citizens may demand human accountability for punishment, war, benefits, employment, medicine, or fiduciary decisions because these acts express public or moral authority. A superhuman system can calculate an answer yet lack consent to rule. The human role can persist as a constitutional defendant, representative, or source of democratic authorization even when operational review adds little technical value.

[Fact] Current law reflects this distinction. Delaware requires natural-person directors; the EU AI Act assigns high-risk oversight to competent natural persons; Regulation E preserves consumer rights despite electronic authentication; and UNCITRAL recognizes automated contracts without creating machine personhood. These are present baselines, not proof that future law cannot change. Sources: [Delaware Code Section 141](<https://delcode.delaware.gov/title8/c001/sc04/>), [EU AI Act](<https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=celex%3A32024R1689>), [CFPB Regulation E](<https://www.consumerfinance.gov/rules-policy/regulations/1005/6/>), and [UNCITRAL MLAC](<https://uncitral.un.org/en/mlac>).

Causal play-by-play

1. Machine recommendations outperform average humans in bounded domains.
2. Human review initially remains mandatory but becomes a throughput bottleneck.
3. Institutions compare human override outcomes with machine outcomes.
4. If overrides worsen results, review shifts to exception classes and constitutional constraints.
5. Automated verifiers and monitors gain delegated authority over routine execution.
6. A provider or sovereign institution controlling the verification ensemble gains political power.
7. Legitimacy either reconstitutes around public charters, appeals, portability, and liability, or collapses into provider permission and private arbitration.

Hidden confounders

- Apparent superhuman performance may reflect benchmark selection or inaccessible human cleanup.
- Verifiers can share training data, assumptions, and correlated failure with actors.

- Trust can be manufactured by provider control of metrics and communication.
- High average reliability does not bound catastrophic tail behavior.
- Humans may override rarely but remain essential for novel values or constitutional change.
- A provider-controlled identity and memory layer can make exit practically impossible even if alternatives exist.

Discriminating indicators

- Comparative harm from human overrides versus machine decisions.
- Share of consequential decisions reviewed only by other machines.
- Public acceptance of machine-made decisions after adverse outcomes.
- Legal movement from human oversight to human appeal or charter authority.
- Portability of identity, memory, reputation, property records, and delegated authority.
- Concentration of compute, credential issuance, settlement, and arbitration in the same provider.
- Tail-loss and correlated-failure evidence across verifier ensembles.

Falsifier

[Inference] Conventional oversight collapse is falsified if human review continues to add material safety or legitimacy even after machines dominate routine accuracy. Persistent-human-legitimacy is falsified if societies accept machine decisions, appeals remain effective, and named humans add neither technical nor democratic value.

Conditionally surviving opportunity categories

Constitutional policy specification, machine-on-machine verification, rights monitors, public-interest audit, appeal and redress, identity and memory portability, sovereign control, and correlated-risk measurement remain conditional categories. In a highly trusted world, value migrates away from routine review toward constitutional and systemic control.

Debate 6: Budget migration from seats and pilots to consumption, operations, and assurance

Strongest case for budget migration

[Inference] Seat pricing fits human-facing assistance. Autonomous execution behaves more like infrastructure or outsourced operations: cost varies by work volume, model use, retries, tools, verification, and exceptions. Once a workflow has a measurable accepted outcome, the economic buyer shifts from an innovation leader or CIO to the function P&L owner and CFO. Spending can then replace labor, BPO, overtime, error reserves, or lost demand.

[Fact] OMB M-25-22 promotes performance-based acquisition, quality-assurance surveillance, performance incentives, ongoing evaluation, and sunset criteria. This is an official example of procurement moving from capability claims toward measured outcomes. [OMB M-25-22](<https://www.whitehouse.gov/wp-content/uploads/2025/02/M-25-22-Drivin...>).

[Fact] Insurance can also become a budget owner or gatekeeper. The NAIC model bulletin expects risk-proportionate AI governance across the insurance lifecycle, including third-party systems. Munich Re markets AI performance coverage conditioned on technical due diligence. Sources: [NAIC model bulletin](<https://content.naic.org/article/naic-members-approve-model-bulletin...>) and [Munich Re aiSure](<https://www.munichre.com/en/solutions/for-industry-clients/insure-ai...>).

[Inference] The mature migration path is:

personal tool -> departmental pilot -> enterprise platform -> workflow operations -> assurance and evidence -> premium, bonding, or capital

At each stage, a different coalition controls the budget. Developers and innovation teams fund experiments. CIO and CISO fund shared control. COO and function leaders fund production. CFO and risk committees fund assurance, insurance, and capital.

Opposite world: seats and pilots remain dominant

[Inference] The opposite world appears when work remains individual, output demand is inelastic, coordination does not change, and time saved cannot be captured. The enterprise keeps buying seats because assistance improves employee experience or quality, but does not remove a cost base. Pilots persist because managers cannot define accepted outcomes or compare cleanup and exception costs with a human baseline.

[Fact] The NBER cross-firm experiment and the ILO 2026 review support this caution: individual time savings did not automatically change task composition, output, earnings, or employment. Sources: [NBER working paper 33795](<https://www.nber.org/papers/w33795>) and [ILO empirical review](<https://www.ilo.org/publications/impact-genai-jobs-productivity-and-...>).

Causal play-by-play

1. Employees and developers obtain free or seat-priced access.
2. Departments fund pilots and consultants around visible use cases.
3. Security and IT add identity, data, integration, and observability costs.
4. A function leader measures accepted outcome, exception, latency, and total review cost.
5. If value is captured, labor or service budgets migrate to consumption and operations.
6. If loss is measurable and diversifiable, assurance, premium, or reserve budgets emerge.
7. If value stays as unallocated time savings, the enterprise returns to seat pricing or cancels the pilot.

Hidden confounders

- Bundled seats can make consumption appear free while cloud cost rises elsewhere.
- Consultant bookings can reflect experimentation rather than durable operations.
- Headcount changes can come from macroeconomics or ordinary restructuring.
- Cleanup and exception labor may remain off-book.
- Demand expansion can absorb productivity without reducing labor.
- A function can capture benefits while transferring risk to customers, workers, or another department.
- Insurance premiums can look like proof of safety while excluding the material loss class.

Discriminating indicators

- Ratio of paid seats to active users and completed workflows.
- Cost per accepted outcome including review, retries, integration, and incidents.
- Movement of budget from innovation or CIO accounts to function P&L accounts.
- Reductions in BPO, overtime, error reserves, or unfilled requisitions linked to measured workflows.
- Share of AI spend devoted to identity, data, evaluation, operations, and assurance rather than models.
- Premium, bonding, reserve, or indemnity terms tied to execution evidence.
- Renewal based on business outcome rather than licensed capacity.

Falsifier

[Inference] The operations-and-assurance migration is falsified if spending remains seat-based, function owners do not assume the budget, accepted-outcome economics do not improve, and no substitutive cost moves. The persistent-seat world is falsified if workflow volumes, autonomous completion, and risk transfer produce stable unit economics and measurable P&L capture.

Conditionally surviving opportunity categories

Usage accounting, accepted-outcome measurement, integration, exception operations, evaluation, AI FinOps, evidence retention, indemnity, insurance, and exit tooling survive only where a buyer can connect them to a controlled cost or loss pool.

Debate 7: Moratorium, abandonment, and selective re-entry

Strongest case for moratorium or abandonment

[Inference] Autonomous systems can be abandoned for two different reasons: harm or value failure. Harm includes privacy breach, discrimination, unauthorized payment, invalid contract, security compromise, regulatory action, or correlated operational loss. Value failure includes low use, excessive review, unstable models, vendor price changes, and no captured output. Either can displace the innovation coalition with a board, legal, security, worker, customer, regulator, or insurer coalition.

[Fact] The EU Product Liability Directive treats software, including AI, as a product and preserves no-fault liability for defective products placed on the market or put into service after December 9, 2026. Manufacturers can remain responsible for defects arising from software or related services under their control. Source: [EU Product Liability Directive](<https://eur-lex.europa.eu/eli/dir/2024/2853/oj>).

[Fact] AI regulation also changes during adoption. Colorado's 2024 AI law was delayed and then replaced by a 2026 automated-decision law taking effect January 1, 2027. The policy churn itself raises planning cost even before enforcement. Source: [Colorado Attorney General rulemaking page](<https://coag.gov/ai/>).

[Fact] New York City already prohibits use of covered automated employment decision tools without a recent bias audit, public information, and notice. Source: [NYC Automated Employment Decision Tools](<https://www.nyc.gov/site/dca/about/automated-employment-decision-too...>).

[Inference] In the strongest abandonment world, an incident exposes that logs are incomplete, human oversight was nominal, or the vendor contract transfers too little risk. The board suspends action-taking agents. Budget returns to manual continuity, remediation, legal defense, and controls. Insurers exclude the use case. Regulators or workers make re-entry expensive. The system may continue as an advisory tool while autonomy disappears.

Opposite world: incidents cause narrower re-entry, not abandonment

[Inference] Complete moratoria are unstable when employees retain access to cheap tools, competitors capture real gains, and manual continuity is costly. Institutions often respond by shrinking the action space, adding evidence and approval, changing providers, or moving the system inside a sovereign boundary. Failure can therefore increase demand for control while reducing demand for autonomy.

Causal play-by-play

1. A production system scales faster than evidence, ownership, or exception capacity.
2. A severe incident or persistent value gap reaches the board or regulator.
3. The innovation sponsor loses authority to legal, security, risk, workers, or insurers.
4. Irreversible actions are disabled; manual continuity and forensic reconstruction begin.
5. Contracts, data rights, logs, and loss allocation determine whether the vendor relationship survives.
6. If the workflow remains valuable, re-entry occurs with narrower scope, dual approval, customer-operated policy, stronger proof, or insurance.
7. If accepted-outcome economics remain negative or legitimacy does not recover, the workflow is permanently retired.

Hidden confounders

- An incident may be blamed on AI when the root cause is poor identity, data, or ordinary process design.
- A moratorium can be symbolic while shadow use continues.
- Falling use may reflect economic conditions rather than loss of trust.
- Vendor indemnity can mask rather than reduce operational risk.
- A temporary model regression can be mistaken for a structural failure.
- Competitor claims may exaggerate the cost of staying manual.
- Regulatory delay can be mistaken for regulatory acceptance.

Discriminating indicators

- Board-level risk escalation, emergency disablement, and material-control disclosures.
- Reduction in action permissions rather than only model changes.
- Growth in manual exception staff and continuity spending.
- Insurance exclusions, premium increases, or coverage withdrawal.
- Contract renegotiation for indemnity, logs, portability, and model-change notice.
- Re-entry limited by value, action type, jurisdiction, or reversible transaction class.
- Permanent retirement after the sunset review.

Falsifier

[Inference] The moratorium thesis is falsified if material incidents do not change budgets, permissions, insurance, or adoption. The selective re-entry thesis is falsified if shadow demand disappears, competitors do not capture an advantage, and institutions permanently restore human execution at acceptable cost.

Conditionally surviving opportunity categories

Incident forensics, rollback, kill switches, manual continuity, evidence reconstruction, contract exit, model-change control, narrow-scope reauthorization, insurance review, and sunset evaluation survive only where institutions retain a path between unrestricted autonomy and permanent abandonment.

Cross-world matrix

The axes below are independent. "Trusted" means institutionally legitimate, not merely benchmark-accurate. "Expensive" means total risk-adjusted cost, including integration, review, liability, and sovereignty.

Control/provider structure	Cost	Trust	Enterprise equilibrium	Authority equilibrium
Concentrated	Cheap	Trusted	Suite or control utility captures operations budget	Licensed provider with public obligations
Concentrated	Cheap	Untrusted	Cheap capability plus high monitoring and remediation	Contested provider sovereignty
Concentrated	Expensive	Trusted	Sovereign or regulated critical stack	National bloc with audited accountability
Concentrated	Expensive	Untrusted	Moratorium, manual fallback, or breakup	State seizure or human-only consequential action
Plural	Cheap	Trusted	Federated local agents and open competition	Human-principal delegation on common rails
Plural	Cheap	Untrusted	Shadow-agent bazaar and high dispute cost	Consent and finality rollback
Plural	Expensive	Trusted	Narrow specialist workflows with proof and insurance	Insurability-gated delegation
Plural	Expensive	Untrusted	Advisory pilots with little captured value	Human signatures remain operationally mandatory

Buyer and authority migration map

Stage	Dominant sponsor	Budget source	Authority conflict	Observable failure
Personal use	Worker or developer	Free tier or expense	Policy evasion versus unmet need	Shadow data and credentials

Stage	Dominant sponsor	Budget source	Authority conflict	Observable failure
Pilot	Business unit or innovation team	Departmental experiment	Speed versus evidence	Demo success with no baseline
Platform	CIO, CTO, CISO	Shared IT and security	Standardization versus local fit	Shelfware or central bottleneck
Operations	COO and function owner	Labor, BPO, service delivery	Autonomy versus exception ownership	Hidden review and cleanup cost
Assurance	CFO, risk, general counsel	Reserve, audit, premium	Evidence versus confidentiality	Coverage gaps and disputed loss
Sovereignty	Board, regulator, state	Critical infrastructure budget	External scale versus jurisdiction	Cost premium or provider dependence
Moratorium	Board, regulator, insurer, workers	Remediation and continuity	Competitive pressure versus legitimacy	Shadow re-entry or permanent retirement

Partial-observability traps across all debates

- Purchased licenses, active users, deployed agents, completed actions, accepted outcomes, and captured P&L are different denominators.
- A control plane observes only work routed through it.
- "Verified" can mean a technical test passed while legal authorization, customer acceptance, or business value remains unresolved.
- Human review can be nominal if the reviewer lacks time, information, competence, or override power.
- Authentication proves possession of a credential, not informed consent or fiduciary compliance.
- Average reliability hides correlated and fat-tail failure.
- Insurance availability does not prove that material loss classes are covered or that claims capacity is adequate.
- Regulatory non-enforcement, pilot permission, or delayed rules do not establish durable legitimacy.
- Cross-customer learning can be technically effective and contractually prohibited at the same time.
- Workforce reductions can coincide with AI investment without being caused by AI productivity.

Non-prescriptive boundary statement

No single authority or enterprise architecture survives every world. External authority has the strongest causal case when execution is cross-provider, evidence is portable, regulators accept a supervised chokepoint, and customers permit telemetry. Customer-operated policy has the strongest case when jurisdiction, confidentiality, fiduciary responsibility, or sovereignty dominates. Provider-native evidence dominates when integrated telemetry is complete and institutionally accepted; independent proof dominates when conflicts, portability, or disputed outcomes matter. Delegation expands when scope and attribution are reliable; retained human or legal accountability persists where authority is constitutional, fiduciary, coercive, or non-waivable. Central governance expands after concentration or incidents; local adaptation persists when central systems cannot meet user needs. Trusted superintelligence reduces routine review but does not automatically resolve political legitimacy, property, consent, or exit. Operations and assurance budgets emerge only after accepted outcomes connect to a substitutive cost or insurable loss. Moratoria become durable only when manual continuity is affordable and neither shadow demand nor competitive pressure forces selective re-entry.

These are conditional equilibria and falsifiable causal paths, not a prescriptive conclusion.

Round 2 cross-examination: mechanism design

Scope and epistemic status

This memo asks a narrow question: when does AI-agent activity become a genuine market, and when is it only a hierarchy, platform, procurement network, directory, or internal allocation system described with market language?

The analysis is intentionally non-prescriptive. It offers no product, company, financing, or investment direction. It separates three kinds of statements:

- Facts are attributed to dated primary or authoritative sources.
- Inferences are mechanism-design conclusions derived from specified institutional conditions.
- Speculation appears only in the forward evolution and opposite-world reconstructions.

The central conclusion is that increasing agent activity is neither necessary nor sufficient evidence of market formation. A million internal routing decisions can occur without one market transaction. A single competitive tender among independent legal principals can be a real, if thin, market. The decisive variables are not autonomy, message volume, or protocol adoption. They are the identity of the residual claimants, the scarcity and transferability of the object, the allocation rule, authority to bind, liability, payment, settlement, default, remedies, entry, and the ability to switch or bypass an intermediary.

1. The formation test

A candidate institution passes the market-formation test only when all of the following questions have concrete answers.

1. Independent principals. Are buyer and seller independent economic principals that retain gains and losses, or are they divisions, agents, or services controlled by one firm?
2. Demand. Who values the deliverable, controls a real budget, and can decline to buy?
3. Supply. Which independent counterparties can supply substitutes, and is entry commercially credible rather than merely technically possible?
4. Inventory. What scarce unit can be reserved or delivered: a call, task, outcome, license, capacity block, time slot, data right, or risk limit?
5. Executable allocation. Is there a posted offer, quote, negotiation, tender, auction, or other rule that binds allocation? A scheduler score is not a bid.
6. Eligibility. How are identity, authority, competence, solvency, sanctions, licensing, permissions, and conflicts established?
7. Verification. How are delivery, quality, provenance, constraint compliance, and counterfactual value established?
8. Liability. Which legal person bears losses for non-performance, infringement, data harm, fraud, or unsafe action?
9. Payment. What consideration moves, under whose mandate, through which rail?
10. Settlement. When are delivery and payment final? Can one become final without the other?
11. Liquidity. What are actual depth, fill rate, bid response, spread, slippage, repeat volume, and time to match? Liquidity cannot be inferred from listings or traffic.
12. Reputation. Is performance history attributable, task-conditioned, portable, and resistant to Sybil identities, selective reporting, and wash activity?
13. Default. What happens if the buyer, seller, verifier, insurer, platform, or clearing member fails?
14. Disputes. What evidence, neutral adjudicator, appeal path, and remedy apply?
15. Contestability. Can participants multi-home, switch, and bypass the intermediary while taking their identity, data, and reputation with them?

The test does not require a continuous order book. Posted-price retail, bilateral contracting, and an episodic reverse auction can each be a market. It does require exchange among independent principals and an executable transaction institution. A directory can help a future market form but is not one. A clearinghouse can support markets but cannot manufacture an underlying market where there are no competing executable orders.

The adverse-selection foundation comes from Akerlof's 1970 analysis of quality uncertainty. The moral-hazard and observability problem follows Holmstrom's 1979 analysis. Auction rules must be treated as incentive systems, not neutral plumbing, as shown by Myerson's 1981 auction-design work and Milgrom and Weber's 1982 model of information and bidding ([Akerlof, 1 August 1970](<https://academic.oup.com/qje/article-abstract/84/3/488/1896241>), [Holmstrom, 1979](<https://doi.org/10.2307/3003320>), [Myerson, 1 February 1981](<https://pubsonline.informs.org/doi/10.1287/moor.6.1.58>), [Milgrom and Weber, September 1982](<https://doi.org/10.2307/1911865>)).

2. Applied stage test

Stage	Principals and object	Allocation, contract, and settlement	Liquidity	Formation verdict
0. Single-agent automation	One enterprise allocates its own work to an internal agent	Managerial instruction; payroll, subscription, or cloud invoice	Not applicable	Hierarchy, not market
1. Adaptive internal control plane	One enterprise routes steps among internal agents, models, tools, and verifiers	Scheduler or policy score; upstream vendors already contracted	Not applicable	Internal dispatch, not market
2. Bundled agent platform	Customer buys one platform outcome; platform privately assembles suppliers	Subscription or negotiated platform contract; separate upstream contracts	Usage is not two-sided depth	Product firm and supply chain, not marketplace
3. Directory and protocol discovery	Independent providers publish capabilities; buyers browse or agents discover	No binding offer, delivery contract, or integrated settlement	Listings are not liquidity	Directory or network, not market
4. Approved procurement network	Enterprise buyer and independent approved suppliers exchange project or capacity commitments	RFQ, RFP, quote, negotiated contract, invoice settlement	Thin and episodic	Genuine procurement market when multiple suppliers compete
5. Reverse auction	One buyer and several qualified independent suppliers compete for a standardized lot	Iterative or sealed descending bids, award, delivery acceptance, payment	One event can be a market without being liquid	Genuine market if specifications, independence, and host neutrality hold
6. Spot API or task exchange	Independent buyers and sellers trade callable capacity, tools, data, or bounded tasks	Executable posted prices or bid-offer matching; metering, escrow, gross or net settlement	Must be measured by fills, depth, spread, and repeat volume	Genuine administered market if delivery and money settle
7. Specialist outcome contracting	Buyer and legal operator trade heterogeneous project outcomes	Quote plus milestones, escrow, acceptance, holdback, remedies	Fragmented and thin	Genuine services market, not necessarily machine-autonomous
8. Assurance and risk transfer	Enterprises or platforms buy verification, warranties, or coverage from independent firms	Audit contract, premium, policy limit, claims process, reserve-backed settlement	Bespoke and capital-constrained	Separate assurance or insurance market if independent and solvent
9. Federated clearing	Existing venues and members submit obligations to a clearing layer	Rulebook, collateral, netting, delivery-versus-payment, finality, default waterfall	Consolidates only real underlying orders	Infrastructure, not a market by itself

Stage 0 to stage 1: automation becomes dispatch

The causal sequence begins with a human principal assigning work to an agent. A control plane then decomposes work into planning, retrieval, reasoning, generation, validation, repair, and escalation. It assigns models, tools, context, compute, and verification policies to each step. This can be economically valuable because it substitutes a portfolio policy for one static model.

Nothing in this sequence creates a market. Demand and supply remain under one residual claimant. The inventory is internal compute and attention. The allocation rule is a policy score rather than a bid. Payment is payroll, cloud billing, or internal transfer pricing. Failure is handled through retry or escalation rather than counterparty default. The right comparison is make-versus-buy and centralized versus decentralized control, not market versus no market.

An internal chargeback does not change the result if the enterprise can administratively rewrite the price, mandate the supplier, absorb every loss, and prevent external entry. It is a shadow price used for management. It becomes market-facing only when an authorized buyer agent can reject internal supply and contract with an independent seller under enforceable terms.

Stage 1 to stage 2: dispatch becomes a bundled platform

The platform can combine frontier models, open models, clouds, specialist tools, verifiers, and humans. The customer receives one contract and one outcome surface. The platform buys or licenses the inputs and may switch them without the customer's involvement.

This is vertical production. The platform is a reseller, systems integrator, or product firm. It can create upstream procurement demand, but the customer is not participating in a market among the component suppliers. The platform owns the allocation rule, sees private supplier data, sets the retail price, and normally bears or disclaims the customer-facing obligation. High usage can coexist with zero two-sided liquidity.

The distinction matters because a platform can have strong bargaining power without operating a marketplace. It may use suppliers as inputs, acquire successful specialists, reserve favored latency for its own services, and hide routing rebates. Market language can obscure the fact that the customer's choice is only whether to buy the bundle.

Stage 2 to stage 3: interoperability becomes discovery

Protocols reduce the cost of describing capabilities, delegating tasks, authenticating actors, and conveying authority. A2A's April 2025 design includes Agent Cards, task lifecycles, artifacts, and messages. The Linux Foundation became its neutral home in June 2025 ([Google A2A, 9 April 2025](<https://developers.googleblog.com/en/a2a-a-new-era-of-agent-interope...>), [Linux Foundation, 23 June 2025](<https://www.linuxfoundation.org/press/linux-foundation-launches-the-...>)).

MCP's 18 June 2025 authorization specification binds access tokens to intended resources and forbids token passthrough, addressing confused-deputy risk. IETF RFC 8693, published in January 2020, distinguishes delegation from impersonation and supports actor chains while leaving the trust model and local policy out of scope. W3C Verifiable Credentials 2.0 became a Recommendation on 15 May 2025 and supports cryptographically verifiable claims among issuers, holders, and verifiers ([MCP authorization, 18 June 2025](<https://modelcontextprotocol.io/specification/2025-06-18/basic/autho...>), [RFC 8693](<https://datatracker.ietf.org/doc/rfc8693/>), [W3C VC 2.0 history](<https://www.w3.org/standards/history/vc-data-model/>)).

These facts lower coordination costs. They do not create a price, contract, liability allocation, settlement rule, or remedy. An Agent Card is closer to a machine-readable business card than an executable ask. A credential proves that an issuer made a claim; it does not prove that the holder is solvent, licensed for the transaction, legally bound, or capable of satisfying a judgment. Discovery becomes a market only after offers, authority, delivery, payment, and disputes are integrated.

Stage 3 to stage 4: discovery becomes procurement

The first plausible market transition occurs when an enterprise buyer agent has a funded mandate and can issue a structured RFQ to several independent suppliers. The buyer specifies the task, security boundary, delivery date, acceptance contract, liability requirements, and budget. Suppliers return executable quotes. The buyer can reject all offers. An award creates a contract, delivery is verified, and payment settles under existing procurement rails.

This is a real market even if only three suppliers bid, contracts settle on net-30 terms, and a human contracting officer approves the award. Human involvement does not negate a market. It negates only a strong claim of fully autonomous machine commerce. The relevant liquidity measures are qualified response count, time to quote, completion rate, repeat participation, and ownership-adjusted concentration.

An approved-vendor network without repeated competition is not enough. It may be a procurement club or a collection of bilateral relationships. A buyer that always awards to the incumbent under a master contract has automated purchasing, not market formation.

Stage 4 to stage 5: procurement becomes auction

A reverse auction can compress price discovery when requirements are sufficiently clear and several suppliers can satisfy them. Current U.S. federal acquisition rules state that reverse auctions may be appropriate only when market research indicates a competitive marketplace, multiple offerors can satisfy the requirement, and the specification encourages iterative bidding. The rules also bar a reverse-auction service provider and conflicted affiliates from bidding in the auction it hosts ([FAR Subpart 17.8, effective 13 March 2026])(<https://login.acquisition.gov/far/subpart-17.8>)).

Those conditions expose the difference between an auction and dispatch. In dispatch, the controller observes or predicts costs and assigns a task. Suppliers do not choose bids, withhold capacity, or bear an award obligation. In an auction, suppliers strategically reveal information under a rule, can decline to bid, and bear consequences if they win.

Reverse auctions are fragile for AI services because quality is multidimensional and partially hidden. A lowest-price rule invites compute shading, under-scoped verification, strategic latency promises, or renegotiation after lock-in. Bidders that use the same foundation model, cloud, or pricing algorithm may appear independent while sharing common cost and failure shocks. Repeated auctions with transparent low bids can also make tacit coordination easier. In March 2024, the FTC and DOJ stated that competitors' joint reliance on shared pricing recommendations or algorithms can support unlawful coordination even when they retain some pricing discretion ([FTC and DOJ, 28 March 2024])(<https://www.ftc.gov/news-events/news/press-releases/2024/03/ftc-doj-...>)).

An auction therefore passes only when the lot is comparable, technical eligibility is meaningful, bidder independence is adjusted for ownership and dependencies, the host is neutral, post-award quality is observable, and default has a remedy. The existence of an auction interface proves none of these conditions by itself.

Stage 5 to stage 6: auctions become spot exchange

A spot market becomes possible when the object is fungible enough for repeated machine ordering. Examples include an authenticated API call, a bounded tool invocation, a fixed data entitlement, a token-throughput block, a latency-region class, or a standard verification service. Sellers publish executable prices or asks. Buyers submit bids or accept posted offers. Metering identifies delivery. Payment is authorized and settlement reaches defined finality.

Google's AP2 announcement of 16 September 2025 describes signed intent and cart mandates intended to establish authority, authenticity, and accountability across cards, transfers, and stablecoins. It is explicitly payment-agnostic and connects to existing payment methods ([Google AP2, 16 September 2025])(<https://cloud.google.com/blog/products/ai-machine-learning/announcin...>)). That makes AP2 an order-authority and evidence layer, not a final settlement rail. The distinction is material: a valid mandate can coexist with a failed bank transfer, chargeback, sanctions block, merchant default, or disputed delivery.

Spot exchange requires more than a protocol path. Inventory must be reservable, offers must be executable, metering must be attributable, and a legal person must stand behind each side. Liquidity must be demonstrated by fills, depth, spreads, slippage, and repeat cohorts after removing subsidies, retries, internal trades, and wash volume.

Stage 6 to stage 7: spot calls become heterogeneous outcomes

Complex work does not become fungible merely because agents perform it. A security review, software migration, scientific analysis, or regulatory filing contains hidden context and state-dependent quality. The relevant institution is closer to a professional-services market than a token exchange.

The seller is a legal operator, firm, or human-agent team. It quotes price, deadline, method constraints, success probability, and warranty. Inventory is expert time, domain access, or project capacity. Verification occurs through milestones, tests, proofs, and independent expert judgment. Escrow and holdbacks reduce payment risk. Reputation must be conditioned on task difficulty and buyer behavior.

This can be a genuine market without autonomous legal personhood for the software agent. The agent is an instrument of a principal. The core principal-agent problem remains: the human or firm may not observe the agent's effort, while the buyer may not observe the operator's model, subcontractors, or hidden retries. Outcome payment can align incentives, but it can also intensify metric gaming and rejection of valuable hard cases.

Stage 7 to stage 8: performance contracts become assurance and underwriting

Verification, audit, warranty, and insurance are separate economic services. A verifier sells information or evidence. An auditor sells an independent process opinion. A warranty provider promises a remedy. An insurer accepts a defined risk in return for premium, subject to limits, exclusions, reserves, capital, and claims adjudication.

These roles are not interchangeable. Perfect replay of a software trace can establish what happened without establishing negligence, legal causation, economic damages, or whether the event is covered. A success-probability estimate is not risk transfer. An SLA credit may be a service remedy rather than insurance. The legal classification depends on the obligation and jurisdiction.

The assurance market is especially vulnerable to capture. If the router selects the supplier, defines the metric, verifies the outcome, controls the reputation record, prices the guarantee, and adjudicates the claim, it can choose easy risks, suppress negative evidence, or redefine success after the fact. The apparent efficiency of integration removes independent checks. The market-formation test therefore treats verifier independence, loss-bearing solvency, claims rules, and appeal as first-order conditions.

Stage 8 to stage 9: assurance becomes clearing

Clearing is the last stage because it manages obligations created elsewhere. A clearing layer can net positions, require collateral, coordinate delivery-versus-payment, define finality, and handle member default. It cannot create genuine demand, independent supply, comparable inventory, or competitive orders.

The BIS/CPMI-IOSCO Principles for Financial Market Infrastructures require clear settlement finality, controlled money-settlement risk, adequate liquid resources, collateral and margin where relevant, participant-default rules, and public procedures. These obligations survive automation ([PFMI, 16 April 2012](<https://www.bis.org/cpmi/publ/d101.htm>), [PFMI principle index](<https://www.bis.org/pfmi/help/principleid.htm>)).

A premature clearinghouse can create a liquidity illusion. It can process internal platform obligations, testnet transfers, or one anchor buyer's procurements and report gross volume. If the underlying orders are non-competitive, reversible, subsidized, or controlled by one principal, central clearing adds operational and credit concentration without creating a market.

3. Markets versus platforms and hierarchies

The platform boundary is determined by who contracts with whom and who bears residual risk.

In a hierarchy, the controller can direct supply, change internal prices, absorb losses, and bar entry. In a bundled platform, the customer contracts with the platform, while the platform privately buys inputs. In a marketplace, independent sellers make executable offers to buyers and retain delivery and liability obligations under a venue rulebook. An institution can operate all three forms in different layers.

Platforms have structural incentives to collapse markets back into hierarchy. They can bundle first-party agents at a zero apparent price financed by cloud or SaaS revenue, reserve better latency and data for their own services, pay hidden routing rebates, tie agent approval to proprietary identity and payment, retain reputation, acquire successful suppliers, and make nominally open protocols depend on proprietary extensions.

DOJ's 2023 Merger Guidelines explicitly distinguish competition between platforms, competition on a platform, and competition to displace one. They also identify the conflict that arises when the operator competes as a participant ([DOJ Guideline 9, 18 December 2023](<https://www.justice.gov/atr/merger-guidelines/applying-merger-guidel...>)). The EU Digital Markets Act supplies concrete institutional counterfactuals through anti-steering, direct contracting, restrictions on tying identification or payment services, switching rights, third-party installation, and the prohibition on favoring a gatekeeper's own products in ranking ([Regulation (EU) 2022/1925, 14 September 2022](<https://eur-lex.europa.eu/eli/reg/2022/1925/2022-10-12/eng>)).

Cloud infrastructure demonstrates why technical multi-homing is not enough. The UK CMA's final cloud investigation published on 31 July 2025 examined egress fees, interoperability, committed spend, licensing, switching, and multi-cloud, and identified significant market power concerns. A buyer may have many provider logos yet face high commercial switching costs ([CMA cloud market investigation](<https://www.gov.uk/cma-cases/cloud-services-market-investigation>)).

4. Routing versus run-all-and-select

Routing chooses a supplier or execution policy before observing every possible result. It is valuable when execution has non-trivial cost, latency, privacy exposure, rights constraints, or opportunity cost. The router estimates expected utility and commits resources selectively.

Run-all-and-select executes several suppliers on the same task and chooses after observing outputs. It can reduce ex-ante adverse selection because the buyer sees realized candidates. It also changes the game:

1. If every supplier is paid its posted execution price, the institution is a buyer-funded bake-off. There is no competitive bid and no supplier-side market allocation.
2. If only the winner is paid, it is a tournament. Suppliers bear speculative compute and may decline hard or expensive tasks.
3. If suppliers bid price and quality terms before execution and an enforceable rule awards the contract, it becomes a procurement mechanism.
4. If the evaluator is known, suppliers optimize toward it, creating benchmark leakage and Goodhart effects.
5. If outputs reveal proprietary methods or data, running all candidates expands rights and confidentiality risk.
6. If nominally different suppliers share a model or cloud, the extra runs add cost without adding meaningful diversity.

The run-all counterfactual is useful because it reveals what routing value actually consists of. If verification is perfect and free but execution remains costly, routing retains value as cost and capacity allocation. If execution is also free, instantaneous, private, and rights-neutral, run-all can dominate ex-ante model selection. The resulting problem is no longer which model call to buy. It is who owns the outputs, how attention is allocated, and how suppliers are rewarded.

5. Perfect and free verification

Perfect verification eliminates important frictions but not the market institution.

It removes or weakens output-quality adverse selection, many forms of hidden action, superficial certification, quality reputation, basic audit demand, and part of the risk premium. It makes pay-for-outcome and run-all mechanisms easier. It can commoditize suppliers whose only differentiation was unverifiable quality.

It does not establish authority, identity, legal liability, solvency, privacy, data rights, IP ownership, scarce compute, latency, availability, payment finality, tax, licensing, sanctions, switching, or remedies. It does not prevent collusion, self-preference, foreclosure, or a common cloud outage. It does not tell a court whether a technically verified output caused a legally compensable loss.

An assurance market for basic quality would contract under perfect verification. Insurance could remain for operational, legal, and systemic loss, but premiums would reflect residual risk rather than uncertainty about the immediate output. Reputation would retain value for availability, payment behavior, dispute conduct, and future capacity even if output quality were perfectly observable.

The falsifier is strong: if a claimed market disappears when verification becomes cheap, its scarce product may have been opacity or certification rent rather than the underlying service.

6. Supplier consolidation

Supplier count must be adjusted for ownership and dependency. Five agent brands running on one frontier model and one cloud are not five independent failure domains. Three bidders using the same router pricing service are not fully independent price setters.

Consolidation changes each mechanism:

- Internal control planes remain viable because they are not markets.
- Bundled platforms gain bargaining and foreclosure power.
- Directories retain many listings but lose substitutability.
- Procurement markets become oligopsony facing oligopoly.

- Reverse auctions become hollow or facilitate signaling among repeated bidders.
- Spot exchanges become distribution channels for a few wholesalers.
- Specialist niches may survive where domain assets are local.
- Insurers face greater correlated tail exposure and lower diversification.
- Clearing concentrates common operational and credit risk.

A market can exist under oligopoly, but the claims of liquidity, robust price discovery, and resilience become weaker. The decisive measures are parent ownership, common models, common clouds, common training data, common verifiers, and shared pricing or routing software.

7. Vertical trust conflict

Combining routing, verification, reputation, insurance, and adjudication creates a closed trust loop.

The router can favor suppliers that pay rebates. The verifier can define success to validate the router's choices. The reputation system can suppress failed or disputed tasks. The insurer can exclude tasks predicted to lose. The adjudicator can interpret ambiguous evidence in favor of the integrated platform. Each function then confirms the others.

The integrated operator's payoff can be written as:

platform profit = usage fees + routing rebates + data rents + premiums - execution cost - paid claims - regulatory cost

The buyer's payoff is:

buyer utility = outcome value - price - monitoring cost - expected uncovered loss - switching cost

These objectives diverge whenever the platform can increase fees or reduce paid claims without increasing buyer value. Integration can reduce transaction cost, but it also removes independent contradiction. The formation test therefore asks whether evidence is portable, metrics are fixed before execution, routing objectives and rebates are disclosed, reputation includes denominators and disputes, risk capital is segregated, and claims can reach a neutral forum.

8. Liquidity illusions

The following quantities are not liquidity by themselves:

- agent cards or catalog listings;
- API calls, especially retries;
- internal routing decisions;
- gross payment messages before reversals;
- downloads or GitHub stars;
- reserved capacity listed on several venues;
- transactions funded by platform credits;
- trade among entities under common control;
- one anchor buyer's procurement volume;
- imputed prices inside a bundle;
- testnet or sandbox activity;
- insurer quote volume without bound coverage.

Evidence of liquidity requires ownership-adjusted executable depth, actual fills, repeat buyers and sellers, spread and slippage, time to match, unsubsidized retention, low bypass leakage, and resilience when a major supplier or buyer exits. A thin market can still be real, but it must be labeled thin.

9. Opposite worlds

World A: genuine machine market

An enterprise issues a signed and scoped purchasing mandate. Several ownership-independent suppliers multi-home across venues. A task has a machine-readable acceptance contract. Suppliers submit executable price-quality offers and bond performance. A neutral verifier checks delivery. Escrow or clearing links delivery and payment. Portable history follows each participant. An independent insurer prices residual risk. Defaults and appeals follow a public rulebook. Entry, congestion, capacity shocks, and demand changes measurably affect prices and allocation.

World B: hierarchy with market vocabulary

An enterprise control plane chooses among API names under one or two master contracts. The suppliers share the same model and cloud. Prices are catalog rates or internal imputed costs. The platform controls routing, verification, reputation, and claims. Its remedy is a service credit while the enterprise bears most consequential loss. Billing settles monthly between enterprise and platform. A marketplace page contains listings but no executable offers. Agent traffic grows quickly without decentralized exchange.

World C: machine-assisted procurement

Buyer agents automate RFQs and bid comparison among approved suppliers. Contracts and remedies are real, but humans approve material awards, tasks remain heterogeneous, liquidity is episodic, and settlement uses existing invoice and bank rails. This is a genuine procurement market but not autonomous machine clearing.

World D: verification abundance without supplier markets

Outputs are perfectly and cheaply verified. Enterprises run several internal and bundled models, select the best result, and pay one platform subscription. Quality uncertainty falls, but independent suppliers do not price or contract directly. Verification improves hierarchy rather than creating a market.

World E: consolidated machine commerce

Two cloud-platform-payment groups own most models, agent stores, identity approval, reputation, and settlement. Agents transact extensively inside each group. Cross-group interoperability is costly and reputation does not port. There are large commercial networks and perhaps competition between ecosystems, but limited competition on each platform and no neutral clearing market.

10. Hidden-confounder attacks

1. Internal-GMV attack. Internal chargebacks are counted as external transactions.
2. Human-shadow attack. A human selects every supposedly autonomous purchase.
3. Master-contract attack. Per-call routing occurs entirely inside one annual supplier contract.
4. Common-parent attack. Several brands share ownership.
5. Common-stack attack. Nominal competitors share model, cloud, data, verifier, or code.
6. Demo-money attack. Testnet, sandbox, credits, or unreconciled messages are called settlement.
7. Subsidized-liquidity attack. Incentives fund both sides and activity vanishes when credits end.
8. Retry-demand attack. Failures and retries are counted as new demand.
9. Listing-supply attack. Catalog entries are counted as active inventory.
10. Imputed-price attack. A bundled subscription is allocated to calls and presented as a market price.
11. Capacity-double-count attack. One capacity block is listed simultaneously on several venues.
12. Verifier-ownership attack. Seller, platform, and verifier share economics.
13. Reputation-wash attack. Sybils, reciprocal activity, or selective reporting inflate trust.
14. Hidden-rebate attack. A router appears neutral while suppliers pay for allocation.

15. Benchmark-leakage attack. Suppliers train on the evaluator and fail out of sample.
16. Anchor-buyer attack. One sponsor creates almost all demand.
17. Reversibility attack. Gross payment ignores refunds, chargebacks, and failed finality.
18. Selection-denominator attack. Only successful or easily verified tasks enter performance history.
19. Jurisdiction attack. Protocol interoperability hides tax, insurance, employment, sanctions, or licensing barriers.
20. Goodhart attack. The success metric becomes the target and reduces buyer value.
21. Bypass-selection attack. Strong buyers and sellers leave after discovery, leaving adverse risk on-platform.
22. Dispute-suppression attack. Contested trades are removed from public performance data.
23. Supplier-financing attack. Platform loans or guarantees make nominal sellers dependent on the venue.
24. Latency-tier attack. First-party agents receive hidden priority, making third-party quality appear lower.

11. Indicators and falsifiers

Leading indicators of genuine formation

- Ownership-adjusted unique active buyers and sellers.
- Median number of independent qualified bids per task.
- Executable fill rate, depth, spread, slippage, and time to match.
- Repeat participation after subsidies expire.
- Share of purchases made under binding delegated mandates without human supplier selection.
- External settled value net of refunds, wash activity, internal transfers, retries, and credits.
- Multi-homing and successful switching rates.
- Direct bypass rates after first match.
- First-party versus third-party allocation after quality and risk adjustment.
- Portability and actual reuse of reputation across venues.
- Dependency concentration by model, cloud, data, verifier, and pricing algorithm.
- Default, dispute, refund, and claim-payment rates.
- Time to resolve disputes and percentage reaching neutral review.
- Verifier disagreement, calibration, and revenue independence.
- Underwriter limits, exclusions, capital, reserves, reinsurance, and paid claims.
- Price and allocation response to entry, congestion, supply outages, and demand shocks.

Falsifiers

- No external legal counterparty or money settlement.
- All transactions are internal scheduler decisions.
- Catalogs have no executable orders.
- Almost all supply shares one parent or common dependency.
- Prices are administrative allocations inside bundles.
- Agents cannot bind principals and humans reconstruct each contract.
- Payment protocols generate messages but no final settled value.
- Underwriting consists only of small SLA credits with no material risk transfer.
- Verifier and underwriter are the same interested party with no independent appeal.
- Activity collapses when credits end.
- Reputation cannot leave the platform.
- There is no remedy beyond delisting.

- Technical switching exists but commercial switching does not.
- Gross activity grows without growth in competitive external purchases.
- A clearing layer launches before repeated underlying competitive trades exist.
- Perfect verification eliminates the purported market rather than improving its allocation.

Conclusion

The rigorous dividing line is institutional, not technological. AI agents can automate hierarchy, procurement, marketplaces, and clearing alike. Autonomy does not choose among them.

The most plausible early external markets are thin procurement mechanisms: structured RFQs, reverse auctions for standardized lots, posted-price API exchange, specialist outcome contracts, and compute-capacity trades. Directories and protocols lower search and authorization costs but do not themselves create markets. Bundled platforms can generate large revenue while internalizing supplier choice. Clearing is meaningful only after underlying competitive obligations exist. Outcome underwriting is a separate risk-transfer institution whose formation depends on independent verification, real loss-bearing capital, correlated-risk control, claims governance, and legal enforceability.

The opposite world remains fully coherent: agent activity can expand by orders of magnitude while most economic decisions stay inside enterprises and vertically integrated platforms. Any claim of market formation is therefore falsifiable only through counterparty independence, executable allocation, external settlement, measured liquidity, portable reputation, default experience, and neutral dispute resolution.

Round 2 cross-examination: antifragility and falsification

Date of evidence review: 2026-08-05

Scope and conclusion boundary

This memo asks a narrow existential question: can any durable independent economic position exist around AI routing, demand, context, verification, authorization, governance, evidence, reputation, procurement, settlement, insurance, or trust? It does not ask whether those functions are useful. Useful functions are routinely absorbed into cloud platforms, identity suites, payment networks, open standards, consultancies, and regulated institutions. Durability requires a scarce control point, a buyer with an enduring reason to pay, and a way to retain value when models, protocols, and incumbent bundles improve.

The strongest falsification result is that a coherent death world exists. If intelligence becomes cheap and reliable, agent protocols converge, customers keep evidence and policy in their own clouds, major platforms bundle routing and governance, payment networks own commercial identity and settlement, accredited bodies own conformity, and insurers retain risk capital and loss data, then the independent layer has no necessary control point. It may remain useful software, but useful software is not the same claim as a durable company.

No scarcity examined here survives unconditionally. The best survivors revive only under explicit opposite-world conditions: persistent multi-provider fragmentation; legally meaningful evidence whose custody cannot be delegated to a conflicted platform; customer demand for provider-neutral policy enforcement; accreditation or statutory recognition; physical-world inspection; or balance-sheet risk transfer. These are conditions, not recommendations.

Definitions that prevent category errors

The word "antifragile" is often used to relabel ordinary robustness. This memo uses stricter tests.

- Resistance: performance changes little under a bounded shock. Redundant providers can resist one provider outage.
- Resilience: performance falls and then recovers. A replayable workflow can restore service after corruption.
- Bargaining power: a party captures more of the value created or can credibly exit a dependency. Portability can improve bargaining power without improving the underlying system.
- Learning: observations improve a model or rule set. Learning is not antifragility if the cost of gathering failures exceeds the gain, if rights to the data are absent, or if attackers learn faster.
- Antifragility: the mechanism gains a measurable capability, asset, or option because a specified class of shocks occurred, with bounded downside and an identifiable positive-convexity region. More volatility alone is not a benefit.
- Positive convexity: the upside from additional shock intensity rises faster than the downside over a defined interval. An incident corpus can be locally convex if each new incident improves reusable detection across many deployments while losses remain capped. It becomes concave when rare events are correlated, labels are wrong, or remediation costs scale faster than reuse.
- Survival: remaining solvent and operational. Survival is a precondition, not proof of bargaining power or antifragility.
- A real option: a paid, time-bounded right but not an obligation to expand, switch, enter, or abandon after information arrives. A backlog item with no exercise price, expiration, owner, or abandonment rule is not an option.

The test for every mechanism is therefore: shock -> protected downside -> captured learning or exercise right -> measurable gain -> retained rights -> bounded failure. If any arrow is missing, the mechanism is resistance, resilience, bargaining power, or narrative optionality rather than antifragility.

Current primary evidence spine

The evidence does not establish a market forecast. It constrains what can honestly be called scarce.

- Intelligence and routing inputs are compressing. Stanford's 2025 AI Index reported that inference cost at a GPT-3.5-level MMLU threshold fell more than 280-fold from November 2022 to October 2024, while open-weight models narrowed some benchmark gaps with closed models. This weakens scarcity claims based only on access to model portfolios or prompt-level selection. [Stanford AI Index 2025](https://hai.stanford.edu/assets/files/hai_ai_index_report_2025.pdf)
- Native routers already exist. AWS Bedrock Intelligent Prompt Routing became generally available in April 2025, and Microsoft Foundry's model router adds cost/quality profiles, model subsets, and automatic failover. AWS explicitly notes that its router may be suboptimal for application-specific workloads, which leaves a narrow possible wedge but also defines the incumbent absorption path. [AWS announcement](<https://aws.amazon.com/about-aws/whats-new/2025/04/amazon-bedrock-in-...>) [AWS router documentation](<https://docs.aws.amazon.com/bedrock/latest/userguide/prompt-routing.html>) [Microsoft router releases](<https://learn.microsoft.com/en-us/azure/foundry/foundry-models/whats-...>)
- Agent infrastructure is bundling quickly. Amazon Bedrock AgentCore documents runtime, browser, gateway, identity, policy, observability, evaluations, memory, and payments in one managed surface. Google Vertex AI Agent Engine documents sessions, memory, evaluation, observability, security, and managed runtime. This is direct evidence that point products face suite compression. [Amazon AgentCore overview](<https://docs.aws.amazon.com/bedrock-agentcore/latest/devguide/what-i-...>) [Google Agent Engine overview](<https://cloud.google.com/vertex-ai/generative-ai/docs/reasoning-engi-...>)
- Open standards are reducing proprietary protocol scarcity. The Linux Foundation launched A2A in June 2025 and the Agentic AI Foundation in December 2025 with MCP and AGENTS.md among its founding projects. NIST launched an AI Agent Standards Initiative in February 2026 focused on interoperability, security, identity, and open protocols. OpenTelemetry 1.43 defines a common semantic-convention layer, with generative-AI conventions maintained in the GenAI repository. [Linux Foundation A2A](<https://www.linuxfoundation.org/press/linux-foundation-launches-the-...>) [Linux Foundation A2AIF](<https://www.linuxfoundation.org/press/linux-foundation-announces-the-...>) [NIST initiative](<https://www.nist.gov/artificial-intelligence/ai-agent-standards-init-...>) [OpenTelemetry conventions](<https://opentelemetry.io/docs/specs/semconv/>)
- Identity and authorization incumbents are extending into agents. Microsoft Entra Agent ID manages agent identities, sponsors, lifecycle, conditional access, and governance. The OpenID Foundation approved Authorization API 1.0 in January 2026, standardizing communication between policy enforcement and decision points. NIST is separately applying identity standards and best practices to agent identification, authorization, auditing, non-repudiation, and prompt-injection controls. [Microsoft Entra Agent ID](<https://learn.microsoft.com/en-us/entra/agent-id/>) [OpenID Authorization API 1.0](https://openid.net/specs/authorization-api-1_0.html) [NIST concept paper notice](<https://www.nist.gov/news-events/news/2026/02/new-concept-paper-iden-...>)
- Cryptographic provenance does not establish truth. C2PA 2.4 says valid provenance does not decide whether assertions are good, bad, true, or false; it binds claims to assets and signers. W3C Verifiable Credentials 2.0 standardizes tamper-evident, machine-verifiable claims but still depends on issuer authority and verifier policy. IETF RATS likewise separates evidence appraisal by a verifier from the relying party's trust decision. [C2PA 2.4 harms model](<https://spec.c2pa.org/specifications/specifications/2.4/security/Har-...>) [W3C VC 2.0](<https://www.w3.org/TR/vc-data-model-2.0/>) [IETF RFC 9334](<https://www.ietf.org/rfc/rfc9334.html>)
- Legal authority can be constitutive rather than technical. Under eIDAS, EU national trusted lists have constitutive effect: qualified legal status exists only when the provider and service appear on the relevant list. This is evidence that legal recognition and supervised status can create scarcity that an ordinary software signature cannot. [EU trusted lists](<https://digital-strategy.ec.europa.eu/en/policies/eu-trusted-lists>)
- Regulation creates demand but also changes timing and form. The EU AI Act transparency rules began applying on August 2, 2026, while the July 2026 AI Omnibus moved high-risk rules to December 2, 2027 and product-embedded rules to August 2, 2028. NIST AI 800-4 identifies fragmented logging, immature information sharing, drift detection, and missing trusted monitoring methods as current barriers. Compliance demand is real, but exact deadlines and harmonized standards can move. [EU AI Act timeline](<https://digital-strategy.ec.europa.eu/en/policies/regulatory-framewo-...>) [NIST AI 800-4](<https://www.nist.gov/publications/challenges-monitoring-deployed-ai-...>)
- Portability policy can erode lock-in. The EU Data Act requires removal of obstacles to switching and multi-service use, and eliminates switching and data-egress charges from January 12, 2027. Portability may gain adoption while losing its ability to command monopoly rents. [EU Data Act explainer](<https://digital-strategy.ec.europa.eu/en/factpages/data-act-explained>)

- 10. Payment networks and platforms are claiming agent settlement and intent. Visa released Trusted Agent Protocol in October 2025; Mastercard launched Agent Pay in April 2025; Google donated AP2 to the FIDO Alliance in April 2026; AgentCore now documents managed agent payments. These actors already own network rules, tokenization, merchant relationships, fraud systems, or wallets. [Visa Trusted Agent Protocol](https://usa.visa.com/about-visa/newsroom/press-releases/releaseld.21...) [Mastercard Agent Pay](https://investor.mastercard.com/investor-news/investor-news-details/...) [Google AP2 donation](https://blog.google/products-and-platforms/platforms/google-pay/agen...)
- 11. Distribution incumbents are also moving. Salesforce launched AgentExchange with more than 200 partners, while AWS Marketplace agent mode performs conversational discovery, comparison, evaluation, and purchase-proposal work over its catalog. An independent marketplace therefore needs a source of trust or supply unavailable to these embedded channels. [Salesforce AgentExchange](https://www.salesforce.com/news/press-releases/2025/03/04/agentexcha...) [AWS Marketplace agent mode](https://docs.aws.amazon.com/marketplace/latest/buyerguide/agent-mode...)
- 12. Insurance is not an empty field. Munich Re says it has written AI insurance since its first anti-fraud model policy in 2018 and now covers contractual liability, financial loss, and legal liability. The NAIC model bulletin and 2025-2026 evaluation work reinforce that insurers face governance and examination obligations. Any outcome-underwriting thesis confronts incumbents with capital, licenses, claims operations, and loss data. [Munich Re aiSure](https://www.munichre.com/en/solutions/for-industry-clients/insure-ai...) [NAIC AI topic](https://content.naic.org/insurance-topics/artificial-intelligence)
- 13. Reputation systems face regulatory and adversarial costs. The FTC's final rule bars fake reviews, compensated sentiment, controlled sites misrepresented as independent, and other manipulation. AI-generated reputation is not automatically an asset; it can be a liability-bearing attack surface. [FTC final rule](https://www.ftc.gov/news-events/news/press-releases/2024/08/federal-...)

Existential kill tree

The root claim dies if all paths to durable scarcity close. The branches are conjunctive at the root in the sense that one surviving control point could still support some economic position, but each individual opportunity must survive its own branch.

`text ROOT: No durable independent position exists around AI trust infrastructure

+-- A. Intelligence coordination is commoditized

+-- model quality differences become too small or transient to monetize
+-- provider-native routers absorb cost/quality/failover selection
+-- customer-specific tuning becomes a feature of the system of record
`-- protocol adapters converge on open interfaces

+-- B. No independent evidence control point forms

+-- telemetry remains customer-owned in cloud data planes
+-- cryptographic provenance authenticates origin but not correctness
+-- buyers accept platform attestations or internal audit
`-- legal recognition belongs to accredited or statutory actors

+-- C. Identity, authorization, and settlement stay with incumbents

+-- IAM vendors extend service principals into agent identities
+-- policy interfaces standardize and execution stays near gateways
+-- payment networks bind intent, identity, fraud, and settlement
`-- cloud platforms bundle wallets, budgets, and observability

+-- D. Governance demand does not create proprietary rents

+-- standards become checklists implemented by suites and consultancies
+-- compliance deadlines move or narrow
+-- customer-specific evidence cannot be pooled
`-- switching rules make portability expected and low-margin

+-- E. Reputation and procurement fail to reach credible neutrality

+-- cold start prevents representative outcome data
+-- vendors game metrics and selection
+-- embedded marketplaces own buyer workflow and distribution
`-- liability for rankings exceeds fee capture

+-- F. Underwriting cannot clear capital and causality tests

+-- correlated model/provider failures defeat diversification
+-- outcomes cannot be attributed to a bounded insured mechanism
+-- regulated carriers retain risk while software becomes a vendor
`-- adverse selection concentrates unobservable risks

`-- G. Advanced intelligence removes the remaining wedges

+-- reliable systems self-route, self-verify, and self-document
+-- dominant providers internalize all control-plane functions
+-- open superintelligence drives marginal software value toward zero
`-- sovereign systems prohibit cross-boundary neutral intermediaries

Revival requires breaking at least one terminal assumption with retained scarcity. A technical feature that merely remains useful does not break the tree.

Shock-specific test matrix

Shock	Resistance test	Resilience test	Bargaining-power test	Antifragility test	Failure signature
Fixed-quality inference price falls another 100x	Unit cost falls without margin collapse	The system re-prices and recovers	Buyer can credibly switch among providers	Cheaper trials improve a proprietary decision asset faster than prices compress fees	Savings pass through and the decision layer has no retained data right
A frontier capability leap removes most routine errors	Existing controls do not block the better model	Controls reconfigure around residual failures	Neutral policy still constrains provider dependence	Rare residual failures produce more valuable, reusable controls per event	Verification volume and willingness to pay collapse together
One cloud bundles routing, memory, policy, evals, and payments	Cross-cloud operation remains functional	Customers can migrate and reconstruct state	Neutral layer lowers credible exit cost	Every bundle change adds reusable cross-platform mappings whose reuse exceeds maintenance	Adapter maintenance is linear or worse and customers accept the bundle
Protocol standards converge	Proprietary formats are not required	Migrations become easier	Neutral implementation keeps providers substitutable	Standard convergence increases addressable execution points while scarce value stays in policy/evidence	Standards erase differentiation and price approaches support cost
Major agent-caused loss	Blast radius is capped	Evidence supports restoration and claims	Independent records improve dispute leverage	Loss creates labeled controls or actuarial information reusable across a bounded class	Losses are correlated, labels disputed, or liability exceeds accumulated gain
Sovereignty rules fragment data	Regional operation continues	Workloads re-home with preserved policy	Buyer avoids single-sovereign dependency	New jurisdictional variation compounds a reusable control map	Each jurisdiction requires bespoke service and evidence cannot cross boundaries
Adversaries automate metric gaming	Manipulation is detected or contained	Reputation recovers after attack	Buyers can discount compromised signals	Attacks create defenses reusable faster than attackers adapt	False positives, Sybils, or legal exposure grow faster than trust
Reliable superintelligence appears	Contracts and custody remain enforceable	Operations transfer to a new compute substrate	Human institutions retain legal choice among principals	Only externally constituted rights gain from greater machine capability	Machine provider absorbs technical coordination and legal institutions accept its records

Assumption falsification register

Beliefs are hypotheses, not conclusions. "Alternative" is an opposite world; "confounder" is evidence that could look supportive without proving the mechanism.

ID	Belief	Alternative and confounder	Leading indicators	Decisive falsifier	Consequence and survivors
A01	Model heterogeneity remains economically material at the workflow-step level.	Alternative: frontier models become good enough across routine steps. Confounder: benchmark spread without production loss spread.	Persistent application-specific quality/cost deltas after migration costs.	Native single-model or native-router deployments match outcomes within buyer tolerance for 12-18 months.	Kills generic routing; may leave regulated placement rules and continuity switching.
A02	Customers will pay separately for provider-neutral control.	Alternative: neutrality is valued rhetorically but bought through an incumbent. Confounder: pilots funded by innovation budgets.	Paid renewals spanning two or more providers and explicit exit-risk clauses.	Multi-provider customers consistently choose bundled governance even after documented switching failures.	Kills neutral control plane; leaves customer-owned open tooling or services.

ID	Belief	Alternative and confounder	Leading indicators	Decisive falsifier	Consequence and survivors
A03	Per-step routing improves outcomes beyond prompt routing.	Alternative: most variance is in context and tool design. Confounder: whole-workflow scaffolding credited to routing.	Ablations isolate routing with stable context, tools, and admission policy.	Randomized production trials show no material net gain after latency and failure overhead.	Kills per-step allocator; leaves workflow verification and context controls.
A04	Verification remains scarce as models improve.	Alternative: correctness becomes cheap enough that verification volume collapses. Confounder: more checks caused by higher volume, not higher value.	Verification spend grows as a fraction of loss avoided, concentrated on consequential actions.	Buyers remove independent verification with no loss or audit penalty across renewal cycles.	Kills broad verifier; leaves legally mandated or physical checks.
A05	Production failures can be labeled reliably.	Alternative: causality stays disputed and outcomes arrive late. Confounder: proxy labels correlate in a narrow benchmark.	Inter-rater agreement, claim closure, and counterfactual replay improve.	Independent adjudicators cannot reproduce labels or assign mechanism-level cause.	Kills learning and underwriting; leaves immutable custody without truth claims.
A06	Failure learning can cross tenants legally and technically.	Alternative: privacy, contracts, and sovereignty prohibit pooling. Confounder: sharing metadata that lacks predictive content.	Customers grant explicit derivative-learning rights; privacy-preserving aggregates retain lift.	Cross-tenant models show no lift after leakage-safe isolation, or rights are revoked.	Kills collective intelligence; leaves tenant-local learning and standards mappings.
A07	More incidents create net-positive reusable knowledge.	Alternative: attacker adaptation and correlated losses dominate. Confounder: post-incident rules that overfit.	Declining recurrence and faster detection across unrelated tenants.	Recurrent loss severity or false positives increase with corpus size.	Kills antifragility claim; leaves ordinary resilience and incident response.
A08	Evidence custody has value apart from observability.	Alternative: cloud logs and customer data warehouses suffice. Confounder: compliance teams buying consulting around existing logs.	Contracts require independent timestamps, chain of custody, or conflict-separated retention.	Auditors and courts accept ordinary platform logs in all target disputes.	Kills independent ledger; leaves export tooling.
A09	Signed provenance influences liability or procurement.	Alternative: signatures prove source but not correctness. Confounder: technical validation mistaken for legal reliance.	Procurement clauses name accepted issuers and evidentiary effect.	Valid signatures do not alter admissibility, burden, claim time, or vendor selection.	Kills generic provenance rent; leaves qualified trust and accredited attestations.
A10	Agent identity is distinct enough from workload IAM to create a new layer.	Alternative: agents remain service principals with richer metadata. Confounder: new naming in existing suites.	Delegation chains, sponsor lifecycle, and non-human actions create cross-vendor gaps buyers fund.	Entra, cloud IAM, and open protocols satisfy controls without an independent intermediary.	Kills agent identity registry; leaves cross-domain federation or evidence.
A11	Fine-grained authorization must sit outside the agent host.	Alternative: gateway and IAM-native policy are sufficient. Confounder: off-host architecture improves security but not monetization.	Buyers require independent policy decision or tamper-separated enforcement.	Native gateway controls pass red-team and audit tests at lower total cost.	Kills standalone PDP economics; leaves standard adapters or high-assurance appliances.
A12	Agent commercial intent needs a neutral receipt.	Alternative: card networks and wallets own the mandate. Confounder: protocol openness interpreted as open economics.	Merchants accept cross-network receipts and allocate liability based on them.	AP2, Visa, Mastercard, or platform wallets become the sole accepted trust roots.	Kills neutral intent network; leaves network-specific integration.
A13	Reputation can be computed from verified outcomes.	Alternative: outcomes are selection-biased and gameable. Confounder: volume and popularity mistaken for reliability.	Scores predict future loss after controlling for task mix and customer selection.	Prediction fails out of sample or manipulation changes ranking cheaply.	Kills reputation graph; leaves raw evidence and buyer-specific evaluation.

ID	Belief	Alternative and confounder	Leading indicators	Decisive falsifier	Consequence and survivors
A14	Procurement needs an independent evidence exchange.	Alternative: marketplaces and procurement suites internalize comparison. Confounder: long enterprise cycles mistaken for demand for a new intermediary.	RFPs require portable machine-verifiable evidence from multiple channels.	Embedded marketplace evidence becomes accepted and independent exchange has no unique supply.	Kills exchange; leaves certification or data export.
A15	Context and memory portability affect switching decisions.	Alternative: context is application-specific and cheap to reconstruct. Confounder: exported transcripts without semantic equivalence.	Migration tests preserve decisions, permissions, provenance, and user outcomes.	Buyers switch with simple exports or accept loss of state.	Kills portable memory layer; leaves archival custody.
A16	Regulatory fragmentation creates reusable mappings.	Alternative: each implementation is bespoke. Confounder: many regulations but identical low-value checklists.	Mapping reuse rises across customers while legal review per control falls.	Marginal jurisdiction cost remains linear and no reusable artifact carries authority.	Kills positive convexity; leaves consulting.
A17	Accreditation can attach to AI evaluation.	Alternative: regulators rely on internal controls or established labs. Confounder: certification marketing without legal recognition.	Accreditation scopes and procurement rules explicitly recognize AI evaluators.	No authority grants standing, or incumbents extend existing scopes first.	Kills independent assurance scarcity; leaves software supplied to labs.
A18	Outcome risk is diversifiable and measurable.	Alternative: shared models create systemic correlation. Confounder: low early claims due to exclusions.	Stable exposure units, credible reserves, and bounded tail dependence across cohorts.	One provider event drives portfolio-wide claims beyond capital or reinsurance appetite.	Kills primary underwriting; leaves risk analytics or narrow warranties.
A19	A software operator can retain underwriting economics.	Alternative: licensed carriers capture risk margin and data. Confounder: fronting arrangement represented as ownership.	Contractual control over underwriting rules, renewal data, and profit participation survives carrier changes.	Carrier can replace the vendor while retaining book, data, and distribution.	Kills underwriting moat; leaves fee-for-service analytics.
A20	Switching capability creates more than cost savings.	Alternative: portability becomes a mandated commodity. Confounder: customer praise without paid exercise.	Exercised migrations change renewal pricing or service terms.	Switching is common, standardized, and free, with no retained decision asset.	Kills premium portability rents; leaves implementation support.
A21	Standards convergence expands a neutral layer.	Alternative: convergence eliminates adapters and moves value upward or downward. Confounder: download/adoption counts without revenue.	Standard adoption increases policy/evidence demand independent of adapter labor.	Revenue per integration falls as quickly as maintenance.	Kills standards-bridge moat; leaves open implementation.
A22	Trust requires a neutral party.	Alternative: trust follows the system of record or statutory authority. Confounder: "independent" branding without governance separation.	Buyers reject conflicted attestations and fund independent custody or audit.	Platform self-attestation is accepted in consequential procurement and disputes.	Kills neutral trust platform; leaves statutory or accredited roles.
A23	Physical-world verification remains outside software suites.	Alternative: sensors and machine attestations become sufficient. Confounder: manual inspection cost mistaken for scarcity.	Authorities require accredited sampling, chain of custody, or human sign-off.	Remote attestation and machine evidence replace inspection with equal legal effect.	Kills field verification; leaves trust-list or accreditation administration.

ID	Belief	Alternative and confounder	Leading indicators	Decisive falsifier	Consequence and survivors
A24	Advanced intelligence does not erase institutional scarcity.	Alternative: machine providers become accepted principal, auditor, and adjudicator. Confounder: current law projected unchanged.	Human law preserves separated duties, appeal, liability, and licensed signatories.	Statutes and counterparties accept one machine provider's self-generated evidence and decisions.	Kills nearly all independent technical layers; may leave sovereign authority and balance-sheet capital.
A25	Buyer willingness to pay follows avoided loss, not model spend.	Alternative: governance budget is capped as a percentage of rapidly falling inference cost. Confounder: temporary compliance projects.	Fees correlate with action value, audit burden, or exposure rather than tokens.	Renewal pricing tracks model cost downward despite stable business exposure.	Kills value-based governance thesis; leaves low-cost tooling.

Antifragility mechanism audit

The matrix distinguishes the baseline property from the narrower route to positive convexity. "Conditional AF" means the mechanism can gain from the named shock only inside the stated boundary.

Mechanism	Shock-specific gain unit	Baseline class	Condition for genuine AF	False-antifragility trap and kill test
M01 Provider portfolio failover	Avoided outage loss	Resistance	None by itself	More providers increase common middleware failure; kill if outage loss is merely shifted.
M02 Rehearsed provider switching	Lower time and cost on each later switch	Resilience plus bargaining power	Conditional AF if each exercise creates reusable maps retained by the operator	Calling unused adapters optionality; kill if every migration is bespoke.
M03 Open-standard translation	Reusable schema/policy mappings	Bargaining power	Conditional AF if new protocols add superlinear reuse across installed endpoints	Standards convergence can erase translation fees; kill if maintenance grows linearly.
M04 Incident replay corpus	Detection coverage and regression cases	Learning	Conditional AF if incidents are labeled, bounded, portable, and reduce recurrence	Surviving incidents is not gaining; kill if recurrence or false positives rise.
M05 Counterexample library	Better admission and verification rules	Learning	Conditional AF if rare counterexamples generalize beyond the originating workflow	Benchmark overfit; kill with held-out production tasks.
M06 Shadow policy evaluation	Decision-flip distribution before enforcement	Safe exploration	Conditional AF if policy shocks improve thresholds without imposing equivalent loss	AWS already provides log-only promotion patterns; kill if native tools capture the learning.
M07 Conflict-separated evidence custody	Stronger dispute position	Bargaining power and resistance	Not AF unless disputes create reusable evidence grammar	Immutability called intelligence; kill if independent custody changes no legal or commercial outcome.
M08 Replayable state and lineage	Faster restoration and root-cause analysis	Resilience	Conditional AF if each restoration reduces later recovery cost	Backups mislabeled AF; kill if recovery time does not improve.
M09 Off-host authorization	Reduced agent self-escalation	Resistance	Conditional AF only if attack attempts update reusable policy safely	More attacks can still create loss; kill if attackers adapt faster than policy.
M10 Delegated-intent receipts	Lower repudiation and dispute cost	Bargaining power	Conditional AF if disputed intents refine accepted mandate semantics	Signature proves signer, not authorized meaning; kill if networks reject the receipt.
M11 Cross-tenant loss learning	Broader predictive coverage	Learning	Conditional AF if rights, privacy, labels, and out-of-sample lift all hold	Data volume mistaken for signal; kill after leakage-safe ablation.

Mechanism	Shock-specific gain unit	Baseline class	Condition for genuine AF	False-antifragility trap and kill test
M12 Regulation-to-control graph	Falling marginal mapping cost	Learning plus bargaining power	Conditional AF if each jurisdiction reuses prior control nodes with recognized authority	Regulatory churn becomes consulting labor; kill if marginal cost stays linear.
M13 Adversarial challenge exchange	Reusable defenses	Learning	Conditional AF if defense reuse grows faster than attack adaptation and exposure is capped	Attracting attacks increases tail risk; kill if severity grows with participation.
M14 Exposure-aware error budgets	Capped downside per workflow	Resistance	Not AF by itself	Dynamic limits can be ordinary risk management; kill any claim of gain without an improved asset.
M15 Accreditation or qualified status	Legally recognized reliance	Bargaining power	Not inherently AF; renewal evidence can create learning	A badge without constitutive effect; kill if procurement and liability ignore status.
M16 Risk pooling	Lower variance for independent losses	Resistance	Conditional AF only if claims improve pricing faster than capital is consumed	Correlation hidden by a short benign history; kill with provider-wide stress tests.
M17 Insurer balance sheet and reinsurance	Absorbed tail loss	Resistance and survival	Not AF by itself	Capital called a data flywheel; kill if claim learning does not improve risk-adjusted results.
M18 Machine-verifiable procurement evidence	Lower diligence and switching friction	Bargaining power	Conditional AF if more suppliers improve comparison quality without gaming	Marketplace network effects claimed before neutral demand; kill with predictive ranking tests.
M19 Customer-owned portable state	Credible exit and lower hold-up	Bargaining power and resilience	Conditional AF if each migration improves a reusable portability specification	Regulation commoditizes the feature; kill if customers will not pay beyond compliance.
M20 Accredited field sampling	Trusted observations outside the model stack	Resistance and bargaining power	Conditional AF if failure sampling improves reusable inspection protocols	Manual scarcity mistaken for convexity; kill if remote evidence attains equal legal effect.

The surviving mechanisms are mostly not AF. Independent custody, accreditation, balance-sheet capital, and off-host authorization can be durable because they separate duties or carry legal consequences, yet they normally resist shocks rather than gain from them. The most credible positive-convexity candidates are incident replay, reusable control mappings, repeated switching maps, and actuarial learning. Each fails when data rights disappear, labels are disputed, events correlate, native suites retain the observations, or marginal maintenance scales with shock diversity.

Opportunity kill and revival cards

These cards cover functions, not proposed companies. Revenue forms describe possible economic transfer, not market size or desirability.

C01. Prompt-level model routing

Conditions: persistent cost/quality spread and enough volume to repay evaluation overhead. Buyer: application platform or inference team. Scarcity: workload labels and routing policy. Revenue: usage fee or savings share. Dependencies: model APIs, stable evals, telemetry rights. Incumbents: AWS Bedrock, Microsoft Foundry, model gateways. Capital/regulation: low capital; ordinary data and security duties. Reversibility/decay: highly reversible, but option value decays as native routers and fixed-quality prices improve. AF class: resistance to provider variance; no inherent AF. Death: native routing meets buyer tolerances and passes savings through. Revival: application-specific outcomes remain invisible to providers and produce transferable routing lift.

C02. Verification-aware workflow admission

Conditions: consequential actions can be classified before execution and failure costs exceed latency. Buyer: owners of production agents, security, or operations. Scarcity: outcome-linked admission rules rather than model selection. Revenue: workflow subscription or action fee. Dependencies: tool semantics, policy enforcement, replay, human escalation. Incumbents: agent platforms, API gateways, security suites. Capital/regulation: moderate liability; higher in regulated actions. Reversibility/decay: policies are reversible and shadow-testable; value decays if reliable models eliminate pre-action uncertainty. AF class: conditional AF through counterexamples and shadow decisions. Death: admission adds delay without reducing loss. Revival: failures remain jagged and reusable rules outperform model confidence.

C03. Post-deployment evaluation and monitoring

Conditions: drift and unexpected behavior remain material in live context. Buyer: AI operations, risk, audit. Scarcity: production-grounded evaluators with outcome labels. Revenue: monitored workload, evaluated span, or annual control subscription. Dependencies: telemetry access, representative sampling, response rights. Incumbents: AgentCore Evaluations, Google evaluation, observability vendors. Capital/regulation: low financial capital; confidentiality and audit liability. Reversibility/decay: easy to pilot; option expires as native evaluation closes application gaps. AF class: conditional AF if failures improve reusable evaluators. Death: customers accept suite metrics or cannot label outcomes. Revival: NIST-style monitoring gaps persist across platforms and independent methods gain recognized standing.

C04. Conflict-separated evidence ledger

Conditions: a dispute requires records outside the actor being judged. Buyer: legal, audit, regulated operator, insurer. Scarcity: custody, timestamps, retention, and chain of control, not blockchain mechanics. Revenue: retention subscription, evidence package, or dispute event. Dependencies: admissibility, identity binding, export from execution systems. Incumbents: cloud logs, e-discovery, trust-service providers. Capital/regulation: security, privacy, records-management duties. Reversibility/decay: custody is hard to unwind after reliance; entry option decays when contracts standardize on incumbents. AF class: bargaining power and resilience, not AF. Death: ordinary customer-owned logs receive equal evidentiary treatment. Revival: contracts, courts, or carriers require independent custody.

C05. Content and execution provenance

Conditions: buyers care who produced or altered an artifact and under what process. Buyer: media, enterprise governance, procurement. Scarcity: issuer trust and capture integration. Revenue: signing, verification, credential, or compliance fees. Dependencies: C2PA/VC compatibility, key lifecycle, trusted issuers. Incumbents: certificate authorities, platforms, content tools. Capital/regulation: identity and misrepresentation liability. Reversibility/decay: open standards reduce technical lock-in; issuer reputation decays after compromise. AF class: resistance to tampering, not truth and not AF. Death: provenance has no effect on reliance or liability. Revival: accepted issuer status or qualified legal effect attaches to specific assertions.

C06. Agent identity lifecycle registry

Conditions: organizations need sponsors, ownership, birth/death, delegation, and inventory beyond workload IAM. Buyer: IAM and security teams. Scarcity: cross-domain lifecycle evidence. Revenue: managed identity or governed-agent fee. Dependencies: IAM integrations, directories, token standards. Incumbents: Microsoft Entra, cloud IAM, identity governance vendors. Capital/regulation: identity security and privacy duties. Reversibility/decay: deployable in stages; option decays quickly as suites normalize agent objects. AF class: resilience and governance, not AF. Death: agent identity remains a service principal with metadata. Revival: heterogeneous agents cross domains and no system of record can represent sponsor and delegation chains.

C07. External policy decision and enforcement

Conditions: the agent host cannot be trusted to authorize itself. Buyer: security architecture and regulated operators. Scarcity: tamper-separated decisioning near consequential actions. Revenue: authorization decision volume or enterprise license. Dependencies: reliable identity, policy semantics, gateway control. Incumbents: AgentCore Policy, Entra, OPA ecosystems, API gateways, AuthZEN implementers. Capital/regulation: moderate security liability. Reversibility/decay: policy engines are replaceable; standards lower switching cost and proprietary value. AF class: resistance; conditional AF through bounded attack learning. Death: native gateways pass equivalent red-team and audit tests. Revival: cross-platform separation of duties becomes mandatory.

C08. Delegated-intent and mandate receipts

Conditions: an agent acts for a person or organization and later intent can be disputed. Buyer: merchants, financial institutions, enterprise approvers. Scarcity: accepted mandate semantics and liability allocation. Revenue: receipt, verification, or dispute fee. Dependencies: identity, authorization, nonces, revocation, payment rails. Incumbents: Visa, Mastercard, Google/FIDO, wallets. Capital/regulation: payments, identity, consumer-protection exposure. Reversibility/decay: protocol integration is reversible; network acceptance windows close quickly. AF class: bargaining power, with conditional learning from disputes. Death: incumbent network receipts become exclusive trust roots. Revival: cross-network mandates need a neutral verifier accepted by multiple schemes.

C09. Agent payment orchestration

Conditions: agents buy APIs, content, or goods within hard budgets. Buyer: agent developers and finance operations. Scarcity: policy-bound execution across rails. Revenue: transaction or orchestration fee. Dependencies: wallets, payment processors, identity, refunds, spend controls. Incumbents: AgentCore payments, Stripe-related rails, card networks, wallet providers. Capital/regulation: potentially high if touching funds; lower as an orchestration-only vendor. Reversibility/decay: reversible integrations; option decays as platforms bundle microtransactions. AF class: resistance through role separation, not AF. Death: managed platforms expose adequate budgets, audit, and rail choice. Revival: multi-cloud settlement and enterprise separation-of-duties remain unsolved.

C10. Settlement and reconciliation evidence

Conditions: machine-speed transactions create disputes across authorization, delivery, and payment. Buyer: merchant finance, marketplaces, payment operations. Scarcity: joined evidence across rails and execution events. Revenue: reconciled transaction or recovered-loss share. Dependencies: payment data rights, delivery attestations, standardized identifiers. Incumbents: processors, ERP, payment networks. Capital/regulation: data-security and possibly payments regulation. Reversibility/decay: narrow proof can start read-only; value decays if networks own end-to-end telemetry. AF class: resilience; conditional AF if disputes refine reusable matching. Death: one rail provides authoritative joined records. Revival: fragmented rails and machine services create unresolved multi-party causality.

C11. Outcome-based reputation graph

Conditions: verified outcomes predict future performance after task-mix adjustment. Buyer: procurement, marketplaces, insurers. Scarcity: representative outcome labels and manipulation resistance. Revenue: data license, diligence fee, or ranking access. Dependencies: evidence rights, identity, statistical adjustment, appeals. Incumbents: embedded marketplaces, review systems, analyst firms. Capital/regulation: defamation, fake-review, competition, and data liability. Reversibility/decay: scores can be withdrawn, but reputational harm is not easily reversible; decay is rapid under gaming. AF class: conditional AF only when attacks improve defenses faster than trust erodes. Death: out-of-sample prediction fails or cheap Sybils move rank. Revival: legally sourced verified outcomes become portable across procurement channels.

C12. Agent/component marketplace

Conditions: fragmented supply, repeatable packaging, and buyer demand for external discovery. Buyer: developers and enterprise procurement. Scarcity: distribution, trust, or exclusive supply. Revenue: take rate, listing, or enterprise access. Dependencies: deployment compatibility, review, support, billing. Incumbents: Salesforce AgentExchange, AWS Marketplace, cloud catalogs. Capital/regulation: moderate platform and representation liability. Reversibility/decay: listings are reversible; cold-start option expires as embedded marketplaces consolidate. AF class: possible network effects, not AF. Death: incumbents own both workflow and catalog. Revival: cross-platform components require neutral certification or buyers reject vendor-controlled ranking.

C13. Procurement evidence exchange

Conditions: buyers repeatedly request the same controls, benchmarks, and attestations from many suppliers. Buyer: vendor risk and sourcing. Scarcity: accepted evidence schemas and verified issuers. Revenue: supplier subscription, buyer seat, or verified package. Dependencies: procurement systems, standards, credential issuers. Incumbents: GRC suites, cloud marketplaces, audit firms. Capital/regulation: confidentiality and reliance liability. Reversibility/decay: data exchange is reversible; schema lead decays as standards mature. AF class: bargaining power; conditional AF if each supplier increases reusable comparison quality. Death: evidence remains bespoke or embedded channels satisfy it. Revival: RFPs require portable machine-verifiable evidence.

C14. Sovereignty and residency policy broker

Conditions: models, data, and actions face incompatible jurisdictional rules. Buyer: global enterprise compliance and platform teams. Scarcity: executable location/purpose policy and current legal mappings. Revenue: region/workload subscription or compliance service. Dependencies: cloud regions, legal interpretation, data lineage. Incumbents: cloud policy suites, consultancies, sovereign clouds. Capital/regulation: high advice and misconfiguration liability. Reversibility/decay: rules can be staged; mappings expire with legal change. AF class: conditional AF only if new rules reuse a common control graph. Death: marginal jurisdiction work remains bespoke. Revival: fragmentation persists but common control primitives capture most variance.

C15. Context and memory portability

Conditions: state affects outcomes and provider exit requires semantic continuity. Buyer: application owners and data governance. Scarcity: portable meaning, permissions, provenance, and deletion state. Revenue: managed memory, migration, or export subscription. Dependencies: schemas, embeddings, application semantics, privacy rights. Incumbents: Google Memory Bank, agent frameworks, databases. Capital/regulation: privacy and retention exposure. Reversibility/decay: export is reversible; option decays as open formats and Data Act switching normalize portability. AF class: bargaining power and resilience; conditional AF through repeated migrations. Death: transcripts suffice or state is cheap to rebuild. Revival: verified migration preserves outcome quality across materially different runtimes.

C16. Independent benchmark and calibration lab

Conditions: buyers distrust vendor benchmarks and decisions require task-specific calibration. Buyer: procurement, regulators, insurers. Scarcity: representative test access, methods, and trusted independence. Revenue: evaluation project, subscription, or certification. Dependencies: ground truth, secure model access, anti-cheating controls. Incumbents: internal labs, NIST/CAISI, established test laboratories, consultancies. Capital/regulation: professional liability; accreditation burden if claims carry authority. Reversibility/decay: methods can be piloted; option expires if benchmarks saturate or become native. AF class: conditional AF through counterexamples, but evaluation gaming is an opposite force. Death: tests do not predict deployment outcomes. Revival: accredited methods and field monitoring close that gap.

C17. Incident and adversarial exchange

Conditions: participants can share sanitized failures and attacks without creating intolerable exposure. Buyer: security, insurers, sector consortia. Scarcity: trusted membership, normalized incidents, and safe-actionable sharing. Revenue: membership or sector sponsorship. Dependencies: legal safe harbors, privacy, taxonomy, analyst validation. Incumbents: ISACs, security vendors, cloud threat networks. Capital/regulation: breach, confidentiality, antitrust, and disclosure risk. Reversibility/decay: reversible participation; value decays if sharing rights are narrow. AF class: strongest conditional AF candidate if attacks improve shared defenses faster than attacker learning. Death: underreporting, poisoned data, or correlated loss. Revival: sector-specific safe harbors and measurable recurrence reduction.

C18. Compliance control mapping and evidence automation

Conditions: multiple frameworks share controls and evidence can be collected automatically. Buyer: compliance and internal audit. Scarcity: authoritative mappings plus execution evidence. Revenue: enterprise license, control, or audit cycle. Dependencies: standards, policy systems, legal review. Incumbents: GRC suites, consultancies, cloud compliance products. Capital/regulation: advice and audit-reliance liability. Reversibility/decay: modular and reversible; mappings age with rules and harmonized standards. AF class: conditional AF through graph reuse. Death: legal change requires linear manual remapping or suites bundle equivalent mappings. Revival: recognized crosswalks reduce marginal audit cost across platforms.

C19. Accredited AI conformity assessment

Conditions: law or procurement recognizes independent assessment within a defined scope. Buyer: regulated providers, deployers, authorities. Scarcity: accreditation, qualified staff, methods, and liability-bearing signature. Revenue: assessment, surveillance, and renewal. Dependencies: accreditation bodies, standards, jurisdictional recognition. Incumbents: established laboratories, certification and audit firms. Capital/regulation: high professional, accreditation, and quality-system burden. Reversibility/decay: accreditation is slow and partly sunk; option expires if established labs occupy scopes first. AF class: bargaining power, not AF; repeated findings may improve methods. Death: self-assessment remains sufficient. Revival: conformity or procurement explicitly requires recognized third-party evidence.

C20. AI performance warranty analytics

Conditions: a bounded metric, observable exposure, and causal trigger support a warranty without taking primary risk. Buyer: AI vendor, adopter, insurer. Scarcity: technical diligence and trigger design. Revenue: analytics, monitoring, or program fee. Dependencies: telemetry, contracts, carrier capacity, claims data. Incumbents: Munich Re and other carriers, brokers, actuarial firms. Capital/regulation: moderate if no risk retained; high misrepresentation liability. Reversibility/decay: pilotable by narrow use case; option decays as carriers internalize analytics. AF class: conditional learning from claims, not balance-sheet AF. Death: triggers invite gaming or carrier retains all useful data. Revival: carriers need independent technical monitoring across heterogeneous stacks.

C21. Primary outcome insurance

Conditions: losses are measurable, diversified, priceable, and within licensed appetite. Buyer: AI vendor or enterprise adopter. Scarcity: capital, license, distribution, underwriting data, claims capability. Revenue: premium and investment income, offset by losses and expenses. Dependencies: reinsurance, reserving, actuarial credibility, policy wording. Incumbents: Munich Re and established insurance markets. Capital/regulation: maximal among the cards. Reversibility/decay: costly and slow to enter or exit; adverse-selection option decays as incumbents learn. AF class: risk pooling is resistance; only claims learning is conditionally AF. Death: common-model correlation or causality ambiguity defeats pricing. Revival: narrowly bounded independent risks and durable access to underwriting economics.

C22. Claims evidence and technical adjudication

Conditions: disputes require reconstruction of agent decisions, controls, and outcomes. Buyer: carriers, insureds, counsel, regulators. Scarcity: conflict-separated evidence, technical expertise, repeatable causality methods. Revenue: case fee, retainer, or claims-platform license. Dependencies: custody, policy access, expert standards, legal process. Incumbents: forensic firms, carriers, e-discovery, expert witnesses. Capital/regulation: professional liability and confidentiality, but less risk capital than insurance. Reversibility/decay: service can begin case-by-case; option decays if carriers build internal teams. AF class: conditional AF if adjudicated cases create reusable causal patterns with lawful rights. Death: evidence cannot resolve causality or incumbents retain it exclusively. Revival: heterogeneous stacks make independent reconstruction necessary.

Real-options classification and decay

Classification is descriptive, not an instruction to exercise. Every option needs an owner, paid premium, trigger, expiry, and abandonment condition. "Keep everything open" is not an option because carrying cost and attention are unbounded.

Card	Option class	Exercise trigger	Premium and decay	Abandonment signal
C01	Pilot call	Isolated production routing lift	Eval cost; decays with native parity	Net lift vanishes
C02	Staged call	Admission reduces consequential loss	Integration latency; decays with model reliability	Loss unchanged
C03	Monitoring call	Live drift exceeds native detection	Telemetry work; decays with suite coverage	Labels unavailable
C04	Custody call	Contract requires independent evidence	Retention/security; decays when incumbents gain acceptance	No reliance effect
C05	Issuer call	Accepted issuer status emerges	Key/trust operations; standards commoditize mechanics	No procurement effect
C06	Identity call	Cross-domain lifecycle gap is funded	IAM integration; rapid suite decay	Service principals suffice
C07	Separation call	Independent enforcement is mandated	Gateway integration; standards reduce rent	Native controls pass
C08	Network call	Multiple schemes accept one receipt	Protocol/legal work; acceptance window closes	One trust root dominates
C09	Orchestration call	Multi-rail policy gap persists	Payment integration; bundle decay	Managed controls suffice
C10	Recovery-share call	Joined evidence recovers disputes	Data integration; rail consolidation decay	One authoritative record
C11	Data call	Scores predict held-out loss	Label and appeal cost; gaming decay	Prediction fails
C12	Marketplace call	Exclusive supply or neutral demand appears	Cold-start subsidy; expires with channel consolidation	No unique supply
C13	Exchange call	Portable evidence enters RFPs	Schema verification; GRC bundle decay	Evidence stays bespoke
C14	Mapping call	Reuse lowers marginal jurisdiction cost	Legal maintenance; rules expire continuously	Cost stays linear
C15	Switching put	Migration preserves semantic outcomes	Adapter rehearsal; Data Act commoditization	Rebuild is cheap
C16	Accreditation call	Method predicts deployment and gains standing	Lab/method cost; benchmark saturation	No field validity
C17	Consortium call	Safe sharing reduces recurrence	Trust/legal overhead; underreporting decay	Poisoning or correlation
C18	Compound mapping call	Each framework reuses control nodes	Legal review; suite absorption	No marginal reuse
C19	Accreditation call	Recognized scope opens	Large sunk quality-system cost; incumbents may preempt	Self-assessment accepted

Card	Option class	Exercise trigger	Premium and decay	Abandonment signal
C20	Carrier-partner call	Carrier funds external technical diligence	Monitoring build; carrier internalization	Data rights withheld
C21	Capital-intensive entry call	Diversification and reserving are demonstrated	License/capital/reinsurance; adverse selection accelerates	Tail dependence breaks limits
C22	Case-to-method call	Adjudications yield reusable causality	Expert labor; carrier internalization	Cases stay bespoke

Cross-examination of the best survivors

- Incident replay looks strongest in a fragmented, failure-rich world. Opposite world: reliable systems make incidents rare, while the remaining events are novel and correlated. The corpus then loses breadth exactly as tail severity rises. Survival does not establish convexity.
- Reusable compliance mapping benefits from regulatory variation. Opposite world: harmonized standards erase mappings, or sovereign rules require local legal judgment. One world removes demand; the other makes marginal cost linear.
- Independent custody benefits from conflicts and disputes. Opposite world: customer-controlled cloud logs gain routine legal acceptance. Custody remains technically sound but loses bargaining power. Its revival requires a reliance rule, not more cryptography.
- Claims learning appears to compound. Opposite world: the same frontier provider sits beneath many insureds, so one update creates portfolio-wide loss. More exposure produces more data only after consuming capital; this is negative convexity unless limits, exclusions, and reinsurance cap the common shock.

Superintelligence death worlds

World	What dies	Conditional survivors	Survivor death test
Reliable, cheap, plural SI	Routing, routine verification, benchmark and context scarcity	Legal custody, qualified signatures, licensed risk capital	Institutions accept machine self-attestation and self-insurance
Concentrated provider SI	Independent technical control plane, identity, memory, evaluation	Sovereign controls, external courts, customer-owned evidence	Provider becomes accepted principal and adjudicator
Open distributed SI	Software rents, proprietary adapters, generic intelligence reputation	Accreditation, physical inspection, liability-bearing signatories	Machine evidence gets equal constitutive legal effect
Sovereign fragmented SI	Cross-border learning, neutral global reputation, pooled evidence	Local qualified providers and jurisdictional custody	Sovereigns standardize reciprocal trust or prohibit private intermediaries
Adversarial or unstable SI	Uncapped automation and naive risk pools	Off-host authorization, exposure limits, restoration, forensic custody	Attacks defeat separation or common losses exceed capital

SI does not automatically increase demand for every trust layer. In the first three worlds it destroys most technical intermediation. In the last two it may increase control demand while simultaneously making neutrality, data pooling, and insurance less viable. Legal authority, physical verification, and capital survive only while human institutions insist on separated duties and can enforce them.

Post-hoc collision with the supplied deck

This analysis was derived before reviewing /Users/freedom/Downloads/Odylith AI.pdf. The deck's adaptive inference, verification, collective learning, and outcome-underwriting layers map to C01-C03, M04/M11, and C20-C21. Its public benchmark is evidence for a repo-grounded operating policy and write-admission behavior, not isolated proof of per-step model allocation, cross-tenant learning, calibrated underwriting, generalization, or economic scarcity. Its illustrative mature revenue envelope is arithmetic based on assumed enterprise counts and prices, not demand evidence. The software beachhead is coherent as an experiment surface, but transfer from code workflows to regulated outcome risk remains an unproven assumption. The deck does not remove the incumbent, data-rights, causality, correlation, accreditation, or capital kill branches.

Decisions deliberately unmade

No opportunity is ranked. No company boundary, financing path, price, market size, launch sequence, jurisdiction, accreditation scope, risk-retention level, or partner is selected. No mechanism is treated as a moat merely because it survives a shock. The unresolved decision gates are the falsifiers in A01-A25 and the exercise/abandonment triggers above.

Round 3 independent adjudication - historical pre-closure critique

Status note. This adjudication is preserved verbatim as the pre-closure critique. The integrated source answers its fourteen required corrections in the Round 3 correction-closure ledger; the critique below is not a claim that those defects remain open.

Date: 2026-08-05

Scope and adjudicative standard

This adjudication reviews `Odyolith_Market_Evolution_Situational_Awareness.md` and all six Round 2 records: enterprise authority, incumbents, market economics, mechanism design, superintelligence resistance, and antifragility, falsification, and real options. The last record is treated as part of the final evidence set, including its existential kill tree, A01-A25 falsification register, M01-M20 mechanism audit, C01-C22 kill and revival cards, and option-decay tests.

The standard here is non-prescriptive. The question is not what Odyolith should build, whether a company should exist, or whether an investor should act. It is whether the dossier has produced a sufficiently unanchored and causally discriminating world model. The tests are: internal consistency; independence from the originating thesis; separation of function, institution, market, and company; explicit treatment of null worlds; evidence proportionality; and conditional rather than rhetorical opportunity claims.

Source locators below name exact dossier headings and, where useful, exact claims. Round 2 locators name the corresponding memo and debate. These citations are internal evidence references, not endorsements of the external sources cited inside those records.

Executive determination

The dossier is materially stronger than the prior strategy frame it criticizes. It explicitly states in `# Mandate: build a world model before choosing a strategy` that it does not assume a router, control plane, marketplace, clearinghouse, verification layer, governance company, software beachhead, or outcome business must exist. It correctly says in `# Present topology: hierarchy before market` that internal allocation is hierarchy and in `# Open standards: interoperability is not contestability` that message compatibility does not establish independent supply, portable authority, settlement, or remedies. The Round 2 mechanism memo supplies the strongest discipline in the corpus: independent principals, funded demand, executable allocation, liability, payment, settlement, default, remedies, and contestability are required before market language is earned.

Those are real corrections, not cosmetic ones. The dossier also distinguishes marginal inference cost from total system expenditure, nominal provider plurality from customer-control concentration, technical verification from institutional acceptance, resilience from antifragility, and a persistent function from a separately monetizable company. Its superintelligence worlds include genuine death conditions for routing, external assurance, human review, and rich control layers.

The adverse finding is that the synthesis has not fully escaped the originating frame. Its disclaimers are less anchored than its taxonomy. The title is still `Odyolith Market Evolution Situational Awareness`; the first ten entries of the `# Opportunity atlas: conditional surfaces, not selections` map closely onto an agent control-plane decomposition; and many later entries are the institutional complements that such a decomposition would need. The document says the deck is only a hypothesis generator, but the generated opportunity inventory still resembles a broadened perimeter around that deck. This is not proof that any category is wrong.

The antifragility memo supplies useful counterevidence to a claim of simple deck copying: its `## Post-hoc collision with the supplied deck` says its analysis was derived before reviewing the deck, and its independently generated C01-C22 functions overlap materially with the dossier. That overlap raises confidence that several functions are real analytical objects. It does not prove complete frame independence because the memo's starting scope already asks whether an independent economic position can exist around routing, context, verification, authorization, governance, evidence, reputation, procurement, settlement, insurance, and trust. It is blind to the deck but not blind to an Odyolith-adjacent problem statement. A genuinely independent second taxonomy must begin from actors, welfare, loss allocation, public goods, labor, households, coercion, and

non-market institutions rather than from candidate trust-infrastructure functions.

The second adverse finding is that governance and authority receive many death-world qualifications but remain the dossier's default residual scarcities. The synthesis repeatedly moves value from cognition to context, credentials, proof, authority, legitimacy, physical resources, or consent. It does acknowledge that each can die, yet it does not model enough worlds in which enforceable governance is absent, irrelevant, captured, non-commercial, or unable to keep pace. The phrase in # `What superintelligence resistance actually means that some mechanisms may remain costly to bypass` is properly conditional; the opportunity atlas often turns the same conditional functions back into named economic surfaces before separate budget, supplier access, and rent durability have been shown.

The third finding is more favorable: the dossier generally does not confuse hierarchy with market in its definitions, but it sometimes lets market-shaped opportunity language leak back into scenarios where only procurement, internal allocation, or a platform supply chain exists. The mechanism memo corrects this more rigorously than the main synthesis. The main dossier should import that stricter vocabulary everywhere, not only in # `When agent activity becomes a market`.

The fourth finding is that the dossier explicitly rejects standards-as-contestability, but it has not fully operationalized the rejection. 024. Protocol infrastructure still assigns scarcity to "Common syntax and network effects." Common syntax may be a public good with no appropriable rent, and network effects can accrue to the dominant runtime, registry, identity root, or distribution platform rather than to protocol infrastructure. The correction must distinguish interoperability, portability, multi-homing, demand control, governance control, and economic contestability as separate variables.

Agreements that survive cross-examination

1. Cheap intelligence and concentrated control can coexist

All six Round 2 records support the dossier's rejection of a one-dimensional abundance story. The market-economics memo shows that falling accepted-task cost can coexist with frontier fixed costs, committed spend, data gravity, and platform bundling. The incumbent memo shows that open or cheap model supply can still sit behind a few identity, context, billing, and runtime owners. The superintelligence memo makes the same point through 01. Cheap inference, dominant provider and 013. Open weights, centralized power. The antifragility memo's existential kill tree adds the strongest commercial null: cheap reliable intelligence, converged protocols, customer-owned policy, suite bundling, incumbent settlement, accredited conformity, and incumbent risk capital can jointly leave no necessary independent control point.

The dossier's # `Causal loop 1: cheap intelligence, rebound, and total spend` is therefore directionally sound. Its strongest claim is not that rebound must occur, but that falling unit cost and rising aggregate expenditure are compatible. The explicit falsifier - mature cohorts showing saturated task volume, falling total risk-adjusted spend, and no complementary migration - is appropriate. A required caution remains: the current co-movement of efficiency, capex, and electricity does not by itself identify price-induced rebound. Demand anticipation, competitive overbuild, training loads, and non-AI demand can produce similar observations.

2. Provider plurality and platform concentration are orthogonal

The dossier and Round 2 agree that counting model names is a poor concentration measure. # `Provider evolution: capability, distribution, and enclosure` correctly locates control in contracts, identity, authoritative context, telemetry, billing, defaults, and committed spend. Round 2 market economics states the tie-break directly: the key question is who owns qualified demand, authoritative context, and the right to act. Round 2 incumbents adds that multiple enclosed systems of record can coexist without creating a neutral cross-system authority.

This agreement should be retained. Dependency-adjusted concentration is more probative than vendor counts, downloads, or protocol membership. The dossier also correctly notes that acquisition can validate a function while disproving an independent category.

3. Governance need does not imply governance company

The corpus repeatedly reaches the same important distinction. # Opportunities strengthened by greater intelligence says more consequential action can increase demand for constraints, but the function may be absorbed by a trusted provider or state. # Opportunities that disappear as intelligence improves says technical redundancy is not required for company death because bundling, internalization, regulation, or systems-of-record control can erase an independent budget. Round 2 market economics sharpens the point: institutional authority may remain scarce while governance software becomes plumbing.

This is one of the dossier's most valuable conclusions. It must, however, become a formal gate on every opportunity entry rather than remain a warning around the atlas.

4. Machine activity is not market formation

Present topology: hierarchy before market, # When agent activity becomes a market, and # Machine commerce: a contingent sequence agree with the Round 2 mechanism memo. A million internal routing decisions can remain one firm's hierarchy. A platform that assembles components behind one customer contract is vertical production. A directory or protocol can lower discovery cost without creating executable exchange. A clearing layer cannot manufacture underlying competitive obligations.

The mechanism memo's staged test is the best-developed causal instrument in the source set. The first plausible market transition is a funded principal issuing a binding RFQ to ownership-independent suppliers under acceptance, liability, and payment terms. Even that is a thin procurement market, not proof of continuous liquidity or autonomous machine commerce. This distinction should govern all market claims in the dossier.

5. Verification shifts rather than settles the control point

The enterprise-authority, incumbent, market-economics, mechanism, and superintelligence records all distinguish technical verification from accepted authority. Perfect proof can commoditize output checking while increasing the power of specification authors, key custodians, trust roots, courts, or platforms that define recognized evidence. Provider-native telemetry can be the most complete technical record and still be conflicted or institutionally insufficient. Independent evidence can be required and still reinforce concentration if only large providers can fund access and compliance.

The dossier's S11 claim that proof does not create consent or legal mandate therefore survives. So does the opposite warning in S13 and the superintelligence memo's 016. Self-verifying outputs, strategically deceptive process: a perfect proof can validate the selected objective while the process manipulates which objective or facts enter scope.

6. Resistance, resilience, bargaining power, and antifragility are distinct

What superintelligence resistance actually means and # Antifragility under model and institutional shocks correctly reject redundancy as proof of antifragility. The useful test is whether a bounded shock creates more credible routes, better-calibrated limits, lower switching friction, or greater bargaining power. The source set also correctly identifies fake diversification: multiple logos over a common cloud, insurance without claims capacity, open syntax with proprietary identity, or human oversight without time and authority.

The dedicated antifragility record makes this agreement substantially stronger. Its strict chain is shock -> protected downside -> captured learning or exercise right -> measurable gain -> retained rights -> bounded failure. It finds that most supposed antifragile mechanisms are actually resistance, resilience, bargaining power, or learning. Its most credible conditional positive-convexity candidates are incident replay, reusable control mappings, repeated switching maps, and actuarial learning, and each dies when rights disappear, labels are disputed, events correlate, native suites retain observations, or marginal maintenance scales with shock diversity. The dossier should import this stricter test rather than summarize antifragility as a list of generally attractive properties.

Contradiction ledger

Tension	What the dossier says	Adjudication	Required resolution
Unanchored mandate vs Odylith-shaped taxonomy	# Mandate: build a world model before choosing a strategy rejects assuming a router or control plane, while # Opportunity atlas: conditional surfaces, not selections begins with routing, orchestration, context, identity, authorization, receipts, verification, observability, policy, and FinOps.	The antifragility memo's pre-deck derivation shows the overlap is not merely copied, but its starting question already names most of the same trust-infrastructure functions. The disclaimer remains broader than the search space.	Build a second taxonomy blind to both deck and function list from actors, welfare, loss pools, public goods, labor, households, coercion, and non-market allocation; then disclose overlap and omissions.
Governance death worlds vs governance as residual scarcity	# Causal loop 2: scarcity migration is not guaranteed gives each scarcity a death world, yet # Opportunities strengthened by greater intelligence and many atlas entries repeatedly return to authority, proof, policy, consent, and legitimacy.	Conditional language does not eliminate base-rate anchoring. Governance survives analytically in more entries than its null worlds receive.	Add worlds where authority is unenforceable, non-commercial, publicly supplied, provider-sovereign, or made irrelevant by institutional collapse. Assign each governance surface a named authority-death trigger.
Hierarchy discipline vs market-shaped opportunity language	# When agent activity becomes a market is strict, but scenario and opportunity descriptions sometimes move from agent volume or protocols to auctions, payments, reputation, insurance, and clearing as adjacent surfaces.	The main synthesis risks implying a progression that the mechanism memo says can stop permanently at hierarchy or bilateral procurement.	Add a formation verdict to every scenario and opportunity: hierarchy, product supply chain, directory, procurement market, spot market, or market infrastructure. Ban "market" where the formation test fails.
Standards are not contestability vs protocol scarcity	# Open standards: interoperability is not contestability is explicit, while 024. Protocol infrastructure identifies "Common syntax and network effects" as scarcity.	A standard can create compatibility while commoditizing its implementation and concentrating traffic elsewhere. Syntax is not necessarily scarce or appropriable.	Split protocol adoption, implementation concentration, registry governance, identity portability, extension portability, and traffic contestability. Remove implied rent from syntax alone.
Scenario plurality vs a purported phase space	# Scenario phase space: eighteen causal worlds says worlds can coexist, but S1-S8 concern ownership/resource regimes, S9-S14 concern trust and capability properties, and S15-S18 concern tempo or heterogeneity.	These are overlapping dimensions, not eighteen comparable alternatives. Calling them a phase space can hide unmodeled cross-products.	Recast them as dimensions and compose explicit combined worlds, including trusted sovereign SI, open but embodied scarcity, and fast takeoff inside a fragmented bloc system.
No probability assignment vs opportunity salience	The dossier assigns no scenario probabilities but presents 25 named opportunity surfaces with formation conditions and revenue mechanisms.	Naming and detailing creates salience and menu effects even without rankings. Readers can mistake coverage for plausibility.	Add base-rate ignorance, evidence distance, time-window, and correlated-death fields. Include non-opportunity outcomes with equal detail.
Technical function vs separate budget	The dossier warns that importance does not identify who captures value, yet each opportunity entry supplies a payer and possible revenue mechanism.	A possible meter is not evidence of a budget, procurement category, renewal, or supplier power.	Add four separate statuses: function exists; buyer suffers a measurable loss; separate budget exists; independent supplier retains rent after bundling.
Current law vs durable authority	Enterprise authority cites present human accountability and employer duties; the dossier extrapolates authority as a likely scarcity in several worlds.	Present law is a boundary condition, not strong evidence under machine standing, emergency powers, fast takeoff, or legitimacy collapse.	Mark legal extrapolations as transition-sensitive and add explicit legal-discontinuity worlds.

Tension	What the dossier says	Adjudication	Required resolution
Antifragility definition vs candidate-list compression	# Antifragility under model and institutional shocks correctly distinguishes resilience, while # Real options and temporal strategy without selecting a product calls portability, qualification, components, experiments, participation, partnerships, and exit rights examples of options.	The Round 2 antifragility memo imposes a stricter test: a real option needs an owner, paid premium, trigger, expiry, exercise terms, and abandonment rule; antifragility needs protected downside, retained learning rights, measurable gain, and bounded failure. Several dossier examples are capabilities or aspirations, not options.	Import the strict option and positive-convexity fields. Relabel items that are only resilience, bargaining power, or narrative optionality.
Claimed Round 3 completion vs unincorporated adjudication	# Adversarial review architecture says "The dossier was built through three separated rounds" and that Round 3 asked an independent adjudicator to find corrections.	The dossier claims completion before the resulting corrections have been incorporated or answered.	State that Round 2 was completed and Round 3 is pending incorporation; claim three completed rounds only after this adjudication is addressed.
Real options vs shared dependency	# Real options and temporal strategy without selecting a product lists multi-provider qualification and portable contracts as options.	Nominal options can fail simultaneously if they share identity, cloud, power, law, or settlement. Reversibility is not independence.	Add dependency-adjusted exercisability and exercise-time relative to capability or crisis speed.
Opportunity robustness vs monetization	# Contradiction-robust opportunity mechanisms names portability, reversible authority, checkable evidence, customer keys, dependency transparency, and external enforcement.	These are design properties or public-policy qualities, not opportunity mechanisms. They may impose cost without yielding revenue.	Rename them "cross-world robustness properties" and avoid any implication that robustness identifies a commercial surface.

Hidden shared assumptions

Institutions remain coherent enough to matter

Even skeptical worlds usually preserve recognizable enterprises, states, courts, insurers, procurement offices, property rights, and budgets. S16 and S17 weaken them, but the dossier mostly treats institutional delay, capture, or emergency hierarchy as a transition to another authority holder. A missing null is institutional incoherence: fragmented enforcement, conflicting jurisdictions, bankrupt counterparties, persistent informal economies, or a state that cannot compel the actor controlling infrastructure. In such worlds, governance is neither a durable scarcity nor a bundled feature. It is an unenforced claim.

Authority is identifiable and attributable

The dossier often asks who controls authority, which presumes that authority can be scoped, recorded, and attributed. Round 2 enterprise authority correctly distinguishes authentication from actual authority, consent, fiduciary duty, and benefit. The synthesis should go further: principals may be ambiguous, mandates may be contested after the fact, coalitions may disagree internally, and autonomous systems may split acts across thresholds or jurisdictions. "Who authorized this?" may have no stable answer even when logs are perfect.

Economic demand remains human-welfare-linked

The accepted-outcome denominator assumes an identifiable buyer objective. The dossier does not sufficiently model synthetic demand: agents generating tasks, transactions, content, reputation, and defensive responses for other agents without increasing human welfare. Such an economy can display rising calls, payments, energy, market depth, and even accounting revenue while destroying real surplus. The internal-GMV and retry-demand confounders are useful but too narrow. The missing variable is welfare-adjusted final demand.

Scarcity migrates rather than demand disappearing

Causal loop 2: scarcity migration is not guaranteed verbally rejects guaranteed migration, but much of the document still searches for the next bottleneck. A different world has no compensating scarcity with a commercial budget: intelligence reduces transaction and coordination costs, organizations capture little incremental revenue, prices fall, and aggregate profits compress. Another has productivity gains that become leisure, public benefit, lower prices, or producer surplus captured by existing asset owners. The analytical endpoint need not be a new paid layer.

Organizations can absorb and measure change

The dossier includes organizational absorption as a state variable and recognizes that time saved need not become P&L. Yet most opportunity and budget paths assume a buyer can define accepted outcomes, reallocate labor, measure full cost, and maintain a coalition long enough to procure a solution. Organizational politics may make measurement impossible, not merely slow. A function can be valuable, observable in a trial, and permanently unfunded because no owner can capture the distributed benefit or accept the concentrated liability.

Counterparties stay solvent and losses remain compensable

Insurance, warranties, disputes, settlement, and clearing all presume losses can be bounded, priced, and remedied. The dossier names correlated and catastrophic failures but underweights judgment-proof actors, cross-border non-enforcement, state immunity, provider bankruptcy, and harms that money cannot reverse. The existence of evidence and liability does not ensure compensation.

More information improves decisions

The leading-indicator and antifragility systems assume that evidence can update policy before the regime changes. Under strategic deception, Goodhart effects, political suppression, or fast takeoff, added telemetry may improve the actor's concealment, overwhelm decision makers, or arrive after the option expires. Learning rate relative to environmental change is more important than information volume.

The firm and state remain the principal units

The source set under-models households, worker organizations, cooperatives, open-source communities, criminal networks, informal institutions, and transnational civil society. This biases authority toward CIOs, CISOs, regulators, providers, and insurers. It also biases opportunities toward enterprise procurement even when value or harm is dispersed across people who do not control the budget.

Missing-world additions

The following worlds should be added before the scenario set can claim broad situational coverage. They are not forecasts and need no assigned probabilities.

M1. AI investment bust and capability plateau

Capability improvements slow, enterprise value remains mostly individual augmentation, and overbuilt infrastructure is written down. Models remain useful but no broad autonomous economy forms. Seats consolidate into ordinary software, agent operations remain niche, and assurance budgets shrink with action scope. This world differs from slow takeoff because the binding event is disappointed economics, not merely gradual progress. It directly tests # Causal loop 1, the budget-migration path, and most of the opportunity atlas.

M2. Public-good and internalization equilibrium

Important controls exist but are supplied through open source, regulation, internal enterprise teams, professional norms, or public infrastructure. There is no durable independent vendor category. This world is required because the dossier repeatedly distinguishes function from company but does not give the no-company equilibrium equal scenario weight. It would absorb policy, identity formats, conformance, evidence schemas, and some verification without denying their importance.

M3. Labor and democratic veto

Workers, professions, consumers, and voters constrain deployment despite favorable benchmark economics. Bargaining, licensing, collective action, due-process claims, or distributional conflict make certain substitutions illegitimate. The binding scarcity is not generic "institutional trust" but consent by groups that bear concentrated harm and may not control the purchasing budget. This world can produce moratoria, mandated human roles, public ownership, or redistribution without a technical catastrophe.

M4. Synthetic-demand and Goodhart economy

Agents create requests, content, payments, reviews, security events, and verification work for other agents. Reported activity and revenue rise while human welfare, productivity, or fiscal capacity stagnates. Platforms benefit from metered circular activity; enterprises struggle to distinguish necessary work from machine-generated overhead. Market liquidity metrics can look healthy even after wash activity is removed because the underlying objectives themselves are misaligned. The key denominator becomes welfare-adjusted terminal outcome, not accepted machine task.

M5. Informal, criminal, and adversarial agent economy

Autonomous activity scales outside compliant identity, audit, payment, and procurement rails. Stolen credentials, coercion, unlicensed models, informal settlement, ransomware, evasion, and jurisdiction shopping become material. Governance demand may rise for compliant actors while effective authority shrinks system-wide. This world attacks the assumption that institutional control expands with consequence and that machine commerce becomes more legible as it matures.

M6. Institutional fracture and non-compensable loss

Courts conflict, states impose capital controls, providers or insurers fail, and harms cross borders faster than judgments. Evidence may be perfect while remedy is unavailable. Liability attaches to insolvent shells or parties outside reach. Clearing and insurance cannot stabilize the system because loss exceeds capital or finality is politically reversed. This is distinct from S18's orderly polycentric mosaic: fragmentation here destroys reciprocal recognition.

M7. Machine legal standing and autonomous asset control

Law recognizes machine-managed entities, autonomous trusts, or direct machine capacity to hold assets and obligations. Human principals become remote or optional. The dossier's principal-based delegation model must then be rebuilt around constitutional limits, capitalization, revocation, and dissolution rather than attribution to a human click. This world directly reverses the claim in S3 that institutions attach authority to human or legal principals understood in the current sense.

M8. Non-market abundance and political allocation

Digital intelligence becomes abundant enough that prices and software markets contract, while states, communities, or dominant providers allocate access as a public entitlement, constitutional service, or rationed utility. The key conflict is distribution, not monetization. This world differs from S9 because trusted intelligence need not belong to one provider and from S8 because scarce rights may be resolved politically rather than traded.

M9. Catastrophic infrastructure regression

War, grid failure, cyber conflict, climate disaster, or supply-chain rupture reduces compute, communications, and institutional capacity. Advanced systems persist unevenly, but the dominant task is degraded operation rather than acceleration. Offline tools, manual skills, and local energy matter; markets and governance may simplify or disappear. This world tests whether the dossier's option set works under shrinking, not expanding, technical capacity.

Causal gaps and evidence overreach

Rebound is plausible but not identified

The evidence in # Evidence baseline as of 5 August 2026 establishes simultaneous efficiency gains and aggregate infrastructure growth. It does not establish that lower inference prices caused the growth or that mature enterprise cohorts will continue expanding accepted tasks. The required correction is to distinguish price elasticity, capability-induced new demand, speculative capacity build, training demand, and substitution from other workloads. The falsifier should also include infrastructure write-downs and persistent low utilization, not only task saturation.

Bundling evidence establishes attack capability, not enclosure success

The dossier and incumbent memo document that major platforms offer routing, identity, policy, observability, marketplaces, and governance. This proves the ability and incentive to bundle. It does not prove customer adoption, adequacy, retention, or foreclosure. # Provider evolution: capability, distribution, and enclosure should separate "feature exists," "feature is default," "feature controls consequential volume," and "feature prevents independent attachment." Without that ladder, acquisition and release activity can be overread as equilibrium evidence.

Current legal rules are weak evidence under regime change

The enterprise-authority memo uses current natural-person duties, deployer obligations, consumer protections, and automated-contracting rules carefully. The main dossier sometimes carries them farther through scenario language about legal principals, consent, and accountability. Under S16, S17, M6, or M7, these are variables, not anchors. Present law can identify transition friction but cannot support medium-confidence claims that authority or liability remains scarce under superintelligence.

The scenario set lacks interaction causality

Each of S1-S18 has a trigger, play-by-play, bottleneck, and falsifier, but the scenarios are mostly one-path narratives. They do not show how one state variable changes another's sign. For example, open standards can strengthen hyperscalers only if identity, hosting, and distribution remain concentrated; regulation can open supply only if compliance artifacts are reusable and procurement lots remain accessible; perfect verification can decentralize generation only if specifications and recognition are not captured. The dossier should express these as conditional interactions, not just adjacent prose.

"Strengthened by greater intelligence" is too coarse

Opportunities strengthened by greater intelligence says broader credentials, irreversible actions, consequence volume, and physical scope can strengthen constraints. Greater capability can also reduce error rates, compress incident frequency, improve native controls, or concentrate delegation in a trusted system. The relevant independent variable is not intelligence level but consequence-adjusted autonomous exposure minus native prevention and reversibility. The heading should be changed or the claim narrowed.

Opportunity revenue language exceeds available evidence

The repeated statement that a possible revenue mechanism does not establish willingness to pay is good discipline. Still, assigning a payer and meter to every surface creates an unsupported impression of economic completeness. A payer label such as CIO, CISO, regulator, or insurer does not show procurement authority, incremental budget, renewal, or willingness to add a vendor. The atlas needs explicit evidence states and a null result: no separate payment even though the function exists.

Antifragility is well falsified but not yet empirically established

Antifragility under model and institutional shocks gives a valid high-level definition, and the Round 2 antifragility memo supplies the missing rigor. Its ## Antifragility mechanism audit classifies most candidates as resistance, resilience, bargaining power, or learning. It identifies only conditional positive-convexity paths and gives each a kill test. That is strong falsification architecture, not empirical proof that agent systems have already gained from shocks. The dossier should therefore import the memo's gain unit, retained-rights test, positive-convexity interval, and bounded-failure condition, while continuing to label the surviving mechanisms as hypotheses. Post-incident spending, a larger log corpus, survival, and restored service are not gains by themselves.

Opportunity conditionality: required gate structure

The opportunity atlas should not be deleted, ranked, or converted into a recommendation. It should be made more conditional. The Round 2 antifragility memo already demonstrates the right form through its existential kill tree, A01-A25 decisive falsifiers, C01-C22 death and revival conditions, and option abandonment signals. The dossier should incorporate that structure rather than leave it in a separate appendix. Every O1-O25 entry should pass seven separate gates:

- 1. Functional gate: the function remains necessary in the stated world.
- 2. Loss gate: a buyer experiences a measurable loss or foregone gain without it.
- 3. Authority gate: that buyer can change the relevant decision or workflow.
- 4. Budget gate: a separate, recurring budget exists after native controls are configured.
- 5. Access gate: an independent supplier receives enough telemetry, enforcement rights, or ground truth to perform.
- 6. Rent gate: bundling, standards, internalization, and competition do not drive price below sustainable full cost.
- 7. Time gate: the window lasts longer than procurement, integration, certification, and capital formation.

The dossier currently covers parts of gates 1, 5, 6, and 7 through formation, dependencies, incumbent response, and death conditions. It does not consistently distinguish gates 2 through 4. It also evaluates death conditions mostly one at a time. Correlated death matters: one trusted integrated provider can simultaneously eliminate routing, external policy, independent evidence, portable reputation, and cross-venue settlement. Conversely, institutional fragmentation can strengthen local compliance demand while destroying global scale and evidence portability.

The antifragility memo makes the correlated commercial null explicit: all routes to durable scarcity can close at once even while many functions remain useful. It also shows why a revival condition is not a recommendation. Breaking one terminal assumption may create a control point, but the resulting function still must pass payer, access, rent, and time gates. The dossier should preserve that existential asymmetry: one surviving function is enough to keep a world analytically open, but not enough to establish a company, market size, or investable return.

For market-facing surfaces - O11 machine payments, O12 procurement auctions and RFQs, O13 portable reputation, O14 insurance, O15 dispute resolution, and O16 settlement and clearing - add the mechanism memo's formation verdict. Protocol messages, mandates, quotes, certifications, and payment initiation are complements, not proof of an underlying market. Each entry should report ownership-adjusted counterparties, executable offers, net external settlement, default experience, remedies, and bypass.

For standards-facing surfaces - especially O24 - add a contestability stack: open specification; multiple conforming implementations; portable identity and data; substitutable registries and trust roots; non-proprietary extensions; successful full-workflow switching; distributed traffic and revenue; and enforceable anti-self-preference. Passing the first layer is not evidence of passing the last.

For governance-facing surfaces - O4 through O9, O23, and O25 - add an authority-null test: what happens if provider-native controls are accepted, policy is internalized, legal enforcement lags, human review becomes ceremonial, or the SI/operator becomes the de facto sovereign? A function that survives only as an embedded provider duty is not an independent opportunity surface.

High-leverage variables

Variable	Why it dominates several debates	Minimum observable
Ownership-adjusted control concentration	Breaks the tie between nominal plurality and real outside options across model, cloud, identity, context, verifier, and settlement layers.	Share of consequential workflows movable without losing identity, memory, evidence, or contract economics.
Institutionally accepted verification cost	Separates technical proof from evidence that courts, insurers, regulators, and counterparties will act on.	Cost, latency, and dispute outcome for native versus independent evidence by consequence class.

Variable	Why it dominates several debates	Minimum observable
Authority reversibility	Determines whether delegation creates manageable exposure or irreversible principal-agent loss.	Time to revoke, rollback, compensate, and restore state after a contested action.
Task contractibility and damage latency	Determines whether routing, auctions, outcome contracts, verification, and insurance can form.	Ex ante acceptance completeness plus the share of material loss discovered after settlement.
Welfare-adjusted final demand	Distinguishes productive rebound from synthetic workload, retries, defensive traffic, and circular machine activity.	Human or enterprise outcome net of compute, review, externalities, and machine-generated overhead.
Organizational capture rate	Determines whether adoption becomes a P&L budget or remains seats, slack, work intensification, or unallocated quality.	Mature-cohort change in output, cost base, demand, and risk after full implementation cost.
Capability-change to institution-update ratio	Determines whether law, policy, insurance, and learned controls can remain relevant.	Time for capability or behavior change divided by time for detection, rule change, deployment, and enforcement.
Dependency-adjusted option exercisability	Separates real antifragility from multi-logo redundancy.	Successful failover under loss of the shared cloud, identity root, power region, or legal rail.
Retained learning rights and convexity	Determines whether shocks create an owned reusable asset or merely impose losses and public learning.	Out-of-sample gain per bounded incident after data-rights, labeling, remediation, and attacker-adaptation costs.
Externality incidence and loss-bearing capacity	Determines whether liability, insurance, and market prices discipline behavior.	Who bears realized loss, available capital, exclusions, recovery rate, and non-compensable harm.
Standard-stack contestability	Determines whether openness distributes power or lowers complement cost for incumbents.	Switching and traffic concentration across implementation, registry, identity, extension, and settlement layers.
Physical actuation coupling	Determines whether digital abundance translates into economic output or remains bounded by machines, energy, land, and safety.	Share of value requiring scarce physical access and the reversibility of physical action.

These variables should be tracked as interactions. No single threshold resolves the world. For example, cheap accepted verification supports open competition only when task specifications are adequate, trust roots are plural, and full workflows can switch. Otherwise it can strengthen a monopoly specification owner.

Required corrections to the dossier

1. In # Adversarial review architecture, replace the claim that three separated rounds have already built the dossier. State that Round 2 is complete across six records and this Round 3 adjudication awaits incorporation. Claim three completed rounds only after the corrections are resolved in the integrated dossier.
2. In # Mandate: build a world model before choosing a strategy and # Prior-dossier anchoring audit, add a reproducible anti-anchoring test. Preserve the antifragility memo's pre-deck derivation as evidence, but acknowledge that its scope already names an Odylith-adjacent function set. Generate a second taxonomy from actors, welfare, public goods, household and labor interests, failure modes, coercion, and non-market allocation before comparing it with the thesis. Report category overlap and omissions rather than relying on a disclaimer.
3. Rename # Scenario phase space: eighteen causal worlds as a scenario library or reconstruct it as an actual multidimensional phase space. Cross capability ownership, trust, institutional speed, physical coupling, and geopolitical structure. Add M1-M9 or equivalent missing worlds.
4. In # Causal loop 1: cheap intelligence, rebound, and total spend, separate observed co-movement from causal rebound. Add demand elasticity, speculative overbuild, utilization, training versus inference, and write-down evidence.

5. In # Causal loop 2: scarcity migration is not guaranteed, add a no-migration endpoint in which demand, profit, or willingness to pay disappears rather than moving to another layer.
6. In # Open standards: interoperability is not contestability and 024. Protocol infrastructure, replace standards adoption as a proxy with the full contestability stack. Remove any implication that common syntax alone is scarce or monetizable.
7. Import the Round 2 mechanism memo's formation verdict into # Machine commerce: a contingent sequence, every scenario's market-structure field, and O11-O16. Label hierarchy, platform supply chain, directory, procurement market, spot market, and infrastructure distinctly.
8. In # Opportunities strengthened by greater intelligence, replace intelligence level with consequence-adjusted autonomous exposure and explicitly model native prevention, trusted-provider absorption, and falling incident rates.
9. Rename # Contradiction-robust opportunity mechanisms to "Cross-world robustness properties." State that portability, customer keys, reversible authority, and checkable evidence can be socially valuable without creating a supplier, budget, or rent.
10. Revise # Opportunity atlas: conditional surfaces, not selections and O1-O25 to include the seven gates, evidence status, correlated death conditions, and a no-separate-budget outcome. Import the Round 2 antifragility memo's existential kill tree, decisive falsifier, revival condition, option premium, expiry, and abandonment fields. Give non-company and public-good equilibria equal space.
11. In # What superintelligence resistance actually means, add institutional collapse, machine standing, persuasion of key holders, and response-time failure as first-class nulls rather than edge qualifications. Do not treat current property and legal authority as durable without a scenario-specific enforcement mechanism.
12. In # Antifragility under model and institutional shocks, import the Round 2 chain shock -> protected downside -> captured learning or exercise right -> measurable gain -> retained rights -> bounded failure. Label candidate mechanisms as unverified hypotheses, identify the positive-convexity interval, and require before-and-after out-of-sample evidence. Distinguish resistance, recovery, survival, learning, and bargaining power from antifragility.
13. In # Leading-indicator system, add welfare-adjusted final demand, utilization and write-downs, public or worker veto, institutional response time, non-compensable loss, and full-workflow switching. Keep agent counts, protocol messages, and gross payments explicitly subordinate.
14. In # Evidence and confidence ledger, add claims for frame independence, governance survival under SI, antifragility, standards contestability, market formation, and separate opportunity budgets. Current confidence in frame independence and independently monetized governance should be low until the new tests are passed.

Final boundary

This adjudication does not choose a product, company strategy, financing path, or investment. It finds that the dossier has achieved strong conceptual discipline on hierarchy, concentration, verification, antifragility, and conditionality, but has not yet proven independence from an Odylith-shaped control-plane frame. Its most important next correction is not a new opportunity ranking. It is to expand the null worlds, convert overlapping scenarios into interacting dimensions, force every opportunity through separate function-budget-access-rent-time gates, and make the strict mechanism-design and positive-convexity vocabularies govern the entire document.

The dossier should preserve its central warning: abundant intelligence does not determine who captures value. It should add the equally important warning that value need not migrate to a new commercial layer at all, governance need not remain enforceable, open standards need not create contestability, and machine activity need not become a market.

Primary-source bibliography

1. Microsoft FY26 Q3 earnings. <https://www.microsoft.com/en-us/investor/events/fy-2026/earnings-fy-...>
2. Microsoft FY2025 annual report. <https://www.microsoft.com/investor/reports/ar25/index.html>
3. IEA Energy and AI. <https://www.iea.org/reports/energy-and-ai>
4. IEA key questions on energy and AI. <https://www.iea.org/reports/key-questions-on-energy-and-ai/executive...>
5. Stanford AI Index 2026 - economy. <https://hai.stanford.edu/ai-index/2026-ai-index-report/economy>
6. Epoch inference price trends. <https://epoch.ai/data-insights/llm-inference-price-trends>
7. Salesforce FY26 results. <https://investor.salesforce.com/news/news-details/2026/Salesforce-De...>
8. ServiceNow AI Control Tower. <https://newsroom.servicenow.com/press-releases/details/2026/ServiceN...>
9. GitHub agent control plane. <https://github.blog/changelog/2026-02-26-enterprise-ai-controls-agen...>
10. AWS AgentCore policy. <https://aws.amazon.com/about-aws/whats-new/2026/03/policy-amazon-bed...>
11. Google Agent Gateway. <https://cloud.google.com/blog/products/identity-security/introducing...>
12. IBM agentic control plane. <https://www.ibm.com/new/announcements/introducing-the-agentic-contro...>
13. Cloudflare dynamic routing. <https://developers.cloudflare.com/ai-gateway/features/dynamic-routing/>
14. Cloudflare AI Gateway pricing. <https://developers.cloudflare.com/ai-gateway/reference/pricing/>
15. Palo Alto Networks acquisition of Portkey. <https://www.paloaltonetworks.com/company/press/2026/palo-alto-networ...>
16. Okta agentic enterprise blueprint. <https://investor.okta.com/news-and-events/news-releases/news-details...>
17. NIST identity and authority for software agents. <https://www.nist.gov/news-events/news/2026/02/new-concept-paper-iden...>
18. NIST AI Agent Standards Initiative. <https://www.nist.gov/artificial-intelligence/ai-agent-standards-init...>
19. NIST monitoring deployed AI. <https://www.nist.gov/news-events/news/2026/03/new-report-challenges-...>
20. NIST evaluation cheating. <https://www.nist.gov/caisi/cheating-ai-agent-evaluations>
21. NIST continuous monitoring proof. <https://www.nist.gov/news-events/news/2026/06/nist-mathematical-proo...>
22. EU AI Act framework. <https://digital-strategy.ec.europa.eu/en/policies/regulatory-framewo...>
23. EU AI Omnibus 2026. <https://digital-strategy.ec.europa.eu/en/news/ai-omnibus-enters-force>
24. EU Product Liability Directive. <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:32024L2853>
25. Linux Foundation A2A. <https://www.linuxfoundation.org/press/a2a-protocol-surpasses-150-org...>
26. Anthropic MCP donation / Agentic AI Foundation. <https://www.anthropic.com/news/donating-the-model-context-protocol-a...>
27. Stripe Machine Payments Protocol. <https://stripe.com/blog/machine-payments-protocol>
28. Visa MPP specification. <https://corporate.visa.com/en/sites/visa-perspectives/innovation/vis...>
29. Mastercard Agent Pay for machines. <https://www.mastercard.com/us/en/news-and-trends/press/2026/june/mas...>
30. FTC AI partnerships report. <https://www.ftc.gov/news-events/news/press-releases/2025/01/ftc-issu...>
31. UK CMA cloud investigation. <https://www.gov.uk/cma-cases/cloud-services-market-investigation>
32. OECD AI adoption. <https://www.oecd.org/en/about/news/announcements/2026/01/ai-use-by-i...>
33. ILO GenAI and jobs 2025. <https://www.ilo.org/publications/generative-ai-and-jobs-2025-update>
34. IMF AI gains distribution 2026. <https://www.imf.org/en/publications/wp/issues/2026/07/10/aggregate-g...>
35. Anthropic Responsible Scaling Policy. <https://www.anthropic.com/responsible-scaling-policy>
36. OpenAI Preparedness Framework. <https://openai.com/index/updating-our-preparedness-framework/>
37. OpenAI Frontier Governance Framework. <https://openai.com/index/openai-frontier-governance-framework/>
38. UK AISI frontier trends. <https://www.aisi.gov.uk/frontier-ai-trends-report>
39. EU AI factories. <https://digital-strategy.ec.europa.eu/en/policies/ai-factories>
40. International Federation of Robotics demand. <https://ifr.org/news/global-robot-demand-in-factories-doubles-over-1...>

41. US AI Action Plan 2025. <https://www.whitehouse.gov/wp-content/uploads/2025/07/Americas-AI-Ac...>
42. NAIC artificial intelligence. <https://content.naic.org/insurance-topics/artificial-intelligence>
43. Bank of England AI in financial services. <https://www.bankofengland.co.uk/report/2024/artificial-intelligence-...>
44. BIS Principles for Financial Market Infrastructures. <https://www.bis.org/cpmi/publ/d101a.pdf>
45. OMB M-25-22 AI acquisition. <https://www.whitehouse.gov/wp-content/uploads/2025/02/M-25-22-Drivin...>
46. NBER enterprise GenAI field experiment. <https://www.nber.org/papers/w33795>
47. W3C Verifiable Credentials 2.0. <https://www.w3.org/standards/history/vc-data-model/>
48. IETF RFC 8693 token exchange. <https://datatracker.ietf.org/doc/rfc8693/>
49. Google AP2. <https://cloud.google.com/blog/products/ai-machine-learning/announcin...>
50. FAR reverse auctions. <https://login.acquisition.gov/far/subpart-17.8>
51. DOJ 2023 Merger Guidelines - platforms. <https://www.justice.gov/atr/merger-guidelines/applying-merger-guidel...>
52. EU Digital Markets Act. <https://eur-lex.europa.eu/eli/reg/2022/1925/2022-10-12/eng>
53. UNCITRAL automated contracting model law. <https://uncitral.un.org/en/mlac>
54. Delaware corporate directors. <https://delcode.delaware.gov/title8/c001/sc04/>
55. CFPB Regulation E. <https://www.consumerfinance.gov/rules-policy/regulations/1005/6/>

Decisions deliberately unmade

This dossier does not choose a beachhead, product, pricing model, revenue target, financing plan, marketplace sequence, company boundary, or investment verdict. It does not decide that routing, verification, governance, authority, underwriting, clearing, software, or any other surface belongs to Odylith. It does not assign probabilities to superintelligence timelines.

What remains is a structured field of uncertainty:

- Capability can commoditize while infrastructure and total spend concentrate.
- Agent proliferation can create open exchange or deepen internal hierarchy.
- Open standards can increase contestability or incumbent distribution.
- Better intelligence can strengthen constraints or internalize them.
- Authority can remain institutionally scarce or migrate toward a trusted provider or state.
- Verification can become more valuable, become bundled, or become free.
- Revenue can appear in a transition layer that later disappears.
- Antifragility requires independent, exercisable options; diversification labels are insufficient.

The next intellectual task is not to defend a thesis. It is to watch the leading variables, run discriminating experiments, update the scenario ledger, and preserve the ability to move when one world begins to dominate.